

Topics: save/load models, evaluate models, models for prediction

```
In [ ]: import tensorflow as tf
import numpy as np
```

```
In [ ]: tf.random.set_seed(1234)
np.random.seed(1234) } useful if we want to modify something in the model  
(with the same feed we can compare)
```

```
In [ ]: # Only if you want to use Google Drive in Colab to save your models
from google.colab import drive
```

```
In [ ]: # Mount Google Drive.
# After running the cell, click to the link and follow the instructions to obtain your authorization code.
# Once the operation is completed you can access your Google Drive home in '/content/drive/My Drive/'  
drive.mount('/content/drive')
```

```
In [ ]: # We can now save files in Drive
with open('/content/drive/My Drive/Keras2/test.txt', 'w') as f:  
    f.write('Hello Google Drive!')
```

```
In [ ]: !cat /content/drive/My\ Drive/Keras2/test.txt
```

The model depend on the initialization, maybe we change the model and the initialization and we obtain better results because of the init.

Let's try to save a model. In this example we will create a model, and we will save the model in memory (Google Drive, if you are using Colab, or your local memory)

This time we create the model in a SEQUENTIAL way (equivalent)

```
In [ ]: # Create base model (e.g., Input -> Hidden -> Out)
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28, 28))) # or as a list
model.add(tf.keras.layers.Dense(units=64, activation=tf.keras.activations.sigmoid))
model.add(tf.keras.layers.Dense(units=10, activation=tf.keras.activations.softmax)) ← classification layer
```

layer	Output Shape	params
Flatten	(None, 784)	0
dense	(None, 64)	50240
dense 1	(None, 10)	650

bases  
 $(784 \times 64 + 64)$   
 $(64 \times 10 + 10)$

```
In [ ]: # Visualize created model as a table
model.summary() →
```

```
In [ ]: # Visualize initialized weights
tf.reduce_sum(model.layers[1].weights[0])
```

```
In [ ]: # Create a random input
random_input = tf.random.uniform([1, 28, 28]) ✓ we're simulating a single image => batchsize = 1
```

```
In [ ]: # Compute the output given the random input (this is used later to compare saved and loaded models)
output = model(random_input)
```

```
In [ ]: # We have two ways to save the model (see slides):
1. Save the entire model.
# When you load the model, it will be already a Model object which can be used for training/evaluation.
# Very useful for example for web/android applications.
# - save
#       model.save('/PATH/TO/YOUR/MODEL')
# - Load
#       Loaded_model = tf.keras.models.load_model('/PATH/TO/YOUR/MODEL')

2. Save only the weights
# In this case we are only saving the weights of the model.
# To load the model (the weights) we need to create the model again and redefine the same model structure of the saved model.
# - save
#       model.save_weights('/PATH/TO/YOUR/MODEL/WEIGHTS/')
# - Load
#       Loaded_model = tf.keras.Sequential()
# ...
#       loaded_model.load_weights('/PATH/TO/YOUR/MODEL/WEIGHTS/')
```

In [ ]: # Let's try the first method
# Save the model
model.save('/content/drive/My Drive/Keras2/saved\_model') # use your local directory if you are not using Drive

In [ ]: del model

In [ ]: # Load the model (with this way, the loaded-model is ready to be used)
loaded\_model = tf.keras.models.load\_model('/content/drive/My Drive/Keras2/saved\_model') # use your local directory if you are not using Drive

In [ ]: # Check the output of the loaded model is the same of the output computed before saving
np.testing.assert\_allclose(loaded\_model(random\_input), output)

1. Save the entire model

## 2. Save only the weights

```
In [ ]: # Let's try the second method (only weights)
loaded_model.save_weights('/content/drive/My Drive/Keras2/saved_model_weights') # use your local directory if you are not using Drive

In [ ]: del loaded_model

In [ ]: # We have to create the model again

new_model = tf.keras.Sequential()
new_model.add(tf.keras.layers.Flatten(input_shape=(28, 28))) # or as a list
new_model.add(tf.keras.layers.Dense(units=64, activation=tf.keras.activations.sigmoid))
new_model.add(tf.keras.layers.Dense(units=10, activation=tf.keras.activations.softmax))

In [ ]: # Check that the weights are not the same of the ones of the original model
tf.reduce_sum(new_model.layers[1].weights[0])

In [ ]: # Load saved weights
new_model.load_weights('/content/drive/My Drive/Keras2/saved_model_weights')

In [ ]: # Check again the weights. Now should be the same.
tf.reduce_sum(new_model.layers[1].weights[0])

# Check the output of the Loaded model is the same of the output computed before saving
np.testing.assert_allclose(loaded_model(random_input), output)
```

this MUST be the same model structure

```
In [ ]: # Once the model is loaded we can use it for example to evaluate the performance on a new test set or to predict an output given specific input.
# Let's try to use a loaded model for evaluation.
# Please notice that in this example I will use a model I trained before. You have to replace it with your trained model.
```

```
In [ ]: (x_train, y_train), (x_test, y_test) = tf.keras.datasets.fashion_mnist.load_data()
```

```
In [ ]: model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28, 28))) # or as a list
model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu))
model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu))
model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu))
model.add(tf.keras.layers.Dense(units=10, activation=tf.keras.activations.softmax))
```

```
In [ ]: model.load_weights('/content/drive/My Drive/Keras2/classification_experiments_/FC_Oct13_22-51-04/ckpts/cp_10.ckpt')
```

```
In [ ]: # You have to compile also before using it for evaluation
```

```
# Optimization params
# ----

# Loss
loss = tf.keras.losses.CategoricalCrossentropy()

# Learning rate
lr = 1e-5
optimizer = tf.keras.optimizers.Adam(learning_rate=lr)
# ----

# Validation metrics
# ----

metrics = ['accuracy']
# ----

# Compile Model
model.compile(optimizer=optimizer, loss=loss, metrics=metrics)
```

```
In [ ]: # evaluate method is similar to the fit one, except that no training is performed.
# The output is computed for each sample (batch) of your test set and the losses/metrics are evaluated.
# Finally, the overall losses and metrics are returned.
```

```
eval_out = model.evaluate(x=x_test / 255.,
                           y=tf.keras.utils.to_categorical(y_test),
                           batch_size=8,
                           verbose=1)
```

name of  
the loaded/  
trained model

```
In [ ]: # In the following an example of how to use the model for prediction. (# evaluation)
# In this example a drawn (literally with Paint :) image is used. Just for fun!
# You will find it in the class folder.
```

```
# Compute output given x
from PIL import Image

shoe_img = Image.open('/content/drive/My Drive/Keras2/shoe.png').convert('L')
```

it prints 2 values:  
[•, •] = [loss on the  
test set, accuracy  
on the test set]

```
In [ ]: shoe_img.resize((256, 256))

In [ ]: shoe_arr = np.expand_dims(np.array(shoe_img), 0) # to obtain the batch dimension
        out_softmax = model.predict(x=shoe_arr / 255.)
        out_softmax # output is the probability distribution (softmax)
```

0 T-shirt/top \ 1 Trouser \ 2 Pullover \ 3 Dress \ 4 Coat \ 5 Sandal \ 6 Shirt \ 7 Sneaker \ 8 Bag \ 9 Ankle boot \

```
In [ ]: # Get predicted class as the index corresponding to the maximum value in the vector probability
        predicted_class = tf.argmax(out_softmax, 1)
        predicted_class
```

## 2/2 Topics: Overfitting, model checkpoints, TensorBoard

```
In [ ]: import tensorflow as tf
import numpy as np

In [ ]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

In [ ]: tf.random.set_seed(1234)
np.random.seed(1234)

In [ ]: from google.colab import drive

In [ ]: drive.mount('/content/drive')

In [ ]: # Load built-in dataset
# -----
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.fashion_mnist.load_data()

In [ ]: # Split in training and validation sets
# e.g., 50000 samples for training and 10000 samples for validation
x_valid = x_train[50000:]
y_valid = y_train[50000:]
x_train = x_train[:50000]
y_train = y_train[:50000]

In [ ]: # Create Training Dataset object
# -----
train_dataset = tf.data.Dataset.from_tensor_slices((x_train, y_train))

# Shuffle
train_dataset = train_dataset.shuffle(buffer_size=x_train.shape[0])

# Normalize images
def normalize_img(x_, y_):
    return tf.cast(x_, tf.float32) / 255., y_

train_dataset = train_dataset.map(normalize_img)

# 1-hot encoding <- for categorical cross entropy
def to_categorical(x_, y_):
    return x_, tf.one_hot(y_, depth=10)

train_dataset = train_dataset.map(to_categorical)

# Divide in batches
bs = 32
train_dataset = train_dataset.batch(bs)

# Repeat
# Without calling the repeat function the dataset will be empty after consuming all the images
train_dataset = train_dataset.repeat()

In [ ]: # Create Validation Dataset
# -----
valid_dataset = tf.data.Dataset.from_tensor_slices((x_valid, y_valid))
valid_dataset = valid_dataset.map(normalize_img) # Normalize images
valid_dataset = valid_dataset.map(to_categorical) # 1-hot encoding
valid_dataset = valid_dataset.batch(bs) # Divide in batches
valid_dataset = valid_dataset.repeat() # Repeat

In [ ]: # Create Test Dataset
# -----
test_dataset = tf.data.Dataset.from_tensor_slices((x_test, y_test))
test_dataset = test_dataset.map(normalize_img) # Normalize images
test_dataset = test_dataset.map(to_categorical) # 1-hot encoding
test_dataset = test_dataset.batch(1) # Divide in batches

In [ ]: # In the following we create a different model based on the 'which_model' param:
# - base model (FC)
# - base model + dropout (FC_Dropout)
# - base model + weight decay (FC_WD)
# - base model + dropout + weight decay (FC_Dropout_WD)

# Please notice that in this example a deeper model is used w.r.t. the one of the previous class
# (with an higher number of parameters).
# This is to highlight the problem of overfitting.

which_model = 'FC'

if which_model == 'FC':
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Flatten(input_shape=(28, 28))) # or as a list
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dropout(0.5))
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dropout(0.5))
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dropout(0.5)) 10 classes
    model.add(tf.keras.layers.Dense(units=10, activation=tf.keras.activations.softmax))
elif which_model == 'FC_Dropout':
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Flatten(input_shape=(28, 28))) # or as a list
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu))
    model.add(tf.keras.layers.Dropout(0.5))
```

```

model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu))
model.add(tf.keras.layers.Dropout(0.5))
model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu))
model.add(tf.keras.layers.Dropout(0.5))
model.add(tf.keras.layers.Dense(units=10, activation=tf.keras.activations.softmax))

elif which_model == 'FC_WD':
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Flatten(input_shape=(28, 28))) # or as a list
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dense(units=10, activation=tf.keras.activations.softmax))

elif which_model == 'FC_Dropout_WD':
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Flatten(input_shape=(28, 28))) # or as a list
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dropout(0.5))
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dropout(0.5))
    model.add(tf.keras.layers.Dense(units=512, activation=tf.keras.activations.relu,
                                    kernel_regularizer=tf.keras.regularizers.l2(0.001)))
    model.add(tf.keras.layers.Dropout(0.5))
    model.add(tf.keras.layers.Dense(units=10, activation=tf.keras.activations.softmax))

```

In [ ]:

```

# Visualize created model as a table
model.summary()

# Visualize initialized weights
# model.weights

```

In [ ]:

```

# Optimization params
# -----
loss      = tf.keras.losses.CategoricalCrossentropy()    # Loss
lr       = 1e-3                                         # Learning rate
optimizer = tf.keras.optimizers.Adam(learning_rate=lr)  # Optimizer

# Validation metrics
# -----
metrics = ['accuracy']

# Compile Model
# -----
model.compile(optimizer=optimizer, loss=loss, metrics=metrics)

```

We do this here  
so we can see  
live what's  
happening

In [ ]:

```

# This cell is useful if you want to visualize Tensorboard in your Colab notebook.
# Please run the cell before training, otherwise you have to wait the training cell is completed before running this one

# If you are using jupyter notebook, you can skip this cell and open Tensorboard with the following:
# From terminal: tensorboard --logdir /PATH/TO/YOUR/EXPERIMENTS/ --port PORT
# Go to 127.0.0.1:PORT in your web browser

%load_ext tensorboard
%tensorboard --logdir /content/drive/My\ Drive/ANN&DL/classification_experiments_\ --port 6009

```

We train the  
model s.t. it'll save  
the model status at  
each training epoch.  
(So we're also able to  
visualize what's  
going on)

In [ ]:

```

import os
from datetime import datetime

cwd = '/content/drive/My Drive/ANN&DL' # use your local directory if you are not using Drive

exp_dir = os.path.join(cwd, 'classification_experiments_')
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

now = datetime.now().strftime('%b%d_%H-%M-%S')

exp_name = which_model

exp_dir = os.path.join(exp_dir, exp_name + '_' + str(now))
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

callbacks = []

# Model checkpoint
# -----
ckpt_dir = os.path.join(exp_dir, 'ckpts')
if not os.path.exists(ckpt_dir):
    os.makedirs(ckpt_dir)

ckpt_callback = tf.keras.callbacks.ModelCheckpoint(filepath=os.path.join(ckpt_dir, 'cp_{epoch:02d}.ckpt'),
                                                    save_weights_only=True) # False saves the entire model
callbacks.append(ckpt_callback)

# Visualize Learning on Tensorboard
# -----
tb_dir = os.path.join(exp_dir, 'tb_logs')
if not os.path.exists(tb_dir):
    os.makedirs(tb_dir)

# By default shows losses and metrics for both training and validation
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=tb_dir,
                                              profile_batch=0,
                                              histogram_freq=1) # if 1 shows weights histograms
callbacks.append(tb_callback)

```

This is how we save  
the model in an  
automatic way  
during trainings

```

# Early Stopping
# -----
early_stop = False
if early_stop:
    es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=10)
    callbacks.append(es_callback)

# -----
model.fit(x=train_dataset,
          epochs=100, ##### set repeat in training dataset
          steps_per_epoch=int(np.ceil(x_train.shape[0] / bs)),
          validation_data=valid_dataset,
          validation_steps=int(np.ceil(x_valid.shape[0] / bs)),
          callbacks=callbacks)

# How to visualize Tensorboard
# 1. tensorboard --logdir EXPERIMENTS_DIR --port PORT      <- from terminal
# 2. Localhost:PORT <- in your browser

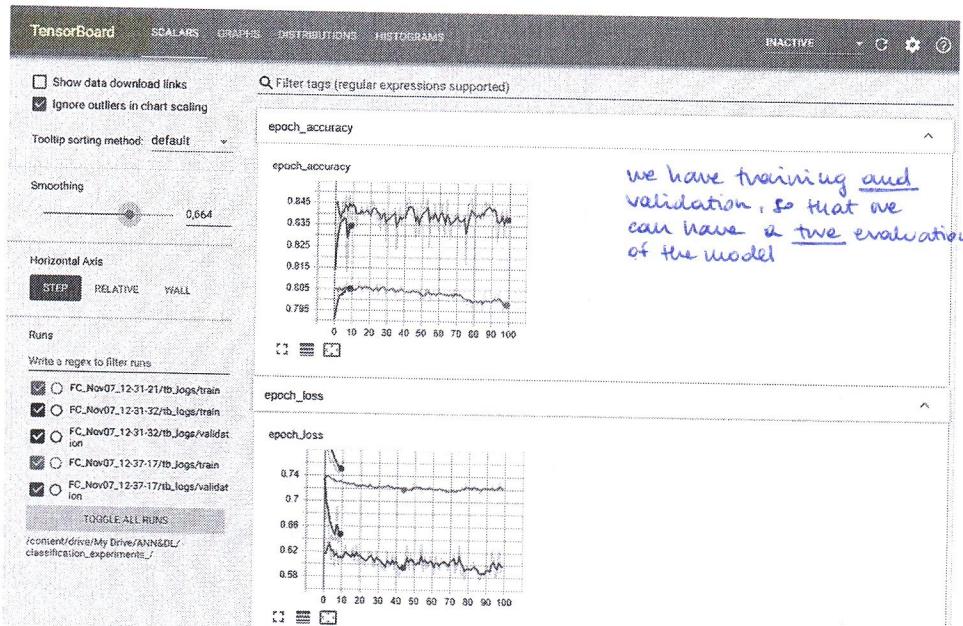
```

Tolerance (= how many epochs we want to wait until we stop the training)

if we want it  $\Rightarrow$  True

here we can specify which is the quantity we rely on for the early stopping (we can choose ACCURACY or LOSS)

To prevent overfitting



### Reduce the overfitting:

- reduce the number of parameters
- **early stopping**
- **dropout** (based on the idea that an ensemble of models helps in preventing overfitting  $\Rightarrow$  if we take the average prediction of different models we are able to obtain better results than a single model) (it's implemented by using some noise in the network. This corresponds to randomly turn off some neurons  $\rightarrow$  This is implemented through: `tf.keras.layers.Dropout(p)` (that comes after the activation function), where p is the probability for each neuron to be turned off.)

- **weight decay** (we can prevent some weights to acquire higher levels. In this way we limit the model complexity. For example we can say that we want  $\min \|w\|_2^2$ : how can we let it know? We add something to the loss function. We introduce also a regularization  $\gamma$  (useful in the case in which our initial weights are already too high)  $\Rightarrow$   $\text{loss} = \sum_{n=1}^N (t_n - q(\theta_n w))^2 + \gamma \|w\|_2^2$ )

- **data augmentation**

