

Topics: Global Average Pooling (GAP) + class activation maps

```
In [1]: from IPython.core.interactiveshell import InteractiveShell  
InteractiveShell.ast_node_interactivity = "all"
```

```
In [12]: import os  
import tensorflow as tf  
import numpy as np
```

```
SEED = 1234
tf.random.set_seed(SEED)

# Get current working directory
print(os.getcwd())
```

```
In [3]: # Run this cell only if you are using Colab with Drive  
from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [4]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

apply\_data\_augmentation = False

4. Czech training TransportCompetition

```
train_data_gen = ImageDataGenerator(rescale=1./255)
```

```
# Create validation and test ImageDataGenerator objects
valid_data_gen = ImageDataGenerator(rescale=1./255)
test_data_gen = ImageDataGenerator(rescale=1./255)
```

```
In [5]: # Run this cell only if you are using Colab with Drive  
!unzip '/content/drive/My Drive/MaskDataset.zip'
```

```
Archive: /content/drive/My Drive/MaskDataset.zip
creating: MaskDataset/
creating: MaskDataset/test/
creating: MaskDataset/test/with_mask/
inflating: MaskDataset/test/with_mask/1-with-ma...
inflating: MaskDataset/test/with_mask/106-with-...
inflating: ..
```

```
In [6]: !ls /content
drive MaskDataset sample data
```

- training
- test
- val

```
In [2]: # Create generators to read images
```

```
#-----  
dataset_dir = os.path.join(os.getcwd(), 'dataset')
```

10.1007/s00339-010-0626-0

**bs = 32**

```
* tiny shape  
img_h = 256  
img_w = 256
```

`num_classes=2`

```
decide_class_indices = True  
if decide_class_indices:
```

```
    classes = ['without_mask', # 0
                'with_mask', # 1
            ]
```

```
    else:  
        classes=None
```

### *# Training*

```
training_dir = os.path.join(dataset_dir, 'train')
```

# Validation

```
# Test
test_dir = os.path.join(dataset_dir, 'test')
test_gen = test_data_gen.flow_from_directory(test_dir,
```

} "in this way we decide the order of the classes  
 (a without mark" will be associated to 0 and  
 "with mark" will be associated to 1)

```

batch_size=bs,
target_size=[img_h, img_w],
classes=classes,
class_mode='categorical',
shuffle=False,
seed=SEED)

```

Found 1315 images belonging to 2 classes.  
 Found 142 images belonging to 2 classes.  
 Found 194 images belonging to 2 classes.

In [8]: # Check how keras assigned the labels  
`train_gen.class_indices`

Out[8]: {'with\_mask': 1, 'without\_mask': 0}

In [9]: # Create Dataset objects  
# -----  
# Training  
`train_dataset = tf.data.Dataset.from_generator(lambda: train_gen,`  
`output_types=(tf.float32, tf.float32),`  
`output_shapes=([None, img_h, img_w, 3], [None, num_classes]))`  
`train_dataset = train_dataset.repeat()`  
# Validation  
# -----  
`valid_dataset = tf.data.Dataset.from_generator(lambda: valid_gen,`  
`output_types=(tf.float32, tf.float32),`  
`output_shapes=([None, img_h, img_w, 3], [None, num_classes]))`  
`valid_dataset = valid_dataset.repeat()`  
# Test  
# -----  
`test_dataset = tf.data.Dataset.from_generator(lambda: test_gen,`  
`output_types=(tf.float32, tf.float32),`  
`output_shapes=([None, img_h, img_w, 3], [None, num_classes]))`  
`test_dataset = valid_dataset.repeat()`

In [10]: # Architecture: Features extraction -> Classifier

```

start_f = 8
depth   = 5
model   = tf.keras.Sequential()
encoder = tf.keras.Sequential()

# Features extraction
for i in range(depth):

    if i == 0:
        input_shape = [img_h, img_w, 3]
    else:
        input_shape=[None]

    # Conv block: Conv2D -> Activation -> Pooling
    encoder.add(tf.keras.layers.Conv2D(filters=start_f,
                                       kernel_size=(3, 3),
                                       strides=(1, 1),
                                       padding='same',
                                       input_shape=input_shape))
    encoder.add(tf.keras.layers.ReLU())

```

```

if i < depth - 1:
    encoder.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2)))
    start_f *= 2

```

```

model.add(encoder)

# GAP
model.add(tf.keras.layers.GlobalAveragePooling2D())

# Classifier
model.add(tf.keras.layers.Dense(units=num_classes, activation='softmax'))

```

In [11]: # Visualize created model as a table  
`model.summary()`

```

# Visualize initialized weights
# model.weights

```

Model: "sequential"

Layer (type)	Output Shape	Param #
sequential_1 (Sequential)	(None, 16, 16, 128)	98384
global_average_pooling2d (Gl)	(None, 128)	0
dense (Dense)	(None, 2)	258

Total params: 98,642  
 Trainable params: 98,642  
 Non-trainable params: 0

If we're at the last convolutional block  
 we're not doing the maxpooling

— instead of the Flatten

the summary is global and so we don't  
 see what it is in the first sequential()

In [12]: `model.layers[-1].weights[0].shape`

Out[12]: `TensorShape([128, 2])`

In [13]: # Optimization params  
# -----  
# Loss

```

loss = tf.keras.losses.CategoricalCrossentropy()
# Learning rate
lr = 1e-3
optimizer = tf.keras.optimizers.Adam(learning_rate=lr)

# Validation metrics
# -----
metrics = ['accuracy']

# Compile Model
model.compile(optimizer=optimizer, loss=loss, metrics=metrics)

```

In [14]:

```
%load_ext tensorboard
%tensorboard --logdir /content/drive/My\ Drive/Keras4/segmentation_experiments_GAP/#

Output hidden; open in https://colab.research.google.com to view.
```

In [15]:

```

import os
from datetime import datetime

cwd = os.getcwd()

exp_dir = os.path.join('/content/drive/My Drive/Keras4', 'segmentation_experiments_GAP')
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

now = datetime.now().strftime('%b%d_%H-%M-%S')

model_name = 'CNN_GAP'

exp_dir = os.path.join(exp_dir, model_name + '_' + str(now))
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

callbacks = []

# Model checkpoint
# -----
ckpt_dir = os.path.join(exp_dir, 'ckpts')
if not os.path.exists(ckpt_dir):
    os.makedirs(ckpt_dir)

ckpt_callback = tf.keras.callbacks.ModelCheckpoint(filepath=os.path.join(ckpt_dir, 'cp_{epoch:02d}.ckpt'),
                                                    save_weights_only=True) # False to save the model directly
callbacks.append(ckpt_callback)

```

```
# Visualize Learning on Tensorboard
# -----
tb_dir = os.path.join(exp_dir, 'tb_logs')
if not os.path.exists(tb_dir):
    os.makedirs(tb_dir)

# By default shows losses and metrics for both training and validation
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=tb_dir,
                                              profile_batch=0,
                                              histogram_freq=1) # if 1 shows weights histograms
callbacks.append(tb_callback)

```

```
# Early Stopping
# -----
early_stop = True
if early_stop:
    es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=10)
    callbacks.append(es_callback)
```

In [16]:

```

model.fit(x=train_dataset,
           epochs=100, ##### set repeat in training dataset
           steps_per_epoch=len(train_gen),
           validation_data=valid_dataset,
           validation_steps=len(valid_gen),
           callbacks=callbacks)

# How to visualize Tensorboard
# 1. tensorboard --logdir EXPERIMENTS_DIR --port PORT      <- from terminal
# 2. Localhost:PORT   <- in your browser

```

```

Epoch 1/100
42/42 [=====] - 9s 216ms/step - loss: 0.6815 - accuracy: 0.5939 - val_loss: 0.6545 - val_accuracy: 0.5845
Epoch 2/100
42/42 [=====] - 9s 210ms/step - loss: 0.5823 - accuracy: 0.6935 - val_loss: 0.6718 - val_accuracy: 0.5915
Epoch 3/100
42/42 [=====] - 9s 208ms/step - loss: 0.5560 - accuracy: 0.7255 - val_loss: 0.4881 - val_accuracy: 0.7887
Epoch 4/100
42/42 [=====] - 9s 206ms/step - loss: 0.4507 - accuracy: 0.8000 - val_loss: 0.2880 - val_accuracy: 0.9437
Epoch 5/100
42/42 [=====] - 9s 206ms/step - loss: 0.3829 - accuracy: 0.8297 - val_loss: 0.2646 - val_accuracy: 0.9296
Epoch 6/100
42/42 [=====] - 9s 206ms/step - loss: 0.2772 - accuracy: 0.8943 - val_loss: 0.1578 - val_accuracy: 0.9587
Epoch 7/100
42/42 [=====] - 9s 204ms/step - loss: 0.1834 - accuracy: 0.9422 - val_loss: 0.1284 - val_accuracy: 0.9366
Epoch 8/100
42/42 [=====] - 8s 199ms/step - loss: 0.2393 - accuracy: 0.9095 - val_loss: 0.0877 - val_accuracy: 0.9789
Epoch 9/100

```

since it's a binary task we could also use a single neuron as classifier (and a sigmoid), in that case we would use the **Binary Crossentropy()** as loss

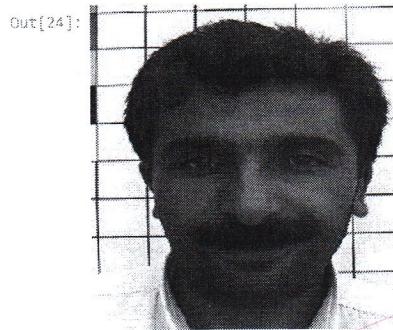
```
42/42 [=====] - 9s 208ms/step - loss: 0.1161 - accuracy: 0.9658 - val_loss: 0.0591 - val_accuracy: 0.9718
Epoch 10/100
42/42 [=====] - 9s 211ms/step - loss: 0.1178 - accuracy: 0.9529 - val_loss: 0.0560 - val_accuracy: 0.9859
Epoch 11/100
42/42 [=====] - 8s 201ms/step - loss: 0.1205 - accuracy: 0.9506 - val_loss: 0.0543 - val_accuracy: 0.9859
Epoch 12/100
42/42 [=====] - 9s 207ms/step - loss: 0.0845 - accuracy: 0.9757 - val_loss: 0.0351 - val_accuracy: 0.9859
Epoch 13/100
42/42 [=====] - 9s 210ms/step - loss: 0.0905 - accuracy: 0.9681 - val_loss: 0.0354 - val_accuracy: 0.9859
Epoch 14/100
42/42 [=====] - 9s 206ms/step - loss: 0.1006 - accuracy: 0.9650 - val_loss: 0.0308 - val_accuracy: 0.9930
Epoch 15/100
42/42 [=====] - 9s 212ms/step - loss: 0.1576 - accuracy: 0.9399 - val_loss: 0.1638 - val_accuracy: 0.9437
Epoch 16/100
42/42 [=====] - 9s 214ms/step - loss: 0.0983 - accuracy: 0.9681 - val_loss: 0.1070 - val_accuracy: 0.9718
Epoch 17/100
42/42 [=====] - 9s 207ms/step - loss: 0.0834 - accuracy: 0.9734 - val_loss: 0.0344 - val_accuracy: 0.9859
Epoch 18/100
42/42 [=====] - 9s 203ms/step - loss: 0.0625 - accuracy: 0.9810 - val_loss: 0.0182 - val_accuracy: 1.0000
Epoch 19/100
42/42 [=====] - 9s 209ms/step - loss: 0.0646 - accuracy: 0.9817 - val_loss: 0.0195 - val_accuracy: 1.0000
Epoch 20/100
42/42 [=====] - 9s 206ms/step - loss: 0.0561 - accuracy: 0.9840 - val_loss: 0.0188 - val_accuracy: 1.0000
Epoch 21/100
42/42 [=====] - 9s 207ms/step - loss: 0.0619 - accuracy: 0.9795 - val_loss: 0.0158 - val_accuracy: 1.0000
Epoch 22/100
42/42 [=====] - 9s 211ms/step - loss: 0.0648 - accuracy: 0.9772 - val_loss: 0.0381 - val_accuracy: 0.9789
Epoch 23/100
42/42 [=====] - 9s 209ms/step - loss: 0.0465 - accuracy: 0.9856 - val_loss: 0.0137 - val_accuracy: 1.0000
Epoch 24/100
42/42 [=====] - 9s 205ms/step - loss: 0.0593 - accuracy: 0.9795 - val_loss: 0.0129 - val_accuracy: 1.0000
Epoch 25/100
42/42 [=====] - 9s 208ms/step - loss: 0.0618 - accuracy: 0.9795 - val_loss: 0.0401 - val_accuracy: 0.9859
Epoch 26/100
42/42 [=====] - 8s 202ms/step - loss: 0.0475 - accuracy: 0.9856 - val_loss: 0.0135 - val_accuracy: 1.0000
Epoch 27/100
42/42 [=====] - 9s 210ms/step - loss: 0.0414 - accuracy: 0.9871 - val_loss: 0.0115 - val_accuracy: 1.0000
Epoch 28/100
39/42 [=====] - ETA: 0s - loss: 0.0540 - accuracy: 0.9808
```

```
In [24]: # Let's visualize the activations of our network
          from PIL import Image

          test_iter = iter(test_dataset)

          # Get a test image with a mask
          test_img, target = next(test_iter)
          test_img = test_img[0]

          # Visualize the image
          Image.fromarray(np.uint8(np.array(test_img)*255.))
```



→ we create a new model by giving the old model inputs and as outputs some layers of the old model. we're creating a new model but the weights are the one we trained before. But now the new model allows to get 2 outputs: the model output and the features map at the end of the encoder.

```
In [25]: np.argmax(target, 1)
```

```
In [26]: out_model = tf.keras.Model(inputs=model.input, outputs=[model.output,
```

output from the last convolutional layer  
(model.layers[0]) is the  
t\_at(-1)) encoder,  
get\_output\_at(-1)  
gets the last conv  
layer)

Now we want to localize the object in the original input space (and so we have to re-size) In [23] In [28]

```
In [27]: resize_feature = tf.keras.layers.experimental.preprocessing.Resizing(  
    256, 256, interpolation="bilinear")  
  
In [28]: softmax_out, last_enc_feature = out_model(tf.expand_dims(test_img, 0))  
pred = tf.argmax(softmax_out, 1) → we obtain the prediction and the encoder  
resized_feature = resize_feature(last_enc_feature) → here we extract the  
predicted class
```

```
In [29]: # Get the weights in input to the output neuron corresponding to the 'with_mask' class  
mask_weights = model.layers[-1].weights[0][:, 1]
```

we take the weights in the fully connected layer (the classifier)

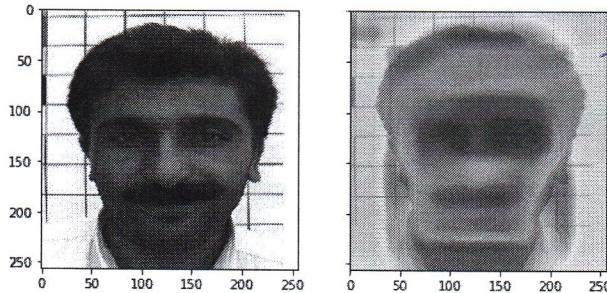
## CLASS ACTIVATION MAP

```
resized_feature = tf.reshape(resized_feature, shape=[256*256, resized_feature.shape[-1]])
cam = tf.linalg.matmul(resized_feature, tf.expand_dims(mask_weights, -1))
cam = tf.reshape(cam, shape=[256, 256])
```

we're transforming  
resized-features into the  
dimension of the matrix:  
 $256 \times 256 \times 128$

```
In [30]: import matplotlib.pyplot as plt
fig, ax = plt.subplots(nrows=1, ncols=2, sharey=True, sharex=True, figsize=plt.figaspect(0.5))
ax[0].imshow(test_img)
ax[1].imshow(test_img, alpha=0.5)
ax[1].imshow(cam, cmap='jet', alpha=0.5)
plt.show()

Out[30]: <matplotlib.image.AxesImage at 0x7fbc0279f7b8>
Out[30]: <matplotlib.image.AxesImage at 0x7fbc027296a0>
Out[30]: <matplotlib.image.AxesImage at 0x7fbc5007c160>
```



without training for object-detection or segmentation, this method is giving (by only inspecting activations in the features map) a way to localize the object we're interested in (more precisely: this method shows us what influences the most in the decision of the classification)

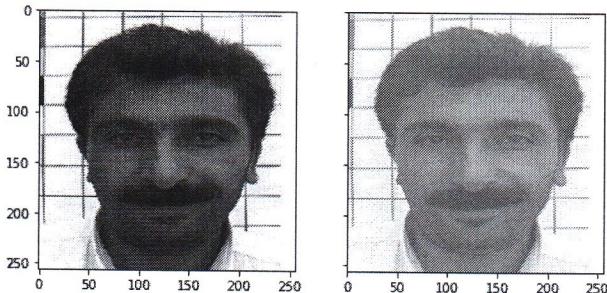
(However it's a very coarse localization)

```
In [33]: # we can obtain a coarse segmentation of the mask by thresholding
# For example we can use thr = 0.5 * max(cam)
```

— in this image we have no mask

```
fig, ax = plt.subplots(nrows=1, ncols=2, sharey=True, sharex=True, figsize=plt.figaspect(0.5))
ax[0].imshow(test_img)
ax[1].imshow(test_img, alpha=0.5)
ax[1].imshow(cam > 0.5 * np.max(cam), cmap='gray', alpha=0.5)
plt.show()
```

```
Out[33]: <matplotlib.image.AxesImage at 0x7fbc50072a58>
Out[33]: <matplotlib.image.AxesImage at 0x7fbff876d748>
```



## Topics: Image segmentation through Encoder-Decoder network (for binary classification) (we distinguish

- background
- foreground

```
In [1]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

```
In [2]: import os
import tensorflow as tf
import numpy as np

# Set the seed for random operations.
# This let our experiments to be reproducible.
SEED = 1234
tf.random.set_seed(SEED)
```

```
In [3]: cwd = os.getcwd()
```

```
In [4]: from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive
```

```
In [5]: !unzip /content/drive/My\ Drive/SegmentationDataset.zip
```

```
Archive: /content/drive/My\ Drive/SegmentationDataset.zip
creating: SegmentationDataset/
creating: SegmentationDataset/training/
creating: SegmentationDataset/training/images/
creating: SegmentationDataset/training/images/img/
inflating: SegmentationDataset/training/images/img/1408.tif
inflating: SegmentationDataset/training/images/img/6857.tif
inflating: ..
```

## Example: Image Segmentation

### Build segmentation

```
In [6]: # ImageDataGenerator
# -----
from tensorflow.keras.preprocessing.image import ImageDataGenerator
apply_data_augmentation = True

# Create training ImageDataGenerator object
# We need two different generators for images and corresponding masks
if apply_data_augmentation:
    train_img_data_gen = ImageDataGenerator(rotation_range=30,
                                             width_shift_range=10,
                                             height_shift_range=10,
                                             zoom_range=0.3,
                                             horizontal_flip=True,
                                             vertical_flip=True,
                                             fill_mode='reflect',
                                             rescale=1./255)
    train_mask_data_gen = ImageDataGenerator(rotation_range=30,
                                              width_shift_range=10,
                                              height_shift_range=10,
                                              zoom_range=0.3,
                                              horizontal_flip=True,
                                              vertical_flip=True,
                                              fill_mode='reflect')
else:
    train_img_data_gen = ImageDataGenerator(rescale=1./255)
    train_mask_data_gen = ImageDataGenerator()

# Create validation and test ImageDataGenerator objects
valid_img_data_gen = ImageDataGenerator(rescale=1./255)
valid_mask_data_gen = ImageDataGenerator()
test_img_data_gen = ImageDataGenerator(rescale=1./255)
test_mask_data_gen = ImageDataGenerator()
```

```
In [7]: !ls /content/SegmentationDataset/training/images
img
```

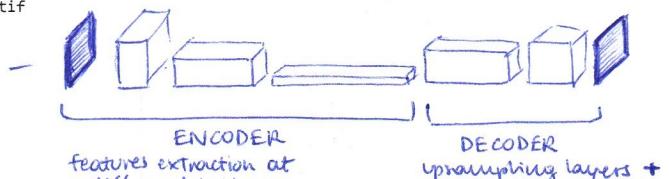
```
In [8]: # Create generators to read images from dataset directory
# -----
dataset_dir = os.path.join(cwd, 'SegmentationDataset')

# Batch size
bs = 16

# img shape
img_h = 256
img_w = 256
num_classes=3

# Training
# Two different generators for images and masks
# ATTENTION: here the seed is important!! We have to give the same SEED to both the generator
# to apply the same transformations/shuffling to images and corresponding masks
training_dir = os.path.join(dataset_dir, 'training')
train_img_gen = train_img_data_gen.flow_from_directory(os.path.join(training_dir, 'images'),
                                                       target_size=(img_h, img_w),
                                                       batch_size=bs,
                                                       class_mode=None,
                                                       shuffle=True,
                                                       interpolation='bilinear',
                                                       seed=SEED)

train_mask_gen = train_mask_data_gen.flow_from_directory(os.path.join(training_dir, 'masks'),
                                                       target_size=(img_h, img_w),
```



ENCODER  
features extraction at  
different scales  
(DOWNSAMPLING)

DECODER  
upsampling layers +  
convolutional layers  
(UPSAMPLING)

Winning approach: U-NET.  
The key point of this architecture is that  
encoder and decoder are connected  
at the corresponding levels.

} ImageDataGenerator for the image  
(RGB image)

} ImageDataGenerator for the mask  
(corresponding to the  
previous image)

Notice that the mask does not need to be  
normalized

```

        batch_size=bs,
        class_mode=None,
        shuffle=True,
        interpolation='bilinear',
        seed=SEED)

train_gen = zip(train_img_gen, train_mask_gen)

# Validation
validation_dir = os.path.join(dataset_dir, 'validation')
valid_img_gen = valid_img_data_gen.flow_from_directory(os.path.join(validation_dir, 'images'),
                                                       target_size=(img_h, img_w),
                                                       batch_size=bs,
                                                       class_mode=None,
                                                       shuffle=False,
                                                       interpolation='bilinear',
                                                       seed=SEED)
valid_mask_gen = valid_mask_data_gen.flow_from_directory(os.path.join(validation_dir, 'masks'),
                                                       target_size=(img_h, img_w),
                                                       batch_size=bs,
                                                       class_mode=None,
                                                       shuffle=False,
                                                       interpolation='bilinear',
                                                       seed=SEED)
valid_gen = zip(valid_img_gen, valid_mask_gen)

Found 500 images belonging to 1 classes.
Found 500 images belonging to 1 classes.
Found 100 images belonging to 1 classes.
Found 100 images belonging to 1 classes.

```

```

In [9]: # Create Dataset objects
# -----
# Training
# -----
train_dataset = tf.data.Dataset.from_generator(lambda: train_gen,
                                              output_types=(tf.float32, tf.float32),
                                              output_shapes=[[None, img_h, img_w, 3], [None, img_h, img_w, 3]])

# This is needed to convert targets from H x W x 3 (RGB) to Label mask H x W x 1.
# Thus, 255 values become True (with y_ > 0), i.e., RGB value [255, 255, 255] -> [True, True, True],
# and then we reduce it to a single value (reduce_any), i.e., RGB value [255, 255, 255] -> [True].
# Finally, cast boolean values to float to obtain 0-1 labels.
def prepare_target(x_, y_):
    return x_, tf.cast(tf.reduce_any(y_ > 0, axis=-1, keepdims=True), tf.float32)

train_dataset = train_dataset.map(prepare_target)
train_dataset = train_dataset.repeat()

# Validation
# -----
valid_dataset = tf.data.Dataset.from_generator(lambda: valid_gen,
                                              output_types=(tf.float32, tf.float32),
                                              output_shapes=[[None, img_h, img_w, 3], [None, img_h, img_w, 3]])
valid_dataset = valid_dataset.map(prepare_target)
valid_dataset = valid_dataset.repeat()

```

```
In [10]: iterator = iter(train_dataset)
```

```

In [11]: # Let's test data generator
# -----
import time
import matplotlib.pyplot as plt

%matplotlib inline

# Assign a color to each class
colors_dict = {}
colors_dict[1] = [255, 255, 255] # foreground
colors_dict[0] = [0, 0, 0] # background

iterator = iter(train_dataset)

fig, ax = plt.subplots(1, 2)

augmented_img, target = next(iterator)
augmented_img = augmented_img[0] # First element
augmented_img = augmented_img * 255 # denormalize

target = np.array(target[0, ..., 0]) # First element (squeezing channel dimension)

print(np.unique(target))

# Assign colors (just for visualization)
target_img = np.zeros([target.shape[0], target.shape[1], 3])

target_img[np.where(target == 0)] = colors_dict[0]
target_img[np.where(target == 1)] = colors_dict[1]

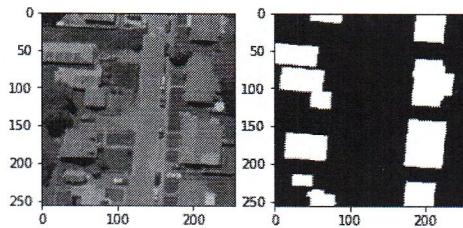
ax[0].imshow(np.uint8(augmented_img))
ax[1].imshow(np.uint8(target_img))

plt.show()

```

[0. 1.]

```
Out[11]: <matplotlib.image.AxesImage at 0x7fd69c1bb7b8>
Out[11]: <matplotlib.image.AxesImage at 0x7fd66724efd0>
```



## Convolutional Neural Network (CNN)

### Encoder-Decoder

In [12]:

```
# Create Model
# -----
def create_model(depth, start_f, num_classes, dynamic_input_shape):
    model = tf.keras.Sequential()

    # Encoder (Same that we used for classification)
    # -----
    for i in range(depth):

        if i == 0:
            if dynamic_input_shape:
                input_shape = [None, None, 3]
            else:
                input_shape = [img_h, img_w, 3]
        else:
            input_shape=[None]

        model.add(tf.keras.layers.Conv2D(filters=start_f,
                                         kernel_size=(3, 3),
                                         strides=(1, 1),
                                         padding='same',
                                         input_shape=input_shape))
        model.add(tf.keras.layers.ReLU())
        model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2)))

        start_f *= 2

    # Bottleneck
    model.add(tf.keras.layers.Conv2D(filters=start_f, kernel_size=(3, 3), strides=(1, 1), padding='same'))
    model.add(tf.keras.layers.ReLU())

    start_f = start_f // 2

    # Decoder
    # -----
    for i in range(depth):
        model.add(tf.keras.layers.UpSampling2D(2, interpolation='bilinear'))
        model.add(tf.keras.layers.Conv2D(filters=start_f,
                                         kernel_size=(3, 3),
                                         strides=(1, 1),
                                         padding='same'))
        model.add(tf.keras.layers.ReLU())

        start_f = start_f // 2

    # Prediction Layer (not a fully connected anymore)
    # -----
    model.add(tf.keras.layers.Conv2D(filters=1,
                                    kernel_size=(1, 1),
                                    strides=(1, 1),
                                    padding='same',
                                    activation='sigmoid'))

return model
```

} this is the  
encoding  
(from now we  
have to re-construct)

In [13]:

```
model = create_model(depth=5,
                     start_f=8,
                     num_classes=1,
                     dynamic_input_shape=False)

# Visualize created model as a table
model.summary()

# Visualize initialized weights
# model.weights
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 256, 256, 8)	224
re_lu (ReLU)	(None, 256, 256, 8)	0
max_pooling2d (MaxPooling2D)	(None, 128, 128, 8)	0
conv2d_1 (Conv2D)	(None, 128, 128, 16)	1168
re_lu_1 (ReLU)	(None, 128, 128, 16)	0
max_pooling2d_1 (MaxPooling2 (None, 64, 64, 16)		0
conv2d_2 (Conv2D)	(None, 64, 64, 32)	4640

re_lu_2 (ReLU)	(None, 64, 64, 32)	0
max_pooling2d_2 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_3 (Conv2D)	(None, 32, 32, 64)	18496
re_lu_3 (ReLU)	(None, 32, 32, 64)	0
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_4 (Conv2D)	(None, 16, 16, 128)	73856
re_lu_4 (ReLU)	(None, 16, 16, 128)	0
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 256)	295168
re_lu_5 (ReLU)	(None, 8, 8, 256)	0
up_sampling2d (UpSampling2D)	(None, 16, 16, 256)	0
conv2d_6 (Conv2D)	(None, 16, 16, 128)	295040
re_lu_6 (ReLU)	(None, 16, 16, 128)	0
up_sampling2d_1 (UpSampling2D)	(None, 32, 32, 128)	0
conv2d_7 (Conv2D)	(None, 32, 32, 64)	73792
re_lu_7 (ReLU)	(None, 32, 32, 64)	0
up_sampling2d_2 (UpSampling2D)	(None, 64, 64, 64)	0
conv2d_8 (Conv2D)	(None, 64, 64, 32)	18464
re_lu_8 (ReLU)	(None, 64, 64, 32)	0
up_sampling2d_3 (UpSampling2D)	(None, 128, 128, 32)	0
conv2d_9 (Conv2D)	(None, 128, 128, 16)	4624
re_lu_9 (ReLU)	(None, 128, 128, 16)	0
up_sampling2d_4 (UpSampling2D)	(None, 256, 256, 16)	0
conv2d_10 (Conv2D)	(None, 256, 256, 8)	1160
re_lu_10 (ReLU)	(None, 256, 256, 8)	0
conv2d_11 (Conv2D)	(None, 256, 256, 1)	9

## Prepare the model for training

In [14]:

```
# Optimization params
# ...
# Loss
# Binary Crossentropy
loss = tf.keras.losses.BinaryCrossentropy()
# learning rate
lr = 1e-3
optimizer = tf.keras.optimizers.Adam(learning_rate=lr)

# Validation metrics (the accuracy does not work very well with segmentation task) => INTERSECTION OF A UNION
# ...
score_th = 0.5
we consider the prob > score_th to be 1, otherwise 0

def IoU(y_true, y_pred):
    y_pred = tf.cast(y_pred > score_th, tf.float32)

    intersection = tf.reduce_sum(y_true * y_pred)
    union = tf.reduce_sum(y_true) + tf.reduce_sum(y_pred) - intersection

    return intersection / union

metrics = ['accuracy', IoU]

# Compile Model
model.compile(optimizer=optimizer, loss=loss, metrics=metrics)
```

we're upsampling till we reach the original input space (then we combine the convolutions to obtain only 1 channel at output)

for each pixel we'll have just 1 value and that value is the output of a sigmoid function, that means that for each pixel we have a probability (from 0 to 1) that that pixel is a background or foreground pixel

that's how we define this metric (IoU)

= percentage of the prediction overlapping the target (if the network correctly predict the class for the pixels (so, if we have highly performance) we expect the prediction to perfectly match the target, otherwise the overlapping will be small. We're NOT considering the background!)

## Training with callbacks

In [15]:

```
import os
from datetime import datetime

cwd = os.getcwd()

exp_dir = os.path.join(cwd, 'drive/My Drive/Keras4/binary_segmentation_experiments')
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

now = datetime.now().strftime('%b%d_%H-%M-%S')

model_name = 'CNN'

exp_dir = os.path.join(exp_dir, model_name + '_' + str(now))
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)
```

```

callbacks = []

# Model checkpoint
# -----
ckpt_dir = os.path.join(exp_dir, 'ckpts')
if not os.path.exists(ckpt_dir):
    os.makedirs(ckpt_dir)

ckpt_callback = tf.keras.callbacks.ModelCheckpoint(filepath=os.path.join(ckpt_dir, 'cp_{epoch:02d}.ckpt'),
                                                    save_weights_only=True) # False to save the model directly
callbacks.append(ckpt_callback)

# Visualize Learning on Tensorboard
# -----
tb_dir = os.path.join(exp_dir, 'tb_logs')
if not os.path.exists(tb_dir):
    os.makedirs(tb_dir)

# By default shows losses and metrics for both training and validation
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=tb_dir,
                                              profile_batch=0,
                                              histogram_freq=0) # if 1 shows weights histograms
callbacks.append(tb_callback)

# Early Stopping
# -----
early_stop = False
if early_stop:
    es_callback = tf.keras.callback.EarlyStopping(monitor='val_loss', patience=10)
    callbacks.append(es_callback)

model.fit(x=train_dataset,
          epochs=100, ##### set repeat in training dataset
          steps_per_epoch=len(train_img_gen),
          validation_data=valid_dataset,
          validation_steps=len(valid_img_gen),
          callbacks=callbacks)

# How to visualize Tensorboard
# 1. tensorboard --logdir EXPERIMENTS_DIR --port PORT      <- from terminal
# 2. Localhost:PORT   <- in your browser

Epoch 1/100
32/32 [=====] - 18s 575ms/step - loss: 0.6493 - accuracy: 0.6722 - IoU: 0.0132 - val_loss: 0.62
91 - val_accuracy: 0.6962 - val_IoU: 0.0000e+00
Epoch 2/100
32/32 [=====] - 17s 538ms/step - loss: 0.6462 - accuracy: 0.6749 - IoU: 0.0000e+00 - val_loss: 0.6157
- val_accuracy: 0.6962 - val_IoU: 0.0000e+00
Epoch 3/100
32/32 [=====] - 17s 537ms/step - loss: 0.6296 - accuracy: 0.6791 - IoU: 0.0000e+00 - val_loss: 0.6114
- val_accuracy: 0.6962 - val_IoU: 0.0000e+00
Epoch 4/100
32/32 [=====] - 18s 553ms/step - loss: 0.6078 - accuracy: 0.6784 - IoU: 0.0000e+00 - val_loss: 0.6163
- val_accuracy: 0.6962 - val_IoU: 0.0000e+00
Epoch 5/100
32/32 [=====] - 18s 564ms/step - loss: 0.6189 - accuracy: 0.6768 - IoU: 0.0000e+00 - val_loss: 0.5837
- val_accuracy: 0.6962 - val_IoU: 0.0000e+00
Epoch 6/100
32/32 [=====] - 18s 557ms/step - loss: 0.5919 - accuracy: 0.6768 - IoU: 0.0000e+00 - val_loss: 0.5876
- val_accuracy: 0.7057 - val_IoU: 0.0567
Epoch 7/100
32/32 [=====] - 18s 551ms/step - loss: 0.5845 - accuracy: 0.7003 - IoU: 0.1890 - val_loss: 0.57
33 - val_accuracy: 0.7077 - val_IoU: 0.3031
Epoch 8/100
32/32 [=====] - 18s 553ms/step - loss: 0.5691 - accuracy: 0.7117 - IoU: 0.2092 - val_loss: 0.55
39 - val_accuracy: 0.7250 - val_IoU: 0.3469
Epoch 9/100
32/32 [=====] - 18s 550ms/step - loss: 0.5495 - accuracy: 0.7267 - IoU: 0.2880 - val_loss: 0.58
81 - val_accuracy: 0.6780 - val_IoU: 0.3805
Epoch 10/100
32/32 [=====] - 18s 551ms/step - loss: 0.5453 - accuracy: 0.7338 - IoU: 0.3065 - val_loss: 0.53
02 - val_accuracy: 0.7482 - val_IoU: 0.3292
Epoch 11/100
32/32 [=====] - 18s 556ms/step - loss: 0.5269 - accuracy: 0.7446 - IoU: 0.3412 - val_loss: 0.54
95 - val_accuracy: 0.7232 - val_IoU: 0.3932
Epoch 12/100
32/32 [=====] - 18s 553ms/step - loss: 0.5383 - accuracy: 0.7288 - IoU: 0.2484 - val_loss: 0.53
22 - val_accuracy: 0.7179 - val_IoU: 0.4000
Epoch 13/100
32/32 [=====] - 18s 550ms/step - loss: 0.5091 - accuracy: 0.7575 - IoU: 0.3833 - val_loss: 0.50
25 - val_accuracy: 0.7549 - val_IoU: 0.3032
Epoch 14/100
32/32 [=====] - 18s 553ms/step - loss: 0.5161 - accuracy: 0.7514 - IoU: 0.3623 - val_loss: 0.51
64 - val_accuracy: 0.7518 - val_IoU: 0.4022
Epoch 15/100
32/32 [=====] - 18s 551ms/step - loss: 0.5026 - accuracy: 0.7586 - IoU: 0.3967 - val_loss: 0.54
36 - val_accuracy: 0.7136 - val_IoU: 0.4075
Epoch 16/100
32/32 [=====] - 18s 554ms/step - loss: 0.4965 - accuracy: 0.7625 - IoU: 0.3771 - val_loss: 0.50
90 - val_accuracy: 0.7626 - val_IoU: 0.3087
Epoch 17/100
32/32 [=====] - 18s 557ms/step - loss: 0.4912 - accuracy: 0.7658 - IoU: 0.4015 - val_loss: 0.49
27 - val_accuracy: 0.7636 - val_IoU: 0.3979
Epoch 18/100
32/32 [=====] - 18s 554ms/step - loss: 0.4831 - accuracy: 0.7758 - IoU: 0.4150 - val_loss: 0.48
26 - val_accuracy: 0.7680 - val_IoU: 0.4119
Epoch 19/100
32/32 [=====] - 18s 552ms/step - loss: 0.4703 - accuracy: 0.7840 - IoU: 0.4494 - val_loss: 0.56
57 - val_accuracy: 0.7247 - val_IoU: 0.4539
Epoch 20/100
32/32 [=====] - 18s 550ms/step - loss: 0.4731 - accuracy: 0.7796 - IoU: 0.4329 - val_loss: 0.47
23 - val_accuracy: 0.7766 - val_IoU: 0.4145
Epoch 21/100
32/32 [=====] - 18s 551ms/step - loss: 0.4658 - accuracy: 0.7810 - IoU: 0.4424 - val_loss: 0.49

```

```

63 - val_accuracy: 0.7629 - val_IoU: 0.4355
Epoch 22/100
32/32 [=====] - 18s 557ms/step - loss: 0.4640 - accuracy: 0.7838 - IoU: 0.4542 - val_loss: 0.49
01 - val_accuracy: 0.7578 - val_IoU: 0.4756
Epoch 23/100
32/32 [=====] - 18s 549ms/step - loss: 0.4711 - accuracy: 0.7848 - IoU: 0.4606 - val_loss: 0.57
82 - val_accuracy: 0.7267 - val_IoU: 0.4760
Epoch 24/100
32/32 [=====] - 18s 550ms/step - loss: 0.4540 - accuracy: 0.7920 - IoU: 0.4751 - val_loss: 0.51
97 - val_accuracy: 0.7532 - val_IoU: 0.4708
Epoch 25/100
32/32 [=====] - 18s 552ms/step - loss: 0.4569 - accuracy: 0.7898 - IoU: 0.4671 - val_loss: 0.46
34 - val_accuracy: 0.7873 - val_IoU: 0.4022
Epoch 26/100
32/32 [=====] - 17s 542ms/step - loss: 0.4409 - accuracy: 0.7979 - IoU: 0.4733 - val_loss: 0.44
50 - val_accuracy: 0.7991 - val_IoU: 0.4496
Epoch 27/100
32/32 [=====] - 18s 550ms/step - loss: 0.4318 - accuracy: 0.8047 - IoU: 0.4989 - val_loss: 0.45
08 - val_accuracy: 0.7983 - val_IoU: 0.4669
Epoch 28/100
32/32 [=====] - 18s 553ms/step - loss: 0.4316 - accuracy: 0.8036 - IoU: 0.4900 - val_loss: 0.45
66 - val_accuracy: 0.7879 - val_IoU: 0.5013
Epoch 29/100
32/32 [=====] - 18s 553ms/step - loss: 0.4194 - accuracy: 0.8138 - IoU: 0.5125 - val_loss: 0.46
55 - val_accuracy: 0.7957 - val_IoU: 0.4531
Epoch 30/100
32/32 [=====] - 18s 550ms/step - loss: 0.4260 - accuracy: 0.8070 - IoU: 0.5094 - val_loss: 0.48
47 - val_accuracy: 0.7662 - val_IoU: 0.5056
Epoch 31/100
32/32 [=====] - 18s 550ms/step - loss: 0.4255 - accuracy: 0.8093 - IoU: 0.5077 - val_loss: 0.48
63 - val_accuracy: 0.7733 - val_IoU: 0.4485
Epoch 32/100
32/32 [=====] - 18s 555ms/step - loss: 0.4178 - accuracy: 0.8120 - IoU: 0.5142 - val_loss: 0.44
66 - val_accuracy: 0.7923 - val_IoU: 0.4975
Epoch 33/100
32/32 [=====] - 18s 552ms/step - loss: 0.4314 - accuracy: 0.8046 - IoU: 0.4935 - val_loss: 0.45
22 - val_accuracy: 0.7928 - val_IoU: 0.4816
Epoch 34/100
32/32 [=====] - 18s 553ms/step - loss: 0.4147 - accuracy: 0.8160 - IoU: 0.5195 - val_loss: 0.43
87 - val_accuracy: 0.8030 - val_IoU: 0.5137
Epoch 35/100
32/32 [=====] - 18s 551ms/step - loss: 0.4091 - accuracy: 0.8180 - IoU: 0.5305 - val_loss: 0.43
69 - val_accuracy: 0.7968 - val_IoU: 0.5182
Epoch 36/100
32/32 [=====] - 18s 556ms/step - loss: 0.3950 - accuracy: 0.8231 - IoU: 0.5409 - val_loss: 0.44
69 - val_accuracy: 0.7938 - val_IoU: 0.5135
Epoch 37/100
32/32 [=====] - 18s 558ms/step - loss: 0.3963 - accuracy: 0.8231 - IoU: 0.5354 - val_loss: 0.44
91 - val_accuracy: 0.7975 - val_IoU: 0.5047
Epoch 38/100
32/32 [=====] - 17s 546ms/step - loss: 0.3979 - accuracy: 0.8222 - IoU: 0.5447 - val_loss: 0.42
38 - val_accuracy: 0.8086 - val_IoU: 0.4832
Epoch 39/100
32/32 [=====] - 18s 550ms/step - loss: 0.3951 - accuracy: 0.8265 - IoU: 0.5442 - val_loss: 0.44
20 - val_accuracy: 0.8017 - val_IoU: 0.5279
Epoch 40/100
32/32 [=====] - 17s 546ms/step - loss: 0.3820 - accuracy: 0.8329 - IoU: 0.5574 - val_loss: 0.41
09 - val_accuracy: 0.8171 - val_IoU: 0.5119
Epoch 41/100
32/32 [=====] - 17s 545ms/step - loss: 0.4093 - accuracy: 0.8161 - IoU: 0.5255 - val_loss: 0.44
09 - val_accuracy: 0.7998 - val_IoU: 0.5130
Epoch 42/100
32/32 [=====] - 17s 543ms/step - loss: 0.3932 - accuracy: 0.8240 - IoU: 0.5311 - val_loss: 0.41
83 - val_accuracy: 0.8101 - val_IoU: 0.5083
Epoch 43/100
32/32 [=====] - 17s 537ms/step - loss: 0.3803 - accuracy: 0.8316 - IoU: 0.5633 - val_loss: 0.42
55 - val_accuracy: 0.8189 - val_IoU: 0.5117
Epoch 44/100
29/32 [=====] - ETA: 1s - loss: 0.3721 - accuracy: 0.8362 - IoU: 0.5690

```

## Test model

### Compute prediction

```

In [16]: iterator = iter(valid_dataset)

In [17]: fig, ax = plt.subplots(1, 3, figsize=(8, 8))
fig.show()
image, target = next(iterator)

image = image[0]
target = target[0, ..., 0]

out_sigmoid = model.predict(x=tf.expand_dims(image, 0))

# Get predicted class as the index corresponding to the maximum value in the vector probability
predicted_class = tf.cast(out_sigmoid > score_th, tf.int32)
predicted_class = predicted_class[0, ..., 0]

print(predicted_class.shape)

# Assign colors (just for visualization)
target_img = np.zeros([target.shape[0], target.shape[1], 3])
prediction_img = np.zeros([target.shape[0], target.shape[1], 3])

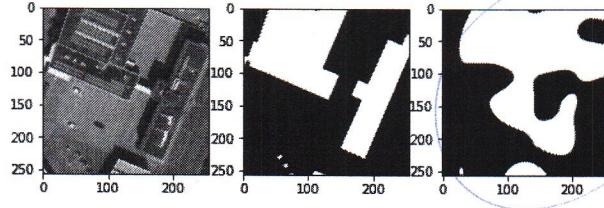
target_img[np.where(target == 0)] = colors_dict[0]
target_img[np.where(target == 1)] = colors_dict[1]

prediction_img[np.where(predicted_class == 0)] = colors_dict[0]
prediction_img[np.where(predicted_class == 1)] = colors_dict[1]

ax[0].imshow(np.uint8(image*255.))
ax[1].imshow(np.uint8(target_img))

```

```
    ax[2].imshow(np.uint8(prediction_img))
    fig.canvas.draw()
    time.sleep(1)
(256, 256)
Out[17]: <matplotlib.image.AxesImage at 0x7fd60b97a9e8>
Out[17]: <matplotlib.image.AxesImage at 0x7fd60b92a2b0>
Out[17]: <matplotlib.image.AxesImage at 0x7fd60b92a390>
```



result obtained on a test image  
(notice that we stopped the training  
when the model was not yet so good)

## Topics : Multi-class Segmentation

```
In [1]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

In [2]: import os

# os.environ["CUDA_VISIBLE_DEVICES"]="-1"
import tensorflow as tf
import numpy as np

# Set the seed for random operations.
# This let our experiments to be reproducible.
SEED = 1234
tf.random.set_seed(SEED)

In [3]: cwd = os.getcwd()

In [4]: from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

In [5]: !unzip /content/drive/My\ Drive/VOCDataset.zip
Archive: /content/drive/My\ Drive/VOCDataset.zip
  creating: VOCDataset/Images/
  creating: VOCDataset/Images/
  inflating: VOCDataset/Images/000032.jpg
  inflating: ...

In [6]: !ls /content/VOCDataset/Images
000032.jpg 001724.jpg 003373.jpg 005144.jpg 006699.jpg 008403.jpg
000033.jpg 001733.jpg ..
```

## Example: Image Segmentation

### Build segmentation

```
In [7]: # ImageDataGenerator
# ----

from tensorflow.keras.preprocessing.image import ImageDataGenerator
apply_data_augmentation = True

# Create training ImageDataGenerator object
# We need two different generators for images and corresponding masks
if apply_data_augmentation:
    img_data_gen = ImageDataGenerator(rotation_range=10,
                                      width_shift_range=10,
                                      height_shift_range=10,
                                      zoom_range=0.3,
                                      horizontal_flip=True,
                                      vertical_flip=True,
                                      fill_mode='reflect')
    mask_data_gen = ImageDataGenerator(rotation_range=10,
                                      width_shift_range=10,
                                      height_shift_range=10,
                                      zoom_range=0.3,
                                      horizontal_flip=True,
                                      vertical_flip=True,
                                      fill_mode='reflect')

In [8]: from PIL import Image
class CustomDataset(tf.keras.utils.Sequence):

    """
    CustomDataset inheriting from tf.keras.utils.Sequence.

    3 main methods:
    - __init__: save dataset params like directory, filenames..
    - __len__: return the total number of samples in the dataset
    - __getitem__: return a sample from the dataset

    Note:
    - the custom dataset return a single sample from the dataset. Then, we use
      a tf.data.Dataset object to group samples into batches.
    - in this case we have a different structure of the dataset in memory.
      We have all the images in the same folder and the training and validation splits
      are defined in text files.

    """

    def __init__(self, dataset_dir, which_subset, img_generator=None, mask_generator=None,
                 preprocessing_function=None, out_shape=[256, 256]):
        if which_subset == 'training':
            subset_file = os.path.join(dataset_dir, 'Splits', 'train.txt')
        elif which_subset == 'validation':
            subset_file = os.path.join(dataset_dir, 'Splits', 'val.txt')

        with open(subset_file, 'r') as f:
            lines = f.readlines()

        subset_filenames = []
        for line in lines:
            subset_filenames.append(line.strip())

        self.which_subset = which_subset
```

```

    self.dataset_dir = dataset_dir
    self.subset_filenames = subset_filenames
    self.img_generator = img_generator
    self.mask_generator = mask_generator
    self.preprocessing_function = preprocessing_function
    self.out_shape = out_shape

    def __len__(self):
        return len(self.subset_filenames)

    def __getitem__(self, index):
        # Read Image
        curr_filename = self.subset_filenames[index]
        img = Image.open(os.path.join(self.dataset_dir, 'Images', curr_filename + '.jpg'))
        mask = Image.open(os.path.join(self.dataset_dir, 'Annotations', curr_filename + '.png'))

        # Resize image and mask
        img = img.resize(self.out_shape)
        mask = mask.resize(self.out_shape, resample=Image.NEAREST)

        img_arr = np.array(img)
        mask_arr = np.array(mask)

        # in this dataset 255 mask label is assigned to an additional class, which corresponds
        # to the contours of the objects. We remove it for simplicity.
        mask_arr[mask_arr == 255] = 0

        mask_arr = np.expand_dims(mask_arr, -1)

        if self.which_subset == 'training':
            if self.img_generator is not None and self.mask_generator is not None:
                # Perform data augmentation
                # We can get a random transformation from the ImageDataGenerator using get_random_transform
                # and we can apply it to the image using apply_transform
                img_t = self.img_generator.get_random_transform(img_arr.shape, seed=SEED)
                mask_t = self.mask_generator.get_random_transform(mask_arr.shape, seed=SEED)
                img_arr = self.img_generator.apply_transform(img_arr, img_t)
                # ImageDataGenerator use bilinear interpolation for augmenting the images.
                # Thus, when applied to the masks it will output 'interpolated classes', which
                # is an unwanted behaviour. As a trick, we can transform each class mask
                # separately and then we can cast to integer values (as in the binary segmentation notebook).
                # Finally, we merge the augmented binary masks to obtain the final segmentation mask.
                out_mask = np.zeros_like(mask_arr)
                for c in np.unique(mask_arr):
                    if c > 0:
                        curr_class_arr = np.float32(mask_arr == c)
                        curr_class_arr = self.mask_generator.apply_transform(curr_class_arr, mask_t)
                        # from [0, 1] to {0, 1}
                        curr_class_arr = np.uint8(curr_class_arr)
                        # recover original class
                        curr_class_arr = curr_class_arr * c
                        out_mask += curr_class_arr
                    else:
                        out_mask = mask_arr
            if self.preprocessing_function is not None:
                img_arr = self.preprocessing_function(img_arr)

        return img_arr, np.float32(out_mask)

```

In [9]:

```

from tensorflow.keras.applications.vgg16 import preprocess_input

img_h = 256
img_w = 256

dataset = CustomDataset('/content/VOCdataset', 'training',
                       img_generator=img_data_gen, mask_generator=mask_data_gen,
                       preprocessing_function=preprocess_input)
dataset_valid = CustomDataset('/content/VOCdataset', 'validation',
                           preprocessing_function=preprocess_input)

```

In [10]:

```

train_dataset = tf.data.Dataset.from_generator(lambda: dataset,
                                              output_types=(tf.float32, tf.float32),
                                              output_shapes=([img_h, img_w, 3], [img_h, img_w, 1]))

train_dataset = train_dataset.batch(32)
train_dataset = train_dataset.repeat()

valid_dataset = tf.data.Dataset.from_generator(lambda: dataset_valid,
                                              output_types=(tf.float32, tf.float32),
                                              output_shapes=([img_h, img_w, 3], [img_h, img_w, 1]))

valid_dataset = valid_dataset.batch(32)
valid_dataset = valid_dataset.repeat()

```

In [11]:

```

!ls /content/VOCdataset
Annotations Images Splits

```

In [12]:

```

# Let's test data generator
# -----
import time
from matplotlib import cm
import matplotlib.pyplot as plt

%matplotlib inline

# Assign a color to each class
evenly_spaced_interval = np.linspace(0, 1, 20)
colors = [cm.rainbow(x) for x in evenly_spaced_interval]

iterator = iter(valid_dataset)

```

```
In [13]: fig, ax = plt.subplots(1, 2)

augmented_img, target = next(iterator)
augmented_img = augmented_img[0] # First element
augmented_img = augmented_img # denormalize

target = np.array(target[0, ..., 0]) # First element (squeezing channel dimension)
print(np.unique(target))

target_img = np.zeros([target.shape[0], target.shape[1], 3])

target_img[np.where(target == 0)] = [0, 0, 0]
for i in range(1, 21):
    target_img[np.where(target == i)] = np.array(colors[i-1])[:3] * 255

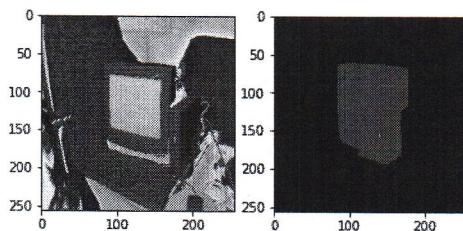
ax[0].imshow(np.uint8(augmented_img))
ax[1].imshow(np.uint8(target_img))

plt.show()
```

[ 0. 20.]

Out[13]: <matplotlib.image.AxesImage at 0x7f8e102dc828>

Out[13]: <matplotlib.image.AxesImage at 0x7f8e103795c0>



```
In [14]: vgg = tf.keras.applications.VGG16(weights='imagenet', include_top=False, input_shape=(img_h, img_w, 3))
vgg.summary()
for layer in vgg.layers:
    layer.trainable = False
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)
58892288/58889256 [=====] - 4s 0us/step  
Model: "vgg16"

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
<hr/>		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

## Convolutional Neural Network (CNN)

### Encoder-Decoder

```
In [15]: def create_model(depth, start_f, num_classes):

    model = tf.keras.Sequential()

    # Encoder
```

```

# -----
model.add(vgg)      ← we take the encoder of the VGG and we have to add
                    only the decoder part
start_f = 256

# Decoder
# -----
for i in range(depth):
    model.add(tf.keras.layers.UpSampling2D(2, interpolation='bilinear'))
    model.add(tf.keras.layers.Conv2D(filters=start_f,
                                    kernel_size=(3, 3),
                                    strides=(1, 1),
                                    padding='same'))
    model.add(tf.keras.layers.ReLU())

    start_f = start_f // 2

# Prediction Layer
# -----
model.add(tf.keras.layers.Conv2D(filters=num_classes,
                                kernel_size=(1, 1),
                                strides=(1, 1),
                                padding='same',
                                activation='softmax'))

return model

```

In [16]:

```

model = create_model(depth=5,
                     start_f=8,
                     num_classes=21)

# Visualize created model as a table
model.summary()

# Visualize initialized weights
# model.weights

```

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 8, 8, 512)	14714688
up_sampling2d (UpSampling2D)	(None, 16, 16, 512)	0
conv2d (Conv2D)	(None, 16, 16, 256)	1179904
re_lu (ReLU)	(None, 16, 16, 256)	0
up_sampling2d_1 (UpSampling2D)	(None, 32, 32, 256)	0
conv2d_1 (Conv2D)	(None, 32, 32, 128)	295040
re_lu_1 (ReLU)	(None, 32, 32, 128)	0
up_sampling2d_2 (UpSampling2D)	(None, 64, 64, 128)	0
conv2d_2 (Conv2D)	(None, 64, 64, 64)	73792
re_lu_2 (ReLU)	(None, 64, 64, 64)	0
up_sampling2d_3 (UpSampling2D)	(None, 128, 128, 64)	0
conv2d_3 (Conv2D)	(None, 128, 128, 32)	18464
re_lu_3 (ReLU)	(None, 128, 128, 32)	0
up_sampling2d_4 (UpSampling2D)	(None, 256, 256, 32)	0
conv2d_4 (Conv2D)	(None, 256, 256, 16)	4624
re_lu_4 (ReLU)	(None, 256, 256, 16)	0
conv2d_5 (Conv2D)	(None, 256, 256, 21)	357

Total params: 16,286,869  
Trainable params: 1,572,181  
Non-trainable params: 14,714,688

## Prepare the model for training

In [17]:

```

# Optimization params
# ----

# LOSS
# Sparse Categorical Crossentropy to use integers (mask) instead of one-hot encoded labels
loss = tf.keras.losses.SparseCategoricalCrossentropy()
# Learning rate
lr = 1e-3
optimizer = tf.keras.optimizers.Adam(learning_rate=lr)

# Here we define the intersection over union for each class in the batch.
# Then we compute the final iou as the mean over classes
def meanIoU(y_true, y_pred):
    # get predicted class from softmax
    y_pred = tf.expand_dims(tf.argmax(y_pred, -1), -1)

    per_class_iou = []

    for i in range(1,21): # exclude the background class 0
        # Get prediction and target related to only a single class (i)
        class_pred = tf.cast(tf.where(y_pred == i, 1, 0), tf.float32)

```

```

    class_true = tf.cast(tf.where(y_true == i, 1, 0), tf.float32)
    intersection = tf.reduce_sum(class_true * class_pred)
    union = tf.reduce_sum(class_true) + tf.reduce_sum(class_pred) - intersection
    iou = (intersection + 1e-7) / (union + 1e-7)
    per_class_iou.append(iou)

    return tf.reduce_mean(per_class_iou)

# Validation metrics
# -----
metrics = ['accuracy', meanIoU]

# Compile Model
model.compile(optimizer=optimizer, loss=loss, metrics=metrics)

```

## Training with callbacks

```

In [18]: import os
from datetime import datetime

cwd = os.getcwd()

exp_dir = os.path.join(cwd, 'drive/My Drive/Keras4/', 'multiclass_segmentation_experiments')
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

now = datetime.now().strftime('%b%d_%H-%M-%S')

model_name = 'CNN'

exp_dir = os.path.join(exp_dir, model_name + '_' + str(now))
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

callbacks = []

# Model checkpoint
# -----
ckpt_dir = os.path.join(exp_dir, 'ckpts')
if not os.path.exists(ckpt_dir):
    os.makedirs(ckpt_dir)

ckpt_callback = tf.keras.callbacks.ModelCheckpoint(filepath=os.path.join(ckpt_dir, 'cp_{epoch:02d}.ckpt'),
                                                    save_weights_only=True) # False to save the model directly
callbacks.append(ckpt_callback)

# Visualize Learning on Tensorboard
# -----
tb_dir = os.path.join(exp_dir, 'tb_logs')
if not os.path.exists(tb_dir):
    os.makedirs(tb_dir)

# By default shows losses and metrics for both training and validation
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=tb_dir,
                                              profile_batch=0,
                                              histogram_freq=0) # if 1 shows weights histograms
callbacks.append(tb_callback)

# Early Stopping
# -----
early_stop = False
if early_stop:
    es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=10)
    callbacks.append(es_callback)

model.fit(x=train_dataset,
          epochs=100, ##### set repeat in training dataset
          steps_per_epoch=len(dataset),
          validation_data=valid_dataset,
          validation_steps=len(dataset_valid),
          callbacks=callbacks)

# How to visualize Tensorboard
# 1. tensorboard --logdir EXPERIMENTS_DIR --port PORT      <- from terminal
# 2. localhost:PORT <- in your browser

Epoch 1/100
209/209 [=====] - 497s 2s/step - loss: 0.8915 - accuracy: 0.7971 - meanIoU: 0.1146 - val_loss: 1.7129 - val_accuracy: 0.7826 - val_meanIoU: 0.1457
Epoch 2/100
5/209 [.....] - ETA: 4:47 - loss: 0.4179 - accuracy: 0.8586 - meanIoU: 0.2225

```

## Test model

### Compute prediction

```

In [ ]: model.load_weights('/content/drive/My Drive/Keras4/multiclass_segmentation_experiments/CNN_Nov27_08-41-36/ckpts/cp_02.ckpt')

In [20]: import time
import matplotlib.pyplot as plt

from PIL import Image
%matplotlib inline

iterator = iter(valid_dataset)

```

```
In [21]: fig, ax = plt.subplots(1, 3, figsize=(8, 8))
fig.show()
image, target = next(iterator)

image = image[0]
target = target[0, ..., 0]

out_sigmoid = model.predict(x=tf.expand_dims(image, 0))

# Get predicted class as the index corresponding to the maximum value in the vector probability
# predicted_class = tf.cast(out_sigmoid > score_th, tf.int32)
# predicted_class = predicted_class[0, ..., 0]
predicted_class = tf.argmax(out_sigmoid, -1)

out_sigmoid.shape

predicted_class = predicted_class[0, ...]

# Assign colors (just for visualization)
target_img = np.zeros([target.shape[0], target.shape[1], 3])
prediction_img = np.zeros([target.shape[0], target.shape[1], 3])

target_img[np.where(target == 0)] = [0, 0, 0]
for i in range(1, 21):
    target_img[np.where(target == i)] = np.array(colors[i-1])[:3] * 255

prediction_img[np.where(predicted_class == 0)] = [0, 0, 0]
for i in range(1, 21):
    prediction_img[np.where(predicted_class == i)] = np.array(colors[i-1])[:3] * 255

ax[0].imshow(np.uint8(image))
ax[1].imshow(np.uint8(target_img))
ax[2].imshow(np.uint8(prediction_img))

fig.canvas.draw()
time.sleep(1)
```

Out[21]: (1, 256, 256, 21)

Out[21]: <matplotlib.image.AxesImage at 0x7f8db8b8e358>

Out[21]: <matplotlib.image.AxesImage at 0x7f8db8aa5eb8>

Out[21]: <matplotlib.image.AxesImage at 0x7f8db8ac1cf8>

