

1/2
Topics

```
In [1]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

```
In [2]: import os
import tensorflow as tf
import numpy as np

# Set the seed for random operations.
# This let our experiments to be reproducible.
SEED = 1234
tf.random.set_seed(SEED)
np.random.seed(SEED)

# Get current working directory
cwd = os.getcwd()
```

```
In [3]: from google.colab import drive
drive.mount('/content/drive/')

Mounted at /content/drive/
```

```
In [4]: !cat /content/drive/My\ Drive/ita.txt
```

```
Output streaming troncato alle ultime 5000 righe.
Please smile. Sorridi, per favore. CC-BY 2.0 (France) Attribution: tatoeba.org #1852296 (CK) & #4140430 (Guybrush88)
Please smile. Sorridi, per piacere. CC-BY 2.0 (France) Attribution: tatoeba.org #1852296 (CK) & #4140431 (Guybrush88)
Please smile. Sorrida, per favore. CC-BY 2.0 (France) Attribution: tatoeba.org #1852296 (CK) & #4140432 (Guybrush88)
Please smile. Sorrida, per piacere. CC-BY 2.0 (France) Attribution: tatoeba.org #1852296 (CK) & #4140433 (Guybrush88)
Please smile. Sorridete, per favore. CC-BY 2.0 (France) Attribution: tatoeba.org #1852296 (CK) & #4140434 (Guybrush88)
Please smile. Sorridete, per piacere. CC-BY 2.0 (France) Attribution: tatoeba.org #1852296 (CK) & #4140435 (Guybrush88)
Put it there. Mettilo lì. CC-BY 2.0 (France) Attribution: tatoeba.org #1553401 (CK) & #5209591 (Guybrush88)
Put it there. Mettila lì. CC-BY 2.0 (France) Attribution: tatoeba.org #1553401 (CK) & #5209592 (Guybrush88)
Put it there. Lo metta lì. CC-BY 2.0 (France) Attribution: tatoeba.org #1553401 (CK) & #5209593 (Guybrush88)
Put it there. La metta lì. CC-BY 2.0 (France) Attribution: tatoeba.org #1553401 (CK) & #5209594 (Guybrush88)
Put it there. Mettetelo lì. CC-BY 2.0 (France) Attribution: tatoeba.org #1553401 (CK) & #5209596 (Guybrush88)
Put it there. Mettetela lì. CC-BY 2.0 (France) Attribution: tatoeba.org #1553401 (CK) & #5209597 (Guybrush88)
```

Neural Machine Translation

Italian/English

Dataset

We have to encode [5]: the words; we can't feed the model with bare words. We could simply do one-hot-encoding but with that every word would be 1 to all the others, while we would want to preserve some closeness in case of words-correlation.

We read line per line and we split eng-ita by the "TAB"

```
# Prepare dataset
# ----

MAX_NUM_SENTENCES = 40000
MAX_NUM_WORDS = 20000

ita_sentences = [] ← encoder input
eng_sentences = [] ← target sentences
eng_sentences_train = [] ← input of the decoder during training

# Simplify the dataset
MAX_LEN = 3 # words
# we take sentences with max 3 words
# Read all lines in translation dataset
count = 0
for line in open(os.path.join('/content/drive/My Drive', 'ita.txt'), encoding='utf-8'):

    if count > MAX_NUM_SENTENCES:
        break

    if '\t' not in line:
        continue

    eng_sentence_, ita_sentence, _ = line.rstrip().split('\t')
    [there are some lines without the TAB; if we encounter them we skip them with "continue"]

    if (len(eng_sentence_.split(' ')) > MAX_LEN or
        len(ita_sentence.split(' ')) > MAX_LEN):
        continue (= go to the next one)

    eng_sentence = eng_sentence_ + '<eos>'
    eng_sentence_train = '<sos>' + eng_sentence_
    [encoding for the END OF SENTENCE]

    ita_sentences.append(ita_sentence)
    eng_sentences.append(eng_sentence)
    eng_sentences_train.append(eng_sentence_train)
    [encoding for the START OF SENTENCE]

    count += 1

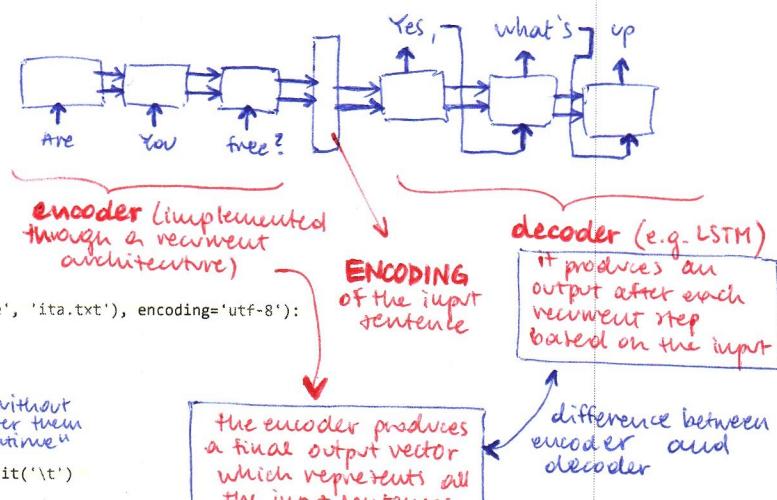
print('Number of sentences:', len(ita_sentences))
Number of sentences: 32499
```

```
In [6]: max(len(sentence.split(' ')) for sentence in eng_sentences)
```

```
Out[6]: 4
```

sentence as input and sentence as output
(a Seq2Seq can work)

the architecture is an encoder-decoder one



Tokenization — The tokenizer allows to obtain automatically the vocabulary of the words given a text

Converts words to integers

```
In [7]:
```

```
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

# Create Tokenizer to convert words to integers
ita_tokenizer = Tokenizer(num_words=MAX_NUM_WORDS)
ita_tokenizer.fit_on_texts(ita_sentences)
ita_tokenized = ita_tokenizer.texts_to_sequences(ita_sentences)

ita_wtoi = ita_tokenizer.word_index
print('Total italian words:', len(ita_wtoi))

max_ita_length = max(len(sentence) for sentence in ita_tokenized)
print('Max italian sentence length:', max_ita_length)

eng_tokenizer = Tokenizer(num_words=MAX_NUM_WORDS, filters='?!.,')
eng_tokenizer.fit_on_texts(eng_sentences+eng_sentences_train)
eng_tokenized = eng_tokenizer.texts_to_sequences(eng_sentences)
eng_tokenized_train = eng_tokenizer.texts_to_sequences(eng_sentences_train)

eng_wtoi = eng_tokenizer.word_index
print('Total english words:', len(eng_wtoi))

max_eng_length = max(len(sentence) for sentence in eng_tokenized)
print('Max english sentence length:', max_eng_length)

num_eng_words = len(eng_wtoi) + 1

Total italian words: 8524
Max italian sentence length: 4
Total english words: 3858
Max english sentence length: 4
```

Padding sequences

we have to pad the sequences in order to have in a batch sentences which are of the same length (required for LSTM)

```
In [8]:
```

```
# Pad to max italian sentence length
ita_encoder_inputs = pad_sequences(ita_tokenized, maxlen=max_ita_length) → (we add some zeros at the beginning to obtain the length we want : length that we want = max_ita_length)

print("Italian encoder inputs shape:", ita_encoder_inputs.shape)

# Pad to max italian sentence length
eng_decoder_inputs = pad_sequences(eng_tokenized_train, maxlen=max_eng_length, padding='post')

print("English decoder inputs shape:", eng_decoder_inputs.shape)

Italian encoder inputs shape: (32499, 4)
English decoder inputs shape: (32499, 4)
```



```
In [9]:
```

```
# Pad to max english sentence length
eng_outputs = pad_sequences(eng_tokenized, maxlen=max_eng_length, padding='post')
```



```
In [10]:
```

```
ita_encoder_inputs
```

```
Out[10]: array([[ 0,  0,  0, 1211],
 [ 0,  0,  0, 2893],
 [ 0,  0,  0, 1693],
 ...,
 [ 0,  3, 2261, 1561],
 [ 0,  1,  861, 158],
 [ 0,  1,  500, 1210]], dtype=int32)
```

Model

Hyperparameter that we have to tune

```
In [11]:
```

```
# Build Encoder-Decoder Model
# -----
EMBEDDING_SIZE = 32 → each word is embedded with 32 values
# ENCODER
# -----
encoder_input = tf.keras.Input(shape=[max_ita_length])
encoder_embedding_layer = tf.keras.layers.Embedding(len(ita_wtoi)+1, EMBEDDING_SIZE, input_length=max_ita_length, mask_zero=True)
encoder_embedding_out = encoder_embedding_layer(encoder_input)
encoder = tf.keras.layers.LSTM(units=128, return_state=True)
encoder_output, h, c = encoder(encoder_embedding_out)
encoder_states = [h, c] → there will initialize the decoder state
# DECODER
# -----
decoder_input = tf.keras.Input(shape=[max_eng_length])
decoder_embedding_layer = tf.keras.layers.Embedding(len(eng_wtoi)+1, EMBEDDING_SIZE)
decoder_embedding_out = decoder_embedding_layer(decoder_input)
decoder_lstm = tf.keras.layers.LSTM(units=128, return_sequences=True, return_state=True)
# Initialize decoder state with final encoder state (initial_state=encoder_states)
decoder_lstm_out, _, _ = decoder_lstm(decoder_embedding_out, initial_state=encoder_states)
decoder_dense = tf.keras.layers.Dense(len(eng_wtoi)+1, activation='softmax')
decoder = decoder_dense(decoder_lstm_out)

# MODEL
model = tf.keras.Model([encoder_input, decoder_input], decoder)
```



```
In [13]:
```

```
model.summary()
#model.weights
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
--------------	--------------	---------	--------------

input_1 (InputLayer)	[(None, 4)]	0	
input_2 (InputLayer)	[(None, 4)]	0	
embedding (Embedding)	(None, 4, 32)	272800	input_1[0][0]
embedding_1 (Embedding)	(None, 4, 32)	123488	input_2[0][0]
lstm (LSTM)	[(None, 128), (None, 82432]	embedding[0][0]	
lstm_1 (LSTM)	[(None, 4, 128), (None, 82432]	embedding_1[0][0] lstm[0][1] lstm[0][2]	
dense (Dense)	(None, 4, 3859)	497811	lstm_1[0][0]

Total params: 1,058,963
Trainable params: 1,058,963
Non-trainable params: 0

Prepare model for training

```
In [14]: # Optimization params
# -----
# Loss
loss = tf.keras.losses.SparseCategoricalCrossentropy()
# Learning rate
lr = 1e-3
optimizer = tf.keras.optimizers.Adam(learning_rate=lr)

# Validation metrics
# -----
metrics = ['accuracy']

# Compile Model
model.compile(optimizer=optimizer, loss=loss, metrics=metrics)

In [15]:
import os
from datetime import datetime

cwd = os.getcwd()

exps_dir = os.path.join('/content/drive/My Drive/KerasRNN', 'translation_experiments')
if not os.path.exists(exps_dir):
    os.makedirs(exps_dir)

now = datetime.now().strftime('%b%d_%H-%M-%S')

exp_name = 'exp'

exp_dir = os.path.join(exps_dir, exp_name + '_' + str(now))
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

callbacks = []

# Model checkpoint
# -----
ckpt_dir = os.path.join(exp_dir, 'ckpts')
if not os.path.exists(ckpt_dir):
    os.makedirs(ckpt_dir)

ckpt_callback = tf.keras.callbacks.ModelCheckpoint(filepath=os.path.join(ckpt_dir, 'cp_{epoch:02d}.ckpt'),
                                                    save_weights_only=True) # False to save the model directly
callbacks.append(ckpt_callback)

# Visualize Learning on Tensorboard
# -----
tb_dir = os.path.join(exp_dir, 'tb_logs')
if not os.path.exists(tb_dir):
    os.makedirs(tb_dir)

# By default shows losses and metrics for both training and validation
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=tb_dir,
                                             profile_batch=0,
                                             histogram_freq=1) # if 1 shows weights histograms
callbacks.append(tb_callback)

# Early Stopping
# -----
early_stop = False
if early_stop:
    es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=10)
    callbacks.append(es_callback)

## Model
model.fit([ita_encoder_inputs, eng_decoder_inputs],
          eng_outputs,
          epochs=100,
          batch_size=128,
          validation_split=0.2,
          callbacks=callbacks)

# How to visualize Tensorboard
# 1. tensorboard --logdir EXPERIMENTS_DIR --port PORT      <- from terminal
# 2. Localhost:PORT  <- in your browser
```

Epoch 1/100
204/204 [=====] - 28s 113ms/step - loss: 6.2915 - accuracy: 0.2422 - val_loss: 4.9048 - val_accuracy: 0.2878

```

Epoch 2/100
204/204 [=====] - 21s 102ms/step - loss: 3.9382 - accuracy: 0.3221 - val_loss: 4.5691 - val_accuracy: 0.3451
Epoch 3/100
204/204 [=====] - 21s 103ms/step - loss: 3.4704 - accuracy: 0.4294 - val_loss: 4.3518 - val_accuracy: 0.3882
Epoch 4/100
204/204 [=====] - 21s 102ms/step - loss: 3.1697 - accuracy: 0.4632 - val_loss: 4.2485 - val_accuracy: 0.4080
Epoch 5/100
204/204 [=====] - 21s 102ms/step - loss: 2.9207 - accuracy: 0.4957 - val_loss: 4.0910 - val_accuracy: 0.4435
Epoch 6/100
204/204 [=====] - 21s 102ms/step - loss: 2.6697 - accuracy: 0.5335 - val_loss: 3.9836 - val_accuracy: 0.4630
Epoch 7/100
204/204 [=====] - 21s 101ms/step - loss: 2.4682 - accuracy: 0.5579 - val_loss: 3.9546 - val_accuracy: 0.4682
Epoch 8/100
204/204 [=====] - 21s 102ms/step - loss: 2.2908 - accuracy: 0.5787 - val_loss: 3.9104 - val_accuracy: 0.4741
Epoch 9/100
204/204 [=====] - 21s 102ms/step - loss: 2.1219 - accuracy: 0.6010 - val_loss: 3.8740 - val_accuracy: 0.4789
Epoch 10/100
204/204 [=====] - 21s 101ms/step - loss: 1.9805 - accuracy: 0.6175 - val_loss: 3.8369 - val_accuracy: 0.4846
Epoch 11/100
194/204 [=====>..] - ETA: 0s - loss: 1.8506 - accuracy: 0.6342

```

Translation (inference)

```

In [19]: # Uncomment this to Load model (use this if you want to restore saved model)
# model.load_weights(os.path.join('/content/drive/My Drive/KerasRNN', 'translation_experiments/exp_Dec04_02-03-59', 'ckptts/cp-58.ckpt'))}

# Modify the model such that the decoder takes predictions as inputs (no teacher forcing)

# ENCODER (remains the same)
# -----
encoder_model = tf.keras.Model(encoder_input, encoder_states)

# DECODER (modified)
# -----
decoder_state_input_h = tf.keras.Input(shape=[128])
decoder_state_input_c = tf.keras.Input(shape=[128])
decoder_state_inputs = [decoder_state_input_h, decoder_state_input_c]

decoder_input_single = tf.keras.Input(shape=[1])
decoder_input_single_embedding = decoder_embedding_layer(decoder_input_single)
decoder_outputs, h, c = decoder_lstm(decoder_input_single_embedding, initial_state=decoder_state_inputs)

decoder_states = [h, c]
decoder_outputs = decoder_dense(decoder_outputs)

decoder_model = tf.keras.Model([decoder_input_single] + decoder_state_inputs,
                               [decoder_outputs]+decoder_states)

```



```

In [20]: # Translation utils
ita_itow = {v:k for k, v in ita_wtoi.items()}
eng_itow = {v:k for k, v in eng_wtoi.items()} } this converts the integers to words

def translate(input_sentence):

    # Prepare input sentence
    input_tokenized = ita_tokenizer.texts_to_sequences([input_sentence])
    input_tokenized = pad_sequences(input_tokenized, maxlen=max_ita_length)

    # Get encoder state
    states_value = encoder_model.predict(input_tokenized)

    # Set first input '<sos>'
    curr_input = np.zeros([1, 1]) # bs x seq_length (1 x 1 at the beginning)
    curr_input[0, 0] = eng_wtoi['<sos>']
    eos = eng_wtoi['<eos>']

    output_sentence = []

    # Cycle until max_eng_length or until the '<eos>' is predicted
    for _ in range(max_eng_length):
        preds, h_, c_ = decoder_model.predict([curr_input]+states_value)
        word_id = np.argmax(preds[0, 0, :])

        if eos == word_id:
            break

        word = ''
        if word_id > 0:
            word = eng_itow[word_id]
            output_sentence.append(word)

        # Update next input with the predicted one
        curr_input[0, 0] = word_id
        # Update state
        states_value = [h_, c_]

    return ' '.join(output_sentence)

```



```

In [21]: print(translate("Ciao"))

welcome

```

Topics: text generator

```
In [1]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

In [2]: import os
import tensorflow as tf
import numpy as np

# Set the seed for random operations.
# This let our experiments to be reproducible.
SEED = 1234
tf.random.set_seed(SEED)
np.random.seed(SEED)

# Get current working directory
cwd = os.getcwd()

In [3]: from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

In [4]: !cat /content/drive/My\ Drive/Cantico_di_Natale.txt # Try to change the text with the one you prefer
```

Marley, prima di tutto, era morto. Niente dubbio su questo. Il registro mortuario portava le firme del prete, del chierico, dell'appaltatore delle pompe funebri e della persona che aveva guidato il mortoro. Scrooge vi aveva apposto la sua: e il nome di Scrooge, su qualunque fogliaccio fosse scritto, valeva tant'oro. Il vecchio Marley era proprio morto per quanto è morto, come diciamo noi, un chiodo di porta.

Badiamo! [...]

→ It uses an LSTM to produce the next character. It does not work on words but on chars. It produces the most probable following char given an input

Text generation - Next character prediction

Charles Dickens

```
In [5]: # Given the text we will:
# 1: divide the text into subsequences of N chars
# 2: build (input, target) pairs, where input is a sequence and target is the
#     same sequence shifted of 1 char
# 3: train a recurrent neural network to predict the next char after each input char
# 4: use the learned model to generate text
```

Dataset

```
# Prepare dataset
# -----
# Read full text
with open(os.path.join('/content/drive/My Drive/Cantico_di_Natale.txt'), 'r') as f:
    full_text = f.read()

full_text_length = len(full_text)
print('Full text length:', full_text_length)

# Create vocabulary
vocabulary = sorted(list(set(full_text)))

print('Number of unique characters:', len(vocabulary))
print(vocabulary)

# Dictionaries for char-to-int/int-to-char conversion
ctoi = {c:i for i, c in enumerate(vocabulary)}
itoc = {i:c for i, c in enumerate(vocabulary)}

# Create input-target pairs
# e.g., given an input sequence
# 'Hell' predict the next characters 'ello'
# Thus,
# extract from the full text sequences of length seq_length as x and
# the corresponding seq_length+1 character as target

# Define number of characters to be considered for the prediction
seq_length = 100

X = [] # will contain all the sequences
Y = [] # will contain for each sequence in X the next characters
X_enc = []
Y_enc = []
# Cycle over the full text
step = 1
for i in range(0, full_text_length - (seq_length), step):
    sequence = full_text[i:i+seq_length]
    target = full_text[i+1:i+seq_length+1]
    X.append(sequence)
    Y.append(target)
    X_enc.append([ctoi[c] for c in sequence])
    Y_enc.append([ctoi[c] for c in target])

X = np.array(X)
Y = np.array(Y)
X_enc = np.array(X_enc)
Y_enc = np.array(Y_enc)
```

print('Number of sequences in the dataset:', len(X))

Full text length: 158965
Number of unique characters: 69
['\n', ' ', '!', '"', '(', ')', ',', '.', '?', '!', 'A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'Z', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'z']

```
'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'y', 'z', 'È', 'à', 'è', 'é', 'ì', 'ò', 'ù', '')]
Number of sequences in the dataset: 158865
```

```
In [6]: print("Input Sequence: {}".format(X[0]))
print("Target Sequence: {}".format(Y[0]))

Input Sequence: Marley, prima di tutto, era morto. Niente dubbio su questo. Il registro mortuario portava le firme d
Target Sequence: arley, prima di tutto, era morto. Niente dubbio su questo. Il registro mortuario portava le firme de

In [7]: # Create data loaders
# -----
# Batch size
bs = 256

# Encode characters. Many ways, for example one-hot encoding.
def char_encode(x_, y_):
    return tf.one_hot(x_, len(vocabulary)), tf.one_hot(y_, len(vocabulary))

# Prepare input x to match recurrent layer input shape
# -> (bs, seq_length, input_size)

# Training
train_dataset = tf.data.Dataset.from_tensor_slices((X_enc, Y_enc))
train_dataset = train_dataset.shuffle(buffer_size=X_enc.shape[0])
train_dataset = train_dataset.map(char_encode)
train_dataset = train_dataset.batch(bs)
train_dataset = train_dataset.repeat()
```

Model

```
In [8]: # Build Recurrent Neural Network
# ----

# We build two models, one for training and one for inference. The two models
# SHARE THE WEIGHTS, thus the inference model will use the learned weights after training

# Training and inference model differ only for the initialization of the state
# in the Lstm. In the inference model the state of the Lstm is obtained through
# input layers. In this way, we can provide the prediction at time t-1 as input
# at time t for the generation of text at inference time.

# Model architecture (2 stacked Lstm layers): Input -> LSTM-1 -> LSTM-2 -> Dense

# Hidden size (state)
h_size = 128

# Model
input_x = tf.keras.Input(shape=(None, len(vocabulary)))

lstm1 = tf.keras.layers.LSTM(
    units=h_size, batch_input_shape=[None, None, len(vocabulary)],
    return_sequences=True, return_state=True, stateful=False)
lstm2 = tf.keras.layers.LSTM(
    units=h_size, return_sequences=True,
    return_state=True, stateful=False)
dense = tf.keras.layers.Dense(units=len(vocabulary), activation='softmax')

x, _, _ = lstm1(input_x)
x, _, _ = lstm2(x)
out = dense(x)

train_model = tf.keras.Model(
    inputs=input_x, outputs=out)

# Inference Model
h1_in = tf.keras.Input(shape=[h_size])
c1_in = tf.keras.Input(shape=[h_size])
h2_in = tf.keras.Input(shape=[h_size])
c2_in = tf.keras.Input(shape=[h_size])

x, h1, c1 = lstm1(input_x, initial_state=[h1_in, c1_in])
x, h2, c2 = lstm2(x, initial_state=[h2_in, c2_in])
out = dense(x)

inference_model = tf.keras.Model(
    inputs=[input_x, h1_in, c1_in, h2_in, c2_in],
    outputs=[out, h1, c1, h2, c2])
```

```
In [9]: train_model.summary()
inference_model.summary()
#model.weights

Model: "model"


| Layer (type)              | Output Shape                    | Param # |
|---------------------------|---------------------------------|---------|
| input_1 (InputLayer)      | [None, None, 69]                | 0       |
| lstm (LSTM)               | [None, None, 128], (None 101376 |         |
| lstm_1 (LSTM)             | [None, None, 128], (None 131584 |         |
| dense (Dense)             | (None, None, 69)                | 8901    |
| Total params: 241,861     |                                 |         |
| Trainable params: 241,861 |                                 |         |
| Non-trainable params: 0   |                                 |         |

Model: "model_1"


| Layer (type) | Output Shape | Param # | Connected to |
|--------------|--------------|---------|--------------|
|--------------|--------------|---------|--------------|


```

input_1 (InputLayer)	[(None, None, 69)]	0	
input_2 (InputLayer)	[(None, 128)]	0	
input_3 (InputLayer)	[(None, 128)]	0	
lstm (LSTM)	[(None, None, 128), 101376]	input_1[0][0] input_2[0][0] input_3[0][0]	
input_4 (InputLayer)	[(None, 128)]	0	
input_5 (InputLayer)	[(None, 128)]	0	
lstm_1 (LSTM)	[(None, None, 128), 131584]	lstm[1][0] input_4[0][0] input_5[0][0]	
dense (Dense)	(None, None, 69)	8901	lstm_1[1][0]
=====			
Total params:	241,861		
Trainable params:	241,861		
Non-trainable params:	0		

Prepare model for training

```
In [10]: # Optimization params
# -----
# Loss
loss = tf.keras.losses.CategoricalCrossentropy()

# Learning rate
lr = 1e-3
optimizer = tf.keras.optimizers.Adam(learning_rate=lr)

# Validation metrics
# -----
metrics = ['accuracy']

# Compile Model
train_model.compile(optimizer=optimizer, loss=loss, metrics=metrics)
```



```
In [11]: import os
from datetime import datetime

cwd = os.getcwd()

exps_dir = os.path.join('/content/drive/My Drive/KerasRNN', 'text_gen_experiments')
if not os.path.exists(exps_dir):
    os.makedirs(exps_dir)

now = datetime.now().strftime('%b%d_%H-%M-%S')

exp_name = 'exp'

exp_dir = os.path.join(exps_dir, exp_name + '_' + str(now))
if not os.path.exists(exp_dir):
    os.makedirs(exp_dir)

callbacks = []

# Model checkpoint
# -----
ckpt_dir = os.path.join(exp_dir, 'ckpts')
if not os.path.exists(ckpt_dir):
    os.makedirs(ckpt_dir)

ckpt_callback = tf.keras.callbacks.ModelCheckpoint(filepath=os.path.join(ckpt_dir, 'cp_{epoch:02d}.ckpt'),
                                                    save_weights_only=True) # False to save the model directly
callbacks.append(ckpt_callback)

# Visualize Learning on Tensorboard
# -----
tb_dir = os.path.join(exp_dir, 'tb_logs')
if not os.path.exists(tb_dir):
    os.makedirs(tb_dir)

# By default shows losses and metrics for both training and validation
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=tb_dir,
                                              profile_batch=0,
                                              histogram_freq=1) # if 1 shows weights histograms
callbacks.append(tb_callback)

# Early Stopping
# -----
early_stop = False
if early_stop:
    es_callback = tf.keras.callback.EarlyStopping(monitor='val_loss', patience=10)
    callbacks.append(es_callback)

# How to visualize Tensorboard
# 1. tensorboard --logdir EXPERIMENTS_DIR --port PORT      <- from terminal
# 2. Localhost:PORT  <- in your browser
```



```
In [12]: # It is difficult to have a validation metric to evaluate quantitatively the quality
# of a generation, since we are generating new text. Thus, for example, we can
# evaluate qualitatively the generation performance by generating some text during validation.
# We can do it through a custom callback which, at the beginning of each epoch (on_epoch_begin),
# will take a list of test sequences and use the trained model to generate some text.
```

```

class TextGenerationCallback(tf.keras.callbacks.Callback):
    def __init__(self, inference_model, start_sequences,
                 generation_length, temperature):
        super(TextGenerationCallback, self).__init__()
        self.inference_model = inference_model
        self.start_sequences = start_sequences
        self.generation_length = generation_length
        self.temperature = temperature

    def sample(self, pred, temperature=1.0):
        # Helper function to sample an index from a probability array
        # Temperature is a parameter that allows to scale the output of the network,
        # thus changing the level of 'exploration' in the output distribution. Try
        # yourself to play with this parameter and to see the difference in the generation.
        pred = np.asarray(pred).astype('float64')
        pred = np.log(pred) / temperature
        exp_pred = np.exp(pred)
        pred = exp_pred / np.sum(exp_pred)
        probas = np.random.multinomial(1, pred, 1)
        return np.argmax(probas)

    def generate_text(self, start_sequence, temperature):
        encoded_input = [ctoi[c] for c in start_sequence]
        encoded_input = tf.one_hot(encoded_input, len(vocabulary))
        encoded_input = tf.expand_dims(encoded_input, 0)

        generated_sequence = start_sequence

        in_states = [tf.zeros([1, h_size]) for _ in range(4)]

        for i in range(self.generation_length):
            output = self.inference_model.predict([encoded_input] + in_states)

            in_states = output[1:]
            pred = output[0][0, -1]
            # To get the final prediction we should use the argmax as usual but if we
            # do so, we will take always the most probable char, so we could incur into a generation loop.
            # Thus, instead of taking the max probability, we sample from the output distribution,
            # giving a chance to explore also other chars during the generation.
            # We do it with an helper function, that is the 'sample' one.
            pred = self.sample(pred, temperature)
            pred_char = itoc[pred]

            generated_sequence += pred_char

            encoded_input = tf.one_hot(pred, len(vocabulary))
            encoded_input = tf.reshape(encoded_input, [1, 1, len(vocabulary)])

        return generated_sequence

    def on_epoch_begin(self, epoch, logs):
        print("Epoch: {}".format(epoch))
        for start_seq in self.start_sequences:
            print("Starting Sequence: {}".format(start_seq))
            for temp in self.temperature:
                print("Temperature: {}".format(temp))
                generated_seq = self.generate_text(start_seq, temp)
                print(generated_seq)
        return

```

```
In [13]: valid_callback = TextGenerationCallback(inference_model=inference_model,
                                             start_sequences=['Un giorno'], generation_length=100,
                                             temperature=[0.2, 0.5, 1.0, 1.2])
callbacks.append(valid_callback)
```

```
In [14]: train_model.fit(x=train_dataset,
                      epochs=100, ##### set repeat in training dataset
                      steps_per_epoch=int(np.ceil(X_enc.shape[0] / bs)),
                      callbacks=callbacks)
```

```

Epoch 1/100
Epoch: 0
Starting Sequence: Un giorno
Temperature: 0.2
Un giornovkSOWEuC-,L!rlLuùlMOÈs
CsS(ETZCCdkùDv5dAs'IwùSéodGÈNUk.yR5MaOZe1NJZÈè.gEfay:èu!mpéubHkZ2ouudaòG;-M1
Temperature: 0.5
Un giornoàZél vùVppq?nr5,'kRÉSè,PsHùrn élGvhVùveùI-(vJefBVD!Hà:hWVpUftE(Aé,?lLRò5àAA;QNGz]phdÈì:G5nCÉ!WU,t I
Temperature: 1.0
Un giornopfèJofg:EUUvèRQT?IUNRnhTùvPoJisBgMOPò2SMU èTVQdopc2w?r'. -P
TisCÈsm;VyBèq!,?i)yUJvA!à(Qep(qm auM
Temperature: 1.2
Un giornoNp.vbùDragEMVAè-hSèP(.PcGzwòp-f;R2yzèL;i(wVk',UCOWmmaWf'ECSAFbwLckDÈlùhkQCnf?àbd -òmfPU
yFAC5pwB,;FO
 6/621 [.....] - ETA: 38s - loss: 4.1992 - accuracy: 0.1082WARNING:tensorflow:Callback method
`on_train_batch_end` is slow compared to the batch time (batch time: 0.0266s vs `on_train_batch_end` time: 0.0317s). Check
your callbacks.
621/621 [=====] - 44s 26ms/step - loss: 2.8979 - accuracy: 0.2129
Epoch 2/100
Epoch: 1
Starting Sequence: Un giorno
Temperature: 0.2
Un giorno di sua di con pitto di sera di suno di suo di che la porto di la sore di conte di sere al pio di co
Temperature: 0.5
Un giorno de pio di porto de fore le sprose, le sconte di soletto i cona che con en a cosse cino con è parela
Temperature: 1.0
Un giorno di sufiò i tropiero anrel litre. La no caccovera pittor a stofata sel hi endaszeto la senitto né so
Temperature: 1.2
Un giornoggio! L cone do veseno. Sun o ché qòrito al lasandii pidere, undo e è queste l'alalà, gnepe ana avet
621/621 [=====] - 34s 26ms/step - loss: 1.9851 - accuracy: 0.3846
Epoch 3/100
Epoch: 2
Starting Sequence: Un giorno

```

Temperature: 0.2
Un giorno di statto alla statto di sore di suo che la spette di sua con la state la si pore di si fuore a con
Temperature: 0.5
Un giorno di sun al fecchio da stesto il veste quella statti una gli che sua costara cami che nella cancarì d
Temperature: 1.0
Un giorno, scostandono. Alduosi glabbe detere quazza fornate, iganchio anghessa, lo Sproo sua gli porto dinqu
Temperature: 1.2
Un giorno, Fera elute ave angiole; si inchiaiva pio gagrizzali so con un gri afpote"
oss'ondbiata orr'vagone 1
621/621 [=====] - 34s 26ms/step - loss: 1.8042 - accuracy: 0.4386
Epoch 4/100
Epoch: 3
Starting Sequence: Un giorno
Temperature: 0.2
Un giorno di suo come si stato e se si seggiere in capo di Scrooge - disse Scrooge e si persa a cosa a cosa a
Temperature: 0.5
Un giorno e il porso che il volte lo stare il vorto il suo non che bisara con la suo del sobbeno a andito la
Temperature: 1.0
Un giorno spegna della raffico! -

Ha con logk Ma piritasté; me che tarore lo state, dute camverlo.

- Più se
Temperature: 1.2
Un giorno davrato quanno nevoa di rormusa sempo né più petterdori, l'ega come verfino. Volo te la drantato. F
621/621 [=====] - 34s 26ms/step - loss: 1.6925 - accuracy: 0.4708
Epoch 5/100
Epoch: 4
Starting Sequence: Un giorno
Temperature: 0.2
Un giorno a compressi di capo di sole stesso se non si pare di quella con un faccio del visito di stanza alla
Temperature: 0.5
Un giorno a culito e sentina con assua che soliva alla più al suo mano un pano alla stesse di biano di scorte
Temperature: 1.0
Un giorno al si fite, mando. Il cirsí uni una grati i somipità, poltesso e che basciò che sgazzando, fecchi d
Temperature: 1.2
Un giorno, fenso suovo. Qui che (profzano per buscio dubbato sololita: la quanto se guiti sebbiusi di trobben
621/621 [=====] - 34s 26ms/step - loss: 1.6082 - accuracy: 0.4959
Epoch 6/100
Epoch: 5
Starting Sequence: Un giorno
Temperature: 0.2
Un giorno di Scrooge di Scrooge di siano di tutto sulla stata per le parta di stato di casa di una dissero a
Temperature: 0.5
Un giorno di facer le parette di non si per sbalire della sarella della pirità di ogni stato il pasare a tarn
Temperature: 1.0
Un giorno di quel lento sulla strelche ebale al trava montrocceiliditi seccoso si dotare, Ti per vituto malo,
Temperature: 1.2
Un giorno predensi.

- Lra peimo. Besa avrescollo, efottà affirito in quanto Crottando? - Acconfinile, se, e
621/621 [=====] - 34s 26ms/step - loss: 1.5361 - accuracy: 0.5180
Epoch 7/100
Epoch: 6
Starting Sequence: Un giorno
Temperature: 0.2
Un giorno di stato il suo fatto di casa di Scrooge di sol solo così sopra di fare che si sentiva a casa di ri
Temperature: 0.5
Un giorno della fare che nella sua bascia di fare che la spiglia; ma non sono farle in quale se dove da caper
Temperature: 1.0
Un giorno ligerte gli si chiale luccia piati, chiema: Ebbele così verre a mane. Bob, Scrooge fattaggazzo e go
Temperature: 1.2
Un giornore la bionvana, per una vecchia Scietella e là qualchi il bocnava si forseno sella saruta saffotta?"
621/621 [=====] - 34s 27ms/step - loss: 1.4736 - accuracy: 0.5366
Epoch 8/100
Epoch: 7
Starting Sequence: Un giorno
Temperature: 0.2
Un giorno di quella stata per le stato in cortere la figlia di più freddi di quale di conto di solito di quel
Temperature: 0.5
Un giorno sempre in un bambino della stesse e sedere davanti in cillere il pensole si spiriti. Il nipote di t
Temperature: 1.0
Un giorno puttro solo quella soppresse s'ardervi quel cario le lentasse.

- gli mia sognir la meso una cotest
Temperature: 1.2
Un giorno; un mio buone primatanza non io nei volesse littuni, da catra matuta chi riveretta per la parso alm
621/621 [=====] - 34s 27ms/step - loss: 1.4160 - accuracy: 0.5532
Epoch 9/100
Epoch: 8
Starting Sequence: Un giorno
Temperature: 0.2
Un giorno al morto come un po' di fanciulla, come se non si prese che solo in casa di Scrooge di Scrooge si s
Temperature: 0.5
Un giorno delle al fuoco per stato il pansola e di fu per vedere di tutto in altro commetto e il cappello. Sc
Temperature: 1.0
Un giorno al sposso. Ah, ah, badessero La sfragativa a c'è sariune."

Senza quando con Natale. Voi vostra rag
Temperature: 1.2
Un giorno alche ero consutino mondato un'altro...

- È Lai giocoscie, lintro, un suoro, mi pendato. Sai dispe
212/621 [=====] - ETA: 10s - loss: 1.3688 - accuracy: 0.5670