

01 Permutation Tests - Two Independent Populations

```
# We sample n1 data from a first population X1 and n2 data from a second population X2
# Test:
# H0: X1 =^d X2
# H1: X1 !=^d X2

#### -----
### Case 1. H0 FALSE
####

# To understand the behaviour of the test, we sample from two populations
# with different means

# Parameters:
n1 <- 10
n2 <- 10
n <- n1 + n2

# Simulation:
set.seed(240279)
x1 <- runif(n1, 0, 4) ~ U([0,4])
x2 <- runif(n2, 0, 4)+3 ~ U([3,7]) }

x_pooled <- c(x1, x2)

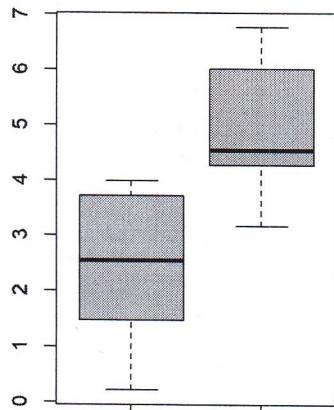
par(mfrow=c(1,2))
boxplot(x1, x2, main='Original data')

# How data change if we apply one random permutation?
permutation <- sample(1:n)

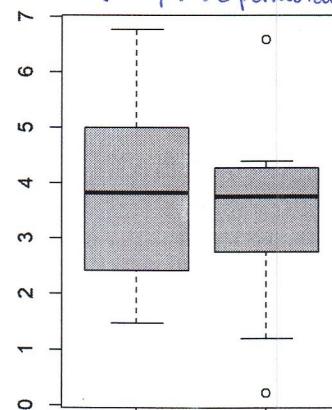
x_perm <- x_pooled[permutation]
x1_perm <- x_perm[1:n1]
x2_perm <- x_perm[(n1+1):n]

boxplot(x1_perm, x2_perm, main='Permuted data')
```

Original data



Permuted data
(one possible permutation)



We see that by permuting data the two distributions seem to change. If this change appears very often while permuting we end with the conclusion that the original data were kind of "extreme" ($\Rightarrow H_1$)

\Rightarrow We want to compare the boxplot of the original data with (theoretically) all the boxplots of all the possible permutations. We want an indicator of how the boxplots are different, for instance we can use the absolute value of the difference of the means

$$(|\bar{x}_1 - \bar{x}_2|)$$

```
abs(mean(x1)-mean(x2))
```

```
## [1] 2.465564
```

```
abs(mean(x1_perm)-mean(x2_perm))
```

```
## [1] 0.4176959
```

(a lot smaller than the original one)

```
# The mean difference is lower. Was that a case?
```

```
# TEST
```

```
# Test statistic: absolute difference between the two means
T0 <- abs(mean(x1) - mean(x2))
```

```
T0
```

```
# [1] 2.465564
```

```
# Cardinality of the permutational space:  
factorial(n)
```

```
## [1] 2.432902e+18 = # possible permutations
```

```
# Number of distinct values of T*:  
factorial(n)/(2*factorial(n1)*factorial(n2))
```

```
## [1] 92378
```

```
# Minimum achievable p-value:  
1/(factorial(n)/(2*factorial(n1)*factorial(n2)))
```

```
## [1] 1.082509e-05
```

```
# CMC to estimate the p-value  
B <- 100000 # Number of permutations  
T_stat <- numeric(B) # Vector where we will store the values of T*
```

```
# To estimate the p-value we use a Loop  
# Inside the Loop, we do the following:  
# 1. choose a random permutation of data (with the command sample)  
# 2. calculate and save the test statistic obtained with the permuted data  
for(perm in 1:B){  
  # permutation:  
  permutation <- sample(1:n)  
  x_perm <- x_pooled[permutation]  
  x1_perm <- x_perm[1:n1]  
  x2_perm <- x_perm[(n1+1):n]  
  # test statistic:  
  T_stat[perm] <- abs(mean(x1_perm) - mean(x2_perm))  
}
```

```
# Permutational distribution of T  
hist(T_stat,xlim=range(c(T_stat,T0)),breaks=30)  
abline(v=T0,col=3,lwd=2)  
  
plot(ecdf(T_stat))  
abline(v=T0,col=3,lwd=2)
```

the test statistic is invariant w.r.t. the permutations within each sample \Rightarrow number of possible stats = $\frac{n!}{2(n_1!)(n_2!)}$

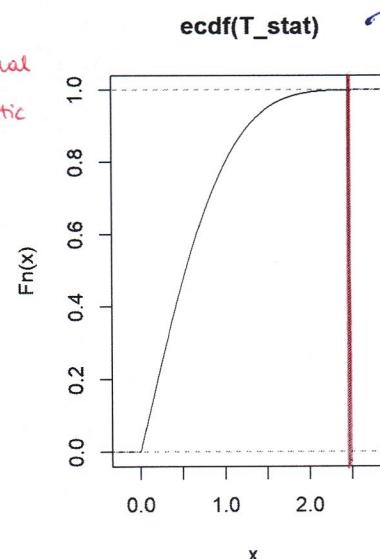
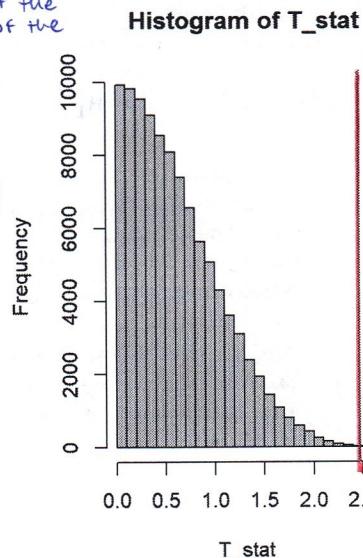
most extreme case, where the sample we have is the most extreme one. This is also the resolution of our p-value. The p-value can assume only:

$$\frac{1}{n!/2n_1!n_2!}, \frac{2}{n!/2n_1!n_2!}, \dots$$

this 2∞ because the test statistic has an absolute value and so also the swapping of the samples will generate the same value of the test statistic ($|\bar{x}_1 - \bar{x}_2| = |\bar{x}_2 - \bar{x}_1|$)

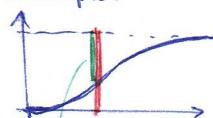
The resolution of the p-value can be an issue: if n is too small (and consequently also n_1 and n_2) the p-value will be able to assume only few discrete values. We want to avoid the discretization of the p-value.
(Case of low resolution of p-value)

estimate of the distribution of the test statistic obtained by permuting data in all the possible permutations)



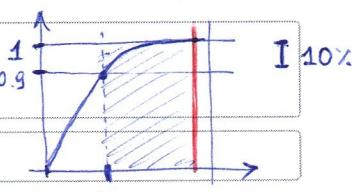
cumulative distribution function of the permutational distribution

the p-value is the length of the segment (on the red line) going from the black line (ecdf) to the "—" line:
example:

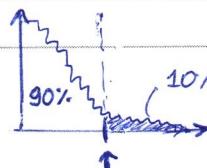


the p-value is the length of this segment

If we want a critical region of the level of 10% (test with prob. of type I error equal to 10%):



equivalent to go to the histogram of Tstat and look for the point that:



from this value it begins the critical region

```
# p-value  
p_val <- sum(T_stat>=T0)/B  
p_val
```

```
## [1] 3e-04
```

our original sample is very extreme
we reject H₀

```

#### -----
### Case 2. H0 TRUE
#### -----
# now we simulate data from two populations with the same distribution

# Simulation:
set.seed(240279)
x1 <- runif(n1,0,4)
x2 <- runif(n2,0,4)+0
x_pooled <- c(x1,x2)

par(mfrow=c(1,2))
boxplot(x1,x2,main='Original data')

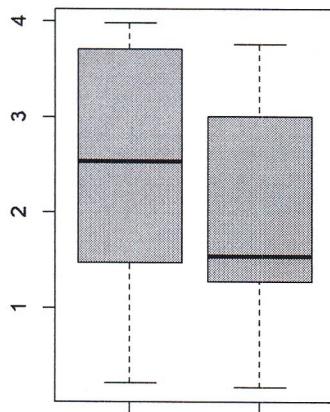
# How data change if we apply one random permutation?
permutation <- sample(1:n)

x_perm <- x_pooled[permutation]
x1_perm <- x_perm[1:n1]
x2_perm <- x_perm[(n1+1):n]

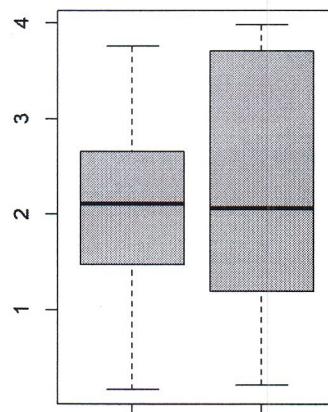
boxplot(x1_perm,x2_perm,main='Permuted data')

```

Original data



Permuted data



this time the difference between the boxplots is not that big - If this is confirmed by a lot of permutations than our original sample is not that extreme.

```
abs(mean(x1)-mean(x2))
```

```
## [1] 0.5344363
```

```
abs(mean(x1_perm)-mean(x2_perm))
```

```
## [1] 0.1823041
```



the permuted sample is even more extreme than the original

```
# TEST
```

```
# Test statistic: absolute difference between the two means
T0 <- abs(mean(x1) - mean(x2))
T0
```

```
## [1] 0.5344363
```

```
# Cardinality of the permutational space:
factorial(n)
```

```
## [1] 2.432902e+18
```

```
# Number of distinct values of T*:
factorial(n)/(2*factorial(n1)*factorial(n2))
```

```
## [1] 92378
```

```
# Minimum achievable p-value:
1/(factorial(n)/(2*factorial(n1)*factorial(n2)))
```

```

## [1] 1.082509e-05

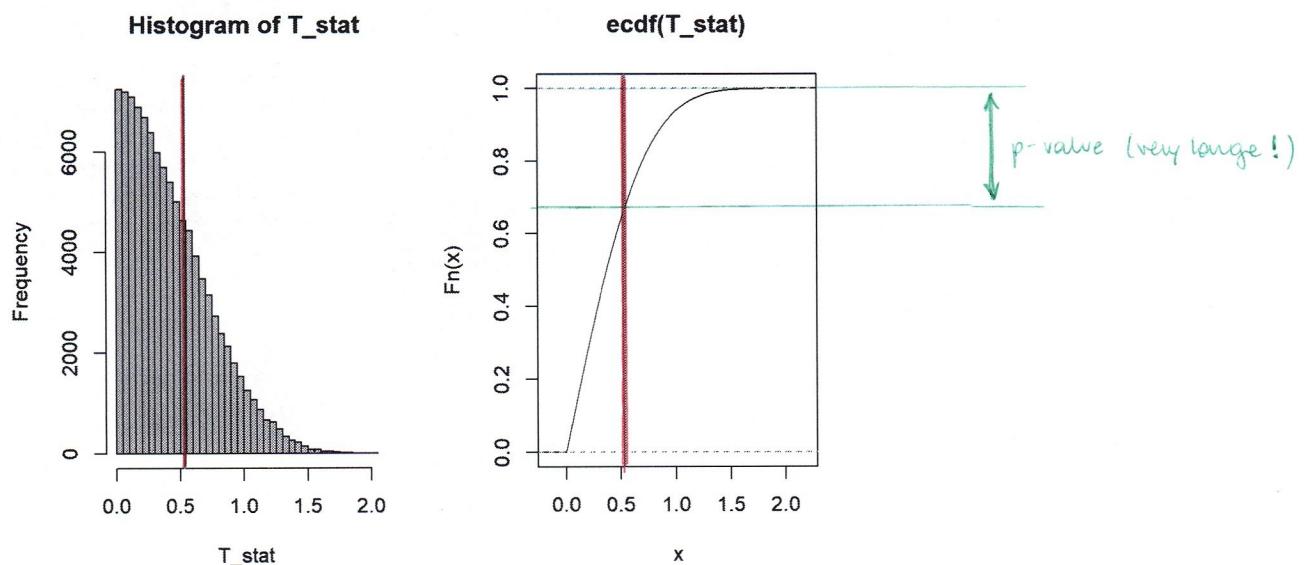
# CMC to estimate the p-value
B <- 100000 # Number of permutations
T_stat <- numeric(B) # Vector where we will store the values of T^

for(perm in 1:B){
  # permutation:
  permutation <- sample(1:n)
  x_perm <- x_pooled[permutation]
  x1_perm <- x_perm[1:n1]
  x2_perm <- x_perm[(n1+1):n]
  # test statistic:
  T_stat[perm] <- abs(mean(x1_perm) - mean(x2_perm))
}

# Permutational distribution of T
hist(T_stat,xlim=range(c(T_stat,T0)),breaks=30)
abline(v=T0,col=3,lwd=2)

plot(ecdf(T_stat))
abline(v=T0,col=3,lwd=2)

```



```

# p-value
p_val <- sum(T_stat>=T0)/B
p_val

```

[1] 0.33116 → this means that 33% of the times that we permuted the data we got a value of the test statistic more extreme than the original one

Notice: as we already pointed out in the theory, the distribution of the T-stat depends on the data! In the two examples we have the same procedure but different data
 ⇒ the two histograms of T-stat are different!
 (one goes to zero ~ 2.5, the other ~ 2.0 / 1.8)

```

### -----
### -----
### Nonparametric Inference
### -----
# Let's reprise again why should I use a nonparametric test, instead of a parametric one...
# the idea is that the real probability of committing a type I error (alpha...) should be
# Less or equal of the nominal one...

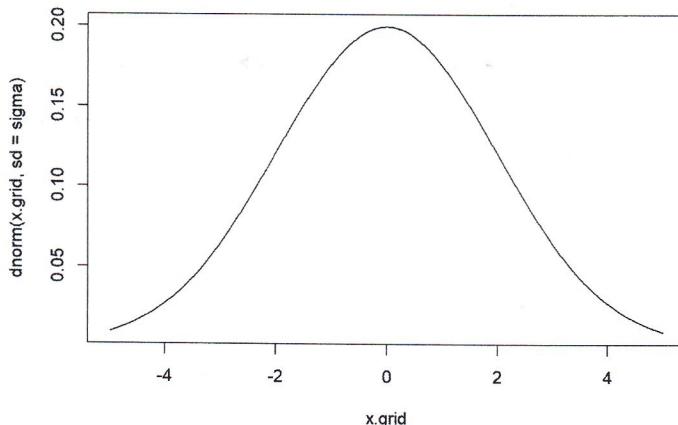
### -----
### Assess Empirical Level of Type-I error
### -----
# Let's work with 2 sample t-tests
library(progress) ← useful for the progress bars (for computations)

rm(list=ls())
alpha = 0.05
n     = 10
B     = 1000
seed  = 2781991
sigma = 2

### Normal data
### -----
x.grid <- seq(-5, 5, by=0.01)
plot(x.grid, dnorm(x.grid, sd=sigma), type='l')

```

Goal: We want to compare two samples from a given distribution (we want to compare their mean)



We estimate the distribution of the p-values in these settings:

```

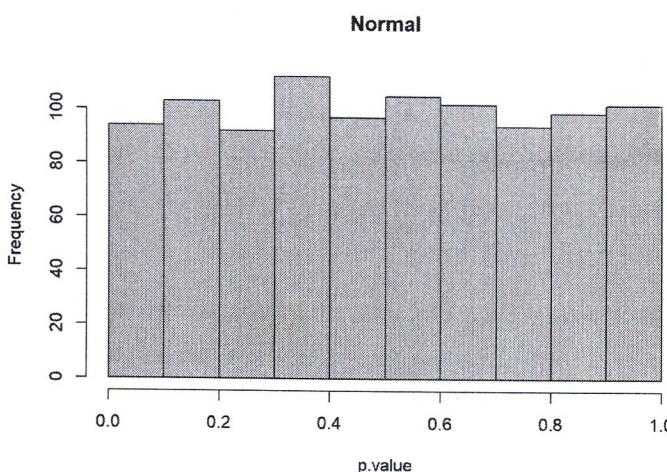
p.value <- numeric(B)

pb = progress_bar$new(total=B)
pb$tick(0)

set.seed(seed)
for(j in 1:B){
  x1 = rnorm(n, sd=sigma)
  x2 = rnorm(n, sd=sigma)
  p.value[j] = t.test(x1, y=x2)$p.value
  pb$tick()
}

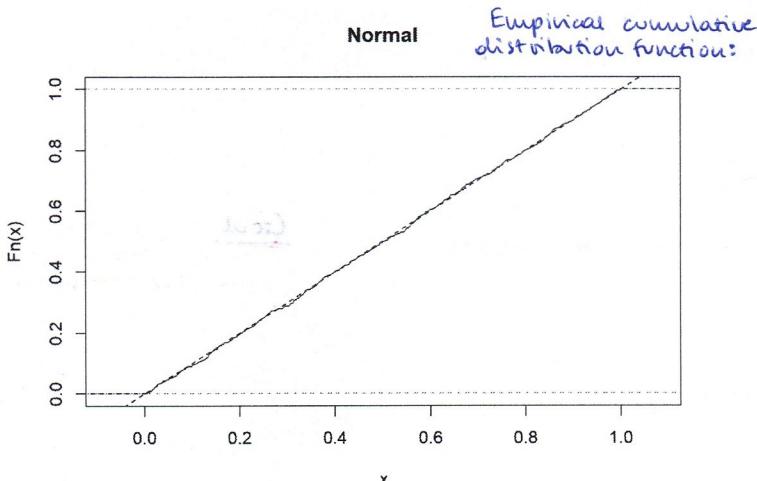
hist(p.value, main = 'Normal')

```



→ The p-values are distributed as a uniform (and so the test is calibrated)

```
plot(ecdf(p.value), main = 'Normal')
abline(0,1, lty=2, col='red')
```



```
estimated.alpha <- sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))
```

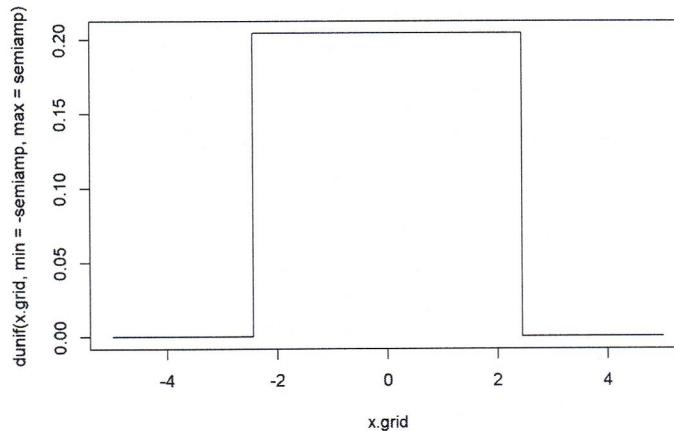
```
## [1] 0.03388273 0.04700000 0.06011727
```

```
### Leptocurtic data (e.g. Uniform)
```

```
## -----
# Let's keep the same variance, for uniformity...
semiamp = sqrt(6)
plot(x.grid, dunif(x.grid,min=-semiamp,max=semiamp), type='l')
```

Empirical type I error rate

We can't refute the hypothesis that the empirical level is equal to the theoretical one

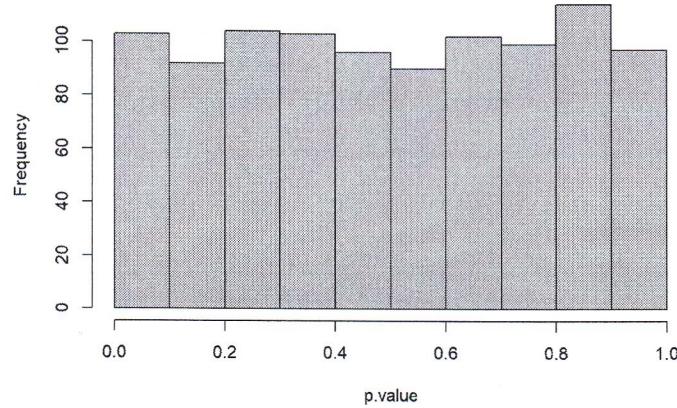


```
p.value <- numeric(B)
pb=progress_bar$new(total=B)
pb$tick(0)
```

```
set.seed(seed)
for(j in 1:B){
  x1 = runif(n,min=-semiamp,max=semiamp)
  x2 = runif(n,min=-semiamp,max=semiamp)
  p.value[j] = t.test(x1,y=x2)$p.value
  pb$tick()
}
```

```
hist(p.value, main = 'Uniform')
```

Uniform

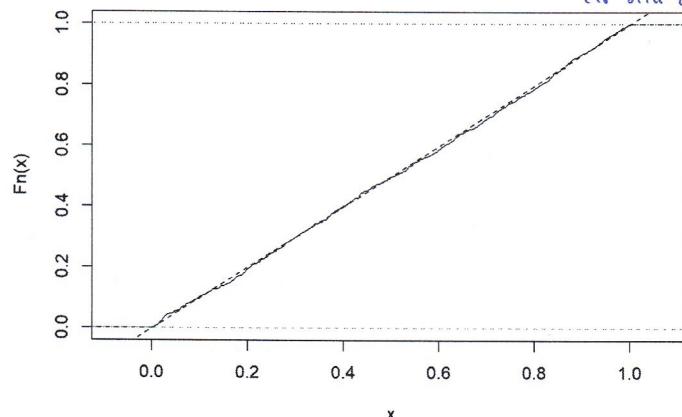


The uniformity of the p-values
still stands (a bit)

```
plot(ecdf(p.value), main = 'Uniform')
abline(0,1, lty=2, col='red')
```

Uniform

Empirical cumulative
distribution function
(it still (nearly) okay)



```
estimated.alpha <- sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))
```

```
## [1] 0.0408699 0.0550000 0.0691301
```

```
# tends not to be a problem with "thin-tailed" data... but!
```

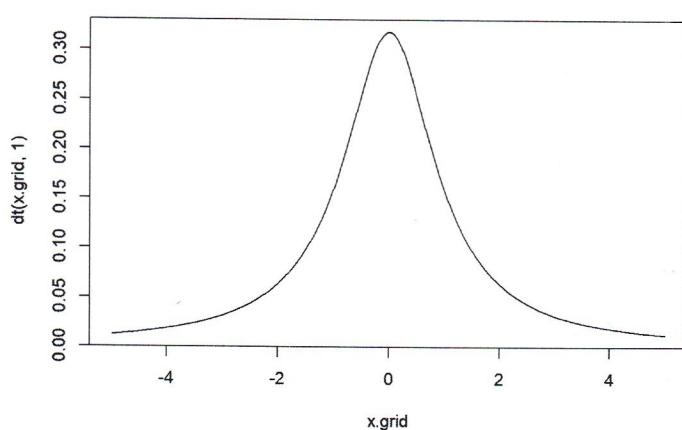
```
### Platicurtic data (e.g. Student's t)
```

(distributions with fat tails)

```
##
```

```
x.grid <- seq(-5, 5, by=0.01)
plot(x.grid, dt(x.grid, 1), type='l')
```

t-student distribution:



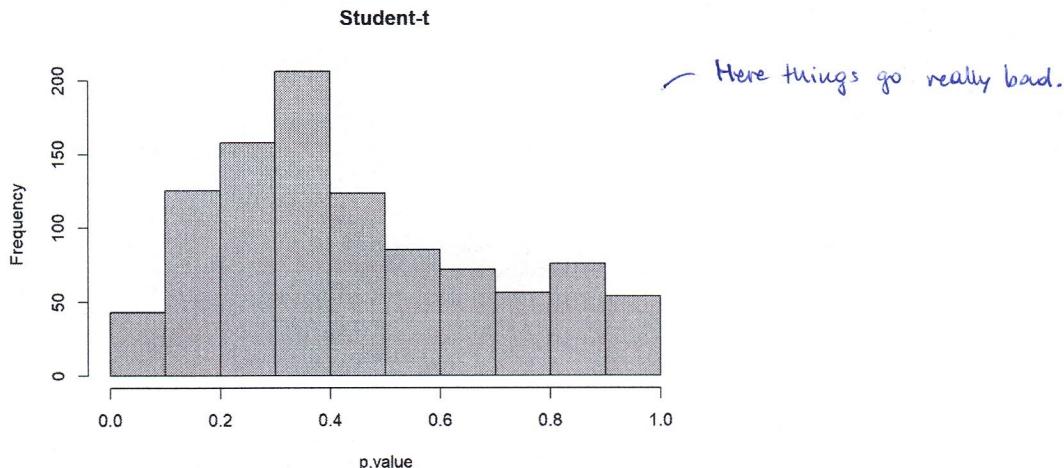
```

p.value <- numeric(B)
pb=progress_bar$new(total=B)
pb$tick(0)

set.seed(seed)
for(j in 1:B){
  x.1 = rt(n, 1)
  x.2 = rt(n,1)
  p.value[j] = t.test(x.1,y=x.2)$p.value
  pb$tick()
}

hist(p.value, main = 'Student-t')

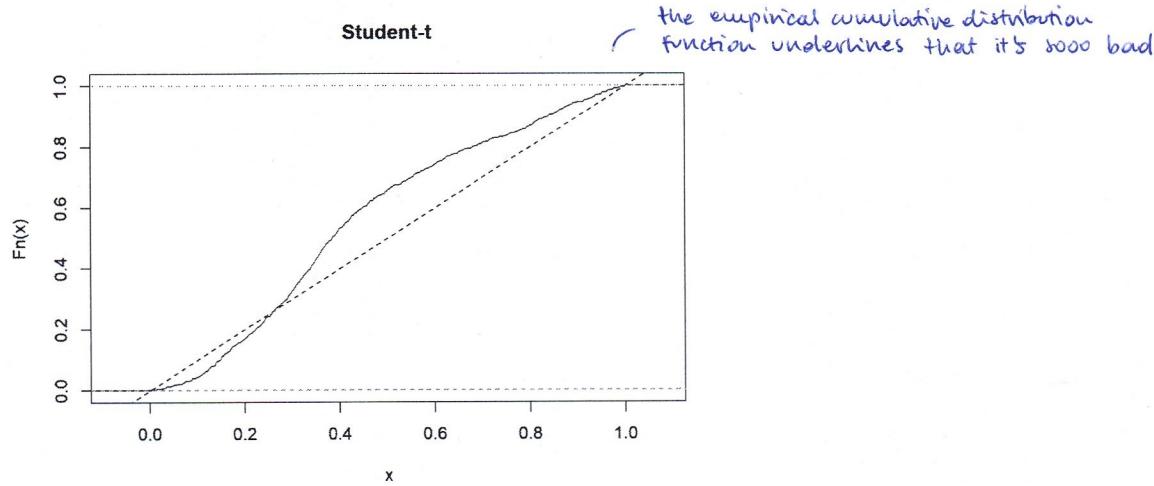
```



```

plot(ecdf(p.value), main = 'Student-t')
abline(0,1, lty=2, col='red')

```



```

estimated.alpha <- sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))

```

```
## [1] 0.009759756 0.018000000 0.026240244
```

very very low

```
# I can actually replicate this with several cases...
# install.packages('stabledist')
library(stabledist)
```

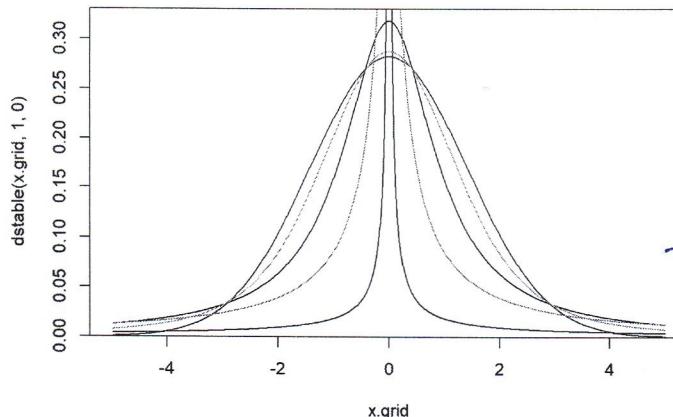
← package used for simulate from "stable" distributions

```
# stabledist is a quite useful package used to simulate from "stable" distributions, where a
# stable distribution intended is in the Levy sense: i.e. if a Linear combination of two
# variables generated from a stable distr P is still distributed as P.
# Parametrisation is done with:
#   alpha = stability, from 0 to 2 (1 → Cauchy, 2 → normal)
#   beta = skewness parameter
#   mu = "mean", mean is actually defined only for alpha>1
#   variance is defined only for alpha=2
# very useful to simulate heavy tailed data
# special cases are...
```

```

x.grid <- seq(-5, 5, by=0.01)
plot(x.grid, dstable(x.grid,1,0), type='l') #cauchy
lines(x.grid, dstable(x.grid,2,0), type='l',col="red") #normal
lines(x.grid, dstable(x.grid,1.5,0), type='l',col="orange")
lines(x.grid, dstable(x.grid,0.5,0), type='l',col="green")
lines(x.grid, dstable(x.grid,0.1,0), type='l',col="blue")

```

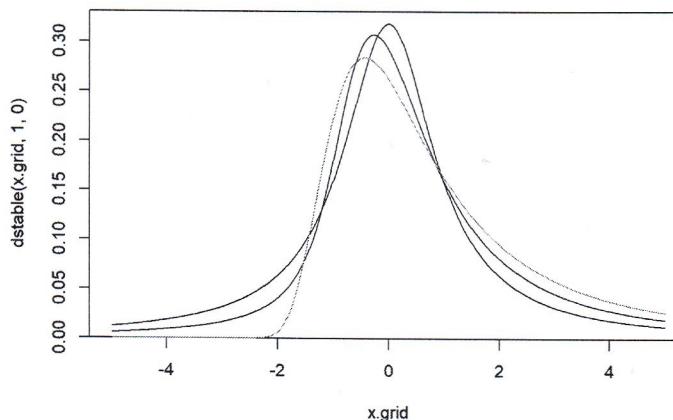


we have some heavy tails distributions
that we can use to check some results
(to check the validity of some results
under "extreme" distributions)

```

plot(x.grid, dstable(x.grid,1,0), type='l') #cauchy
lines(x.grid, dstable(x.grid,1,0.5), type='l',col="red") #normal
lines(x.grid, dstable(x.grid,1,1), type='l',col="orange")

```



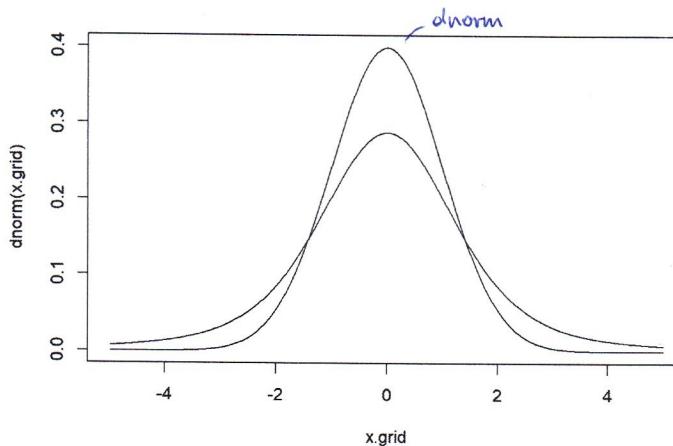
Let's try to run a test on a stable with alpha=1.5 (mean, no variance) 1.
and on a alpha=0.5 (no mean, no variance) 2.

1.

```

x.grid <- seq(-5, 5, by=0.01)
plot(x.grid, dnorm(x.grid), type='l')
lines(x.grid, dstable(x.grid,1.5,0), type='l',col='red')

```



```

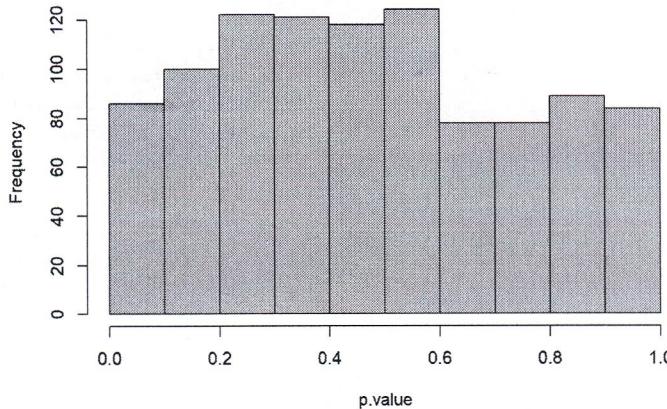
p.value <- numeric(B)
pb=progress_bar$new(total=B)
pb$tick(0) } this is for the progress bar

set.seed(seed)
for(j in 1:B){
  x.1 = rstable(n,1.5,0)
  x.2 = rstable(n,1.5,0)
  p.value[j] = t.test(x.1,y=x.2)$p.value
} if we add: pb$tick() ← it prints the progress bar

hist(p.value, main = 'Stable - Alpha=1.5')

```

Stable - Alpha=1.5

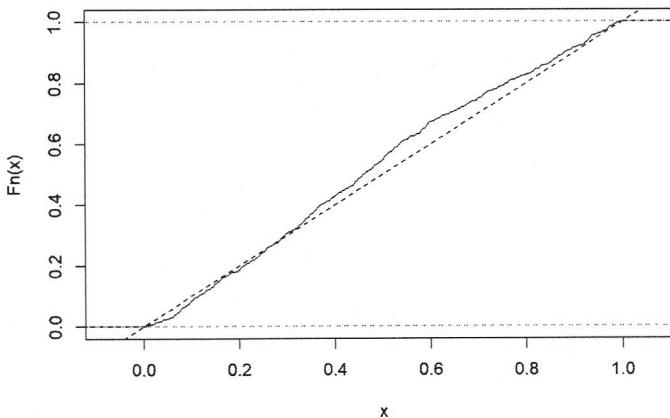


```

plot(ecdf(p.value), main = 'Alpha=1.5')
abline(0,1, lty=2, col='red')

```

Alpha=1.5



```

estimated.alpha <- sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))

```

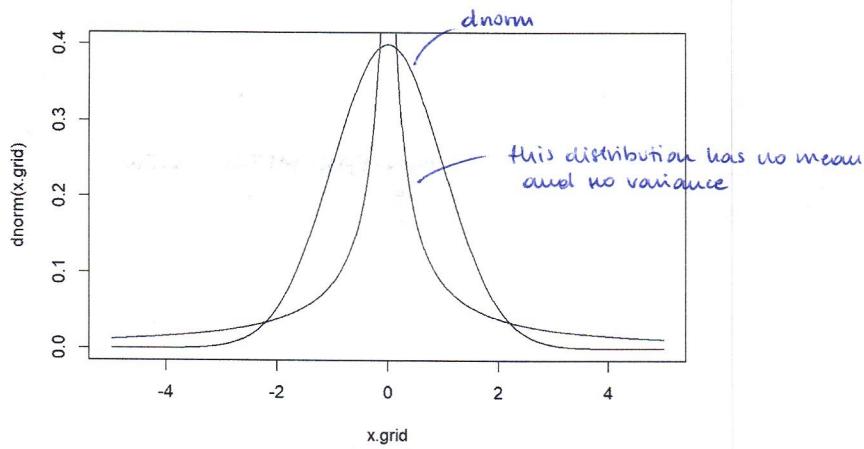
```
## [1] 0.01532345 0.02500000 0.03467655
```

```

plot(x.grid, dnorm(x.grid), type='l')
lines(x.grid, dstable(x.grid,.5,0), type='l', col='red')

```

Example for
which we'll try
nonparametric
settings := Ex.



```

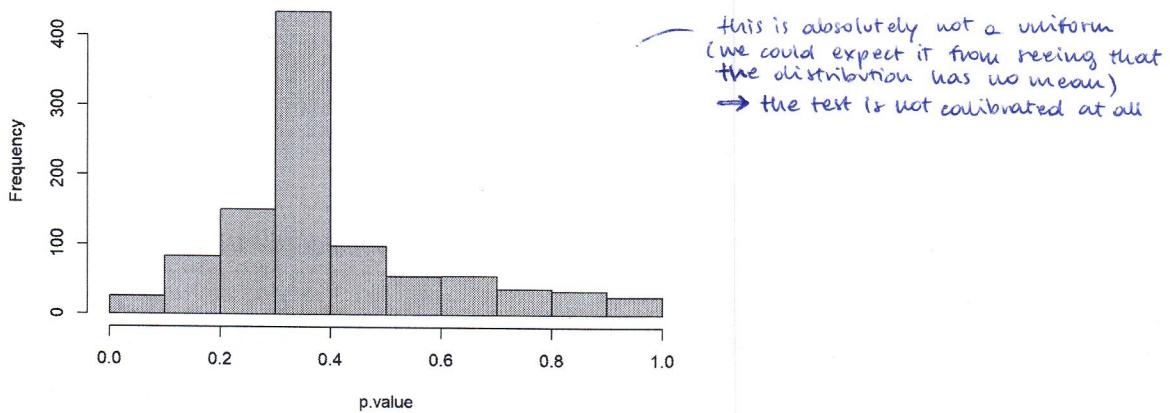
p.value <- numeric(B)
pb=progress_bar$new(total=B)
pb$tick(0)

set.seed(seed)
for(j in 1:B){
  x.1 = rstable(n,.5,0)
  x.2 = rstable(n,.5,0)
  p.value[j] = t.test(x.1,y=x.2)$p.value
  pb$tick()
}

hist(p.value, main = 'Stable - Alpha = 0.5')

```

Stable - Alpha = 0.5



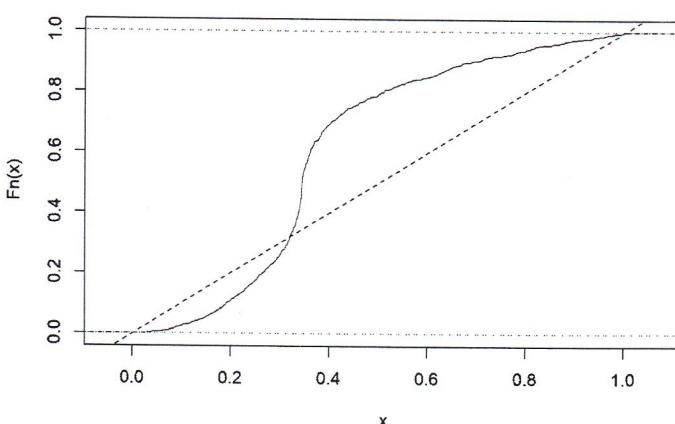
```

plot(ecdf(p.value), main = 'Alpha = 0.5')
abline(0,1, lty=2, col='red')

```

Alpha = 0.5

↙ this does not have any taste at all



```

estimated.alpha <- sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))

```

```
## [1] 0.001213513 0.006000000 0.010786487
```

```

# Let's keep this last example as a testbed...
# t-test goes bananas, does not cover at all.
# Let's try to compare the situation with nonparametric tests...
# Let's start with Wilcoxon...

```

```

p.value <- numeric(B)
pb=progress_bar$new(total=B)
pb$tick(0)

```

```

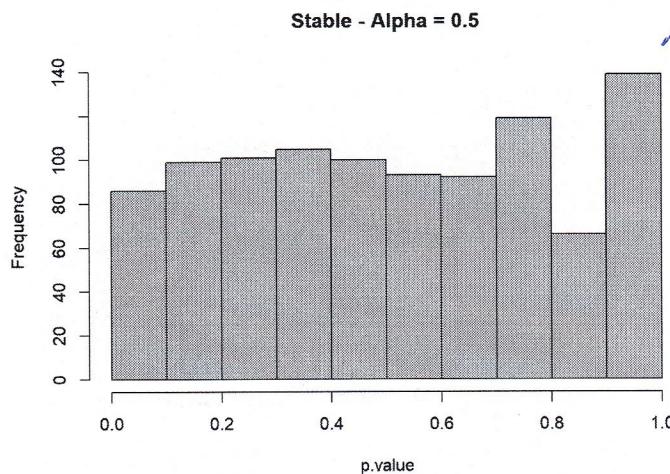
set.seed(seed)
for(j in 1:B){
  x.1 = rstable(n,.5,0)
  x.2 = rstable(n,.5,0)
  p.value[j] = wilcox.test(x.1,y=x.2,correct = T)$p.value
  pb$tick()
}

```

```
hist(p.value, main = 'Stable - Alpha = 0.5')
```

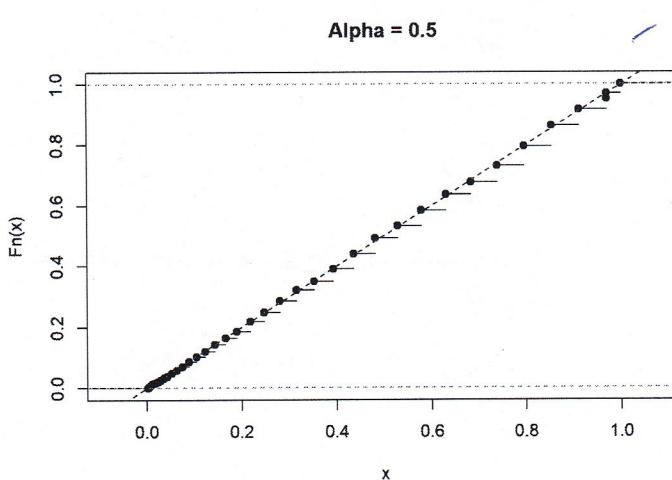
NON-PARAMETRIC VIEW

- Wilcoxon test
- Permutational test



✓ still some issues in the calibration
(but it's just something related to the quantization of the p-value)

```
plot(ecdf(p.value), main = 'Alpha = 0.5')
abline(0,1, lty=2, col='red')
```



✓ the p-value is quantized, so it's not a continuous function like before (# t-test)
(the issues with the histogram are related to that)

```

estimated.alpha <- sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))

```

```
## [1] 0.02530064 0.03700000 0.04869936
```

← way better than the t-test

Permutation test (MEAN)

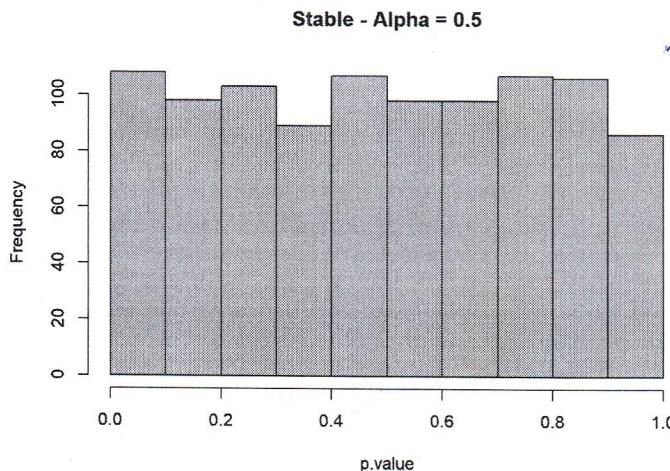
```
# quite unfortunately for you, there are no good implementations in R for permutation tests...
# so we need to "weaponise" prof.

# Vantinis code
perm_t_test=function(x,y,iter=1e4){
  T0      = abs(mean(x)-mean(y))
  T_stat  = numeric(iter)
  x_pooled = c(x,y)
  n       = length(x_pooled)
  n1      = length(x)
  for(perm in 1:iter){
    # permutation:
    permutation <- sample(1:n)
    x_perm      <- x_pooled[permutation]
    x1_perm     <- x_perm[1:n1]
    x2_perm     <- x_perm[(n1+1):n]
    # test statistic:
    T_stat[perm] <- abs(mean(x1_perm) - mean(x2_perm))
  }
  # p-value
  p_val <- sum(T_stat>=T0)/iter
  return(p_val)
}

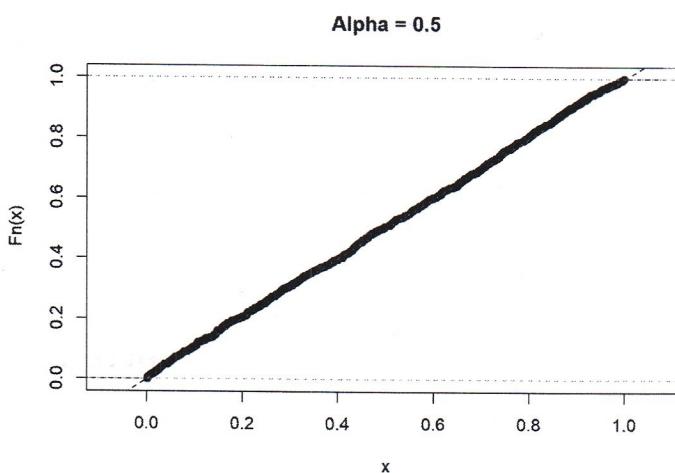
p.value <- numeric(B)
pb=progress_bar$new(total=B)
pb$tick(0)

set.seed(seed)
for(j in 1:B){
  x.1 = rstable(n,.5,0)
  x.2 = rstable(n,.5,0)
  p.value[j] = perm_t_test(x.1,x.2, iter=1000)
  pb$tick()
}

hist(p.value, main = 'Stable - Alpha = 0.5')
```



```
plot(ecdf(p.value), main = 'Alpha = 0.5')
abline(0,1, lty=2, col='red')
```



```

estimated.alpha <- sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))

```

```
## [1] 0.04174957 0.05600000 0.07025043
```

→ we cannot refuse the hypothesis of the empirical α being equal to the theoretical one

```

## -----
### Assess Empirical Level of Type-II error
## -

```

```
# How to choose between the two (at least in this specific setting?) Power calculations!
# i.e. we fix the data and the significance level, cycle over a grid of effect sizes, and see what is the most powerful in identifying things...
```

```

delta_grid=c(1,2,3,4,5)
set.seed(seed)
power_wilcox=numeric(length(delta_grid))
for(ii in 1:length(delta_grid)){
  p.value <- numeric(B)
  pb=progress_bar$new(total=B)
  pb$tick(0)
  delta=delta_grid[ii]
  for(j in 1:B){
    x.1 = rstable(n,.5,0)
    x.2 = rstable(n,.5,0,delta=delta)
    p.value[j] <- wilcox.test(x.1,y=x.2,correct = T)$p.value
    pb$tick()
  }
  estimated.power <- sum(p.value < alpha)/B
  power_wilcox[ii]=estimated.power
}

```

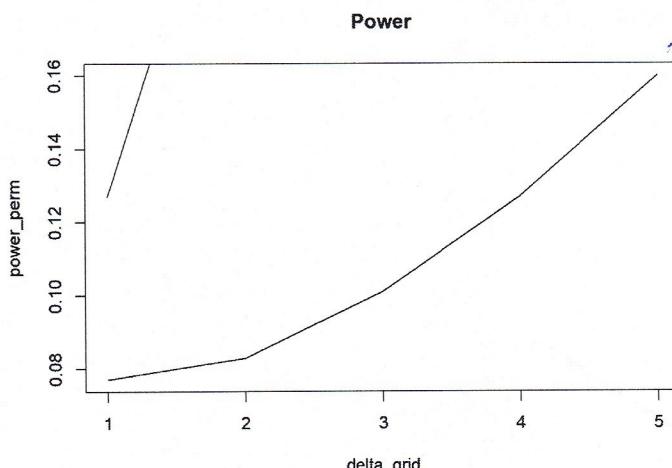
→ if we add here: "power_wilcox" we get [0.127, 0.243, 0.324, 0.408, 0.425]

```

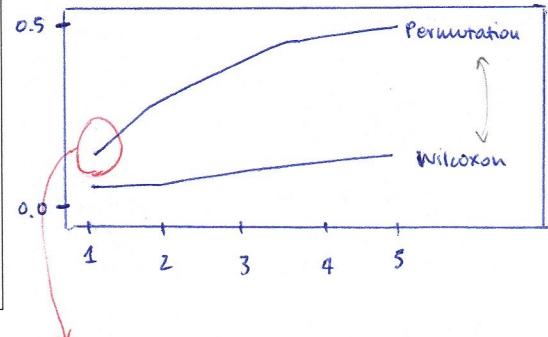
power_perm=numeric(length(delta_grid))
set.seed(seed)
for(ii in 1:length(delta_grid)){
  p.value <- numeric(B)
  pb=progress_bar$new(total=B)
  pb$tick(0)
  delta=delta_grid[ii]
  for(j in 1:B){
    x.1 = rstable(n,.5,0)
    x.2 = rstable(n,.5,0,delta=delta)
    p.value[j] = perm_t_test(x.1,x.2,iter=1000)
    pb$tick()
  }
  estimated.power <- sum(p.value < alpha)/B
  power_perm[ii]=estimated.power
}

```

```
plot(delta_grid,power_perm,type='l', main='Power')
lines(delta_grid,power_wilcox,col='red')
```



Setting up the right limits:



However here it's still too small. That's probably because we're using as a test statistic something that consider the mean, when our distribution DOES NOT ALLOW A MEAN!

→ We change the test statistic into something that consider the MEDIAN instead of the MEAN.

Note: qua mi sa che quello in alto in realtà è Wilcoxon, comunque sia il problema che si ha tra media e mediana rimane (si mischia e Wilcoxon perché quando andiamo di mediana (e non media) il permutation migliore, il Wilcoxon rimane)

Conclusion: moving under nonparametric settings we improved and we're now able to test something that we were not able to test before with a control of the type I error

We know that the t-test for EX went crazy, the (nonparametric) Wilcoxon was already better, while the perm. was good.

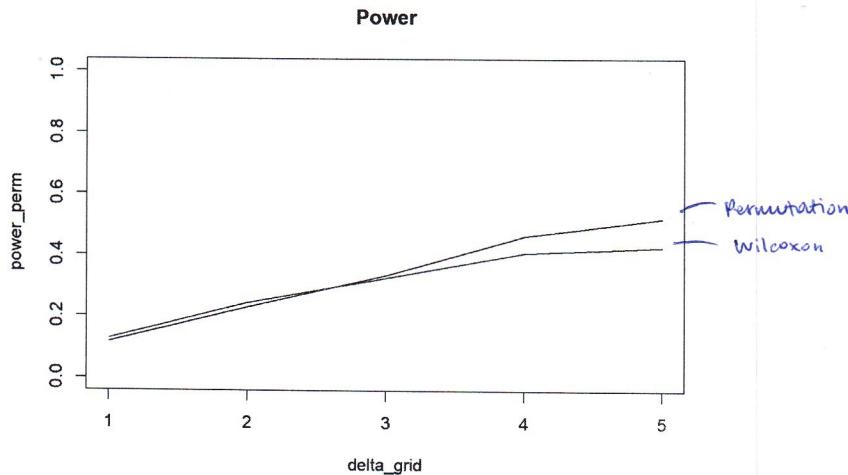
Permu. test (MEDIAN)

```
# power with a t-test structured like this is abysmal...
# Let's work on a median test...

perm_median_test=function(x,y,iter=1e4){
  T0      = abs(median(x)-median(y))
  T_stat  = numeric(iter)
  x_pooled = c(x,y)
  n       = length(x_pooled)
  n1      = length(x)
  for(perm in 1:iter){
    # permutation:
    permutation <- sample(1:n)
    x_perm     <- x_pooled[permutation]
    x1_perm   <- x_perm[1:n1]
    x2_perm   <- x_perm[(n1+1):n]
    # test statistic:
    T_stat[perm] <- abs(median(x1_perm) - median(x2_perm))
  }
  # p-value
  p_val <- sum(T_stat>=T0)/iter
  return(p_val)
}
```

```
power_perm=numeric(length(delta_grid))
set.seed(seed)
for(ii in 1:length(delta_grid)){
  p.value <- numeric(B)
  pb=progress_bar$new(total=B)
  pb$tick(0)
  delta=delta_grid[ii]
  for(j in 1:B){
    x.1 = rstable(n,.5,0)
    x.2 = rstable(n,.5,0,delta=delta)
    p.value[j] <- perm_median_test(x.1,x.2,iter=1000)
    pb$tick()
  }
  estimated.power <- sum(p.value < alpha)/B
  power_perm[ii]=estimated.power
}
```

```
plot(delta_grid,power_perm,type='l', main='Power',ylim=c(0,1))
lines(delta_grid,power_wilcox,col='red')
```



```
# what happens instead, if I run a nonparametric test when i could've used a parametric one?
```

```
# T-test vs nonparametric alternatives...
delta_grid=seq(.1,1,by=.2)
power_perm=numeric(length(delta_grid))
set.seed(seed)
for(ii in 1:length(delta_grid)){
  p.value <- numeric(B)
  pb=progress_bar$new(total=B)
  pb$tick(0)
  delta=delta_grid[ii]
  for(j in 1:B){
    x.1 = rnorm(100)
    x.2 = rnorm(100,mean=delta)
    p.value[j] = perm_t_test(x.1,x.2,iter=1000)
    pb$tick()
  }
  estimated.power <- sum(p.value < alpha)/B
  power_perm[ii]=estimated.power
}
```

Settings: normal distributions shifted by:
 δ -grid := [0.1, 0.3, 0.5, 0.7, 0.9]

Approaches:

- permutation test (mean)
- Wilcoxon test
- (standard) t-test

```

set.seed(seed)
power_wilcox=numeric(length(delta_grid))
for(ii in 1:length(delta_grid)){
  p.value <- numeric(B)
  pb=progress_bar$new(total=B)
  pb$tick(0)
  delta=delta_grid[ii]
  for(j in 1:B){
    x.1 = rnorm(100)
    x.2 = rnorm(100,mean=delta)
    p.value[j] <- wilcox.test(x.1,y=x.2,correct = T)$p.value
    pb$tick()
  }
  estimated.power <- sum(p.value < alpha)/B
  power_wilcox[ii]=estimated.power
}

```

```

set.seed(seed)
power_t=numeric(length(delta_grid))
# I can actually compute in an analytical fashion, but let's go with sim also here.
for(ii in 1:length(delta_grid)){
  p.value <- numeric(B)
  pb=progress_bar$new(total=B)
  pb$tick(0)
  delta=delta_grid[ii]
  for(j in 1:B){
    x.1 = rstable(n,.5,0)
    x.2 = rstable(n,.5,0,delta=delta)
    p.value[j] <- t.test(x.1,y=x.2)$p.value
    pb$tick()
  }
  estimated.power <- sum(p.value < alpha)/B
  power_t[ii]=estimated.power
}

```

```

set.seed(seed)
power_median=numeric(length(delta_grid))
for(ii in 1:length(delta_grid)){
  p.value <- numeric(B)
  pb=progress_bar$new(total=B)
  pb$tick(0)
  delta=delta_grid[ii]
  for(j in 1:B){
    x.1 <- rnorm(100)
    x.2 = rnorm(100,mean=delta)
    p.value[j] <- perm_median_test(x.1,x.2,iter=1000)
    pb$tick()
  }
  estimated.power <- sum(p.value < alpha)/B
  power_median[ii]=estimated.power
}

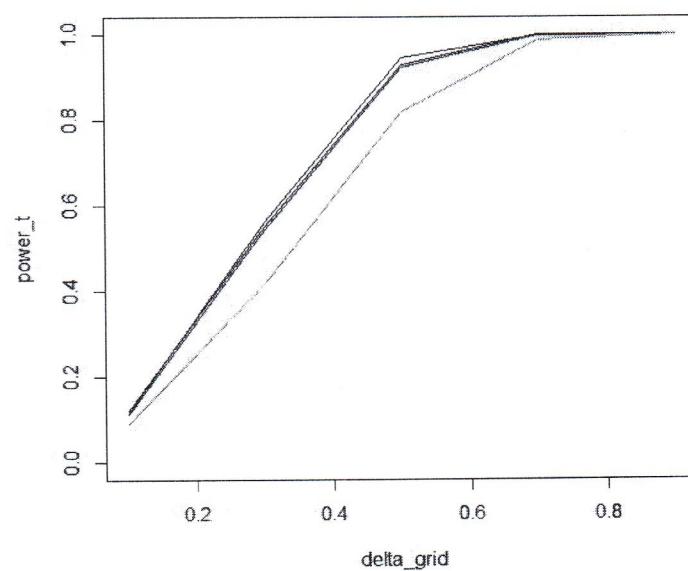
```

```

plot(delta_grid,power_t,ylim=c(0,1), main='Power, normal case',type='l')
lines(delta_grid,power_wilcox,col='red')
lines(delta_grid,power_perm, col='blue')
lines(delta_grid,power_median,col='green')

```

Power, normal case



Remember: if assumptions are met we want to use parametric settings (for the sake of computation). If the assumptions are not met then we should always go for non-parametric settings.

```

library(progress)
library(stabledist)

### -----
### HOMEWORK
### Permutation test for equality of variances
### -----
### Show that it's valid.

### -----
### H0 : var(group_1) = var(group_2)
### H1 : var(group_1) > var(group_2)
### Test stat = (s^2)_1 / (s^2)_2
### -----

```

```

### 1. Function
### -----
perm_var_test = function(x,y,iter=1e3){
  T0      = var(x)/var(y)
  T_stat  = numeric(iter)
  x_pooled = c(x,y)
  n       = length(x_pooled)
  n1      = length(x)
  for(perm in 1:iter){
    # permutation:
    permutation <- sample(1:n)
    x_perm   <- x_pooled[permutation]
    x1_perm  <- x_perm[1:n1]
    x2_perm  <- x_perm[(n1+1):n]
    # test statistic:
    T_stat[perm] <- var(x1_perm)/var(x2_perm)
  }
  # p-value
  p_val <- sum(T_stat>=T0)/iter
  return(p_val)
}

```

```

### 2. Type I error (validity)
### -----

```

```

B      = 1000
n      = 10
alpha  = 0.05
p.value = numeric(B)
set.seed(123)

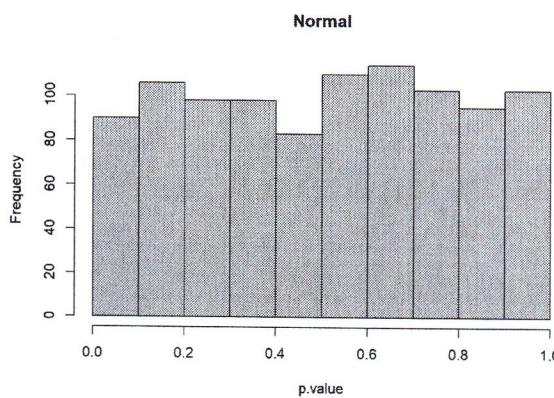
```

```

### 2.1. Gaussian
### -----
for(j in 1:B){
  x.1 = rnorm(n)
  x.2 = rnorm(n)
  p.value[j] = perm_var_test(x.1, x.2, iter=1000)
}

hist(p.value, main='Normal')

```



```

plot(ecdf(p.value), main='Normal', xlim=c(0,1))
abline(0,1, lty=2, col='red')

```

PARAMETRIC SETTINGS:

VARIANCE OF TWO POPULATIONS
(F-test, gaussian assumptions)

$$H_0: \sigma_1^2 = \sigma_2^2$$

$$H_1: \sigma_1^2 > \sigma_2^2$$

$$\sigma_1^2 < \sigma_2^2$$

$$\sigma_1^2 \neq \sigma_2^2$$

$$F = \frac{S_1^2}{S_2^2}$$

pop 1 = n_1

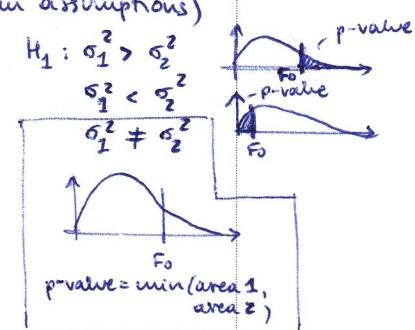
pop 2 = n_2

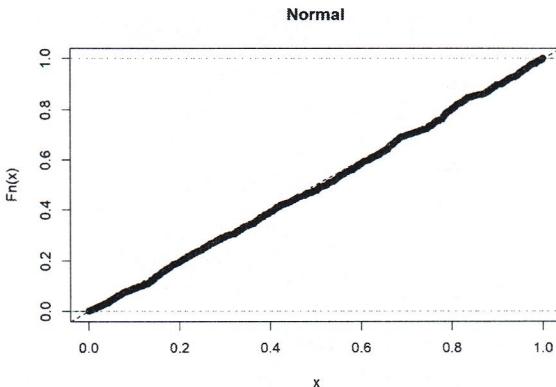
under H_0 :

$$F \sim F(n_1 - 1, n_2 - 1)$$

Hint: always call group 1 the group that we suppose to have LARGEST variance

Note: this test is very sensitive to gaussianity hypothesis violation. If H_0 are slightly violated this test goes crazy
→ non-param tests





```

estimated.alpha = sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha * (1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha * (1-estimated.alpha)/B)*qnorm(0.975))

## [1] 0.02956759 0.04200000 0.05443241

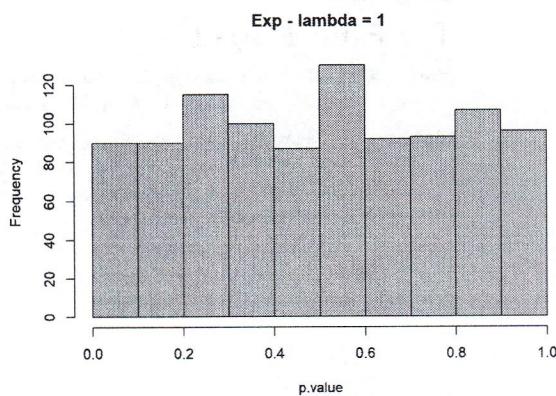
```

```

### 2.2. Exponential
### -----
for(j in 1:B){
  x.1 = rexp(n)
  x.2 = rexp(n)
  p.value[j] = perm_var_test(x.1, x.2, iter=1000)
}

hist(p.value, main='Exp - lambda = 1')

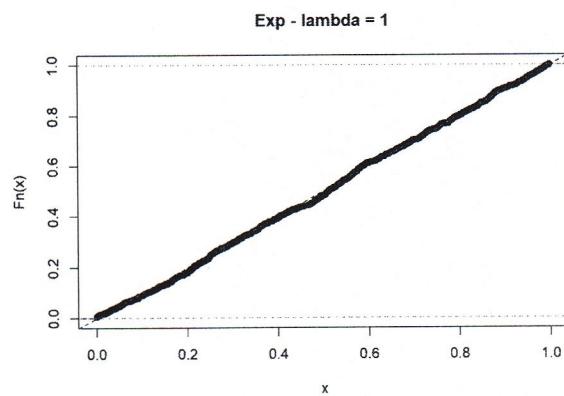
```



```

plot(ecdf(p.value), main='Exp - lambda = 1', xlim=c(0,1))
abline(0,1, lty=2, col='red')

```



```

estimated.alpha = sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha * (1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha * (1-estimated.alpha)/B)*qnorm(0.975))

## [1] 0.03301622 0.04600000 0.05898378

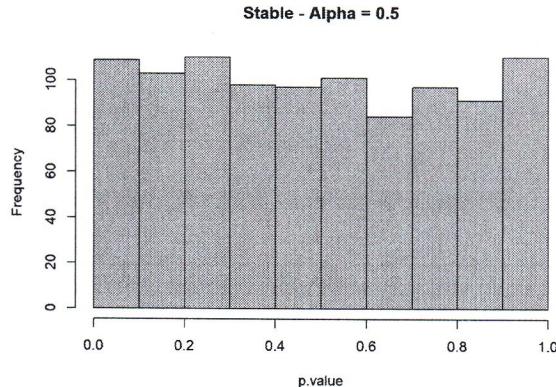
```

```

### 2.3. Stable with alpha = 0.5 (no mean, no variance)
###      (we hope to test somehow the variability of data)
### -----
for(j in 1:B){
  x.1 = rstable(n, .5, 0)
  x.2 = rstable(n, .5, 0)
  p.value[j] = perm_var_test(x.1, x.2, iter=1000)
}

hist(p.value, main='Stable - Alpha = 0.5')

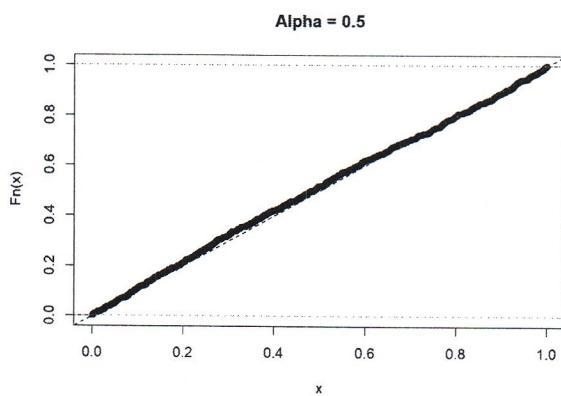
```



```

plot(ecdf(p.value), main='Alpha = 0.5', xlim=c(0,1))
abline(0,1, lty=2, col='red')

```



```

estimated.alpha = sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha * (1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha * (1-estimated.alpha)/B)*qnorm(0.975))

```

```

## [1] 0.03649188 0.05000000 0.06350812

```

```
### -----
### Correction of the assignment
### -----
# Assignment: come up with a permutation test to evaluate equality in dispersion parameters
# in a two population, independent setting.
# Let's put ourselves in a relatively simple case:
# H0: same dispersion vs. H1: dispersion(y)>dispersion(x)

library(progress)
rm(list=ls())

perm_f_test = function(x,y,iter=1e3){
  T_stat = numeric(iter)
  x_pooled = c(x,y)
  n = length(x_pooled)
  n1 = length(x)
  T0 = var(y)/var(x)
  for(perm in 1:iter){
    # permutation:
    permutation <- sample(1:n)
    x_perm <- x_pooled[permutation]
    x1_perm <- x_perm[1:n1]
    x2_perm <- x_perm[(n1+1):n]
    # test statistic:
    T_stat[perm] <- var(x2_perm)/var(x1_perm)
  }
  # p-value
  p_val <- sum(T_stat>=T0)/iter
  return(p_val)
}

perm_f_test(rnorm(10),rnorm(10, sd=10))

```

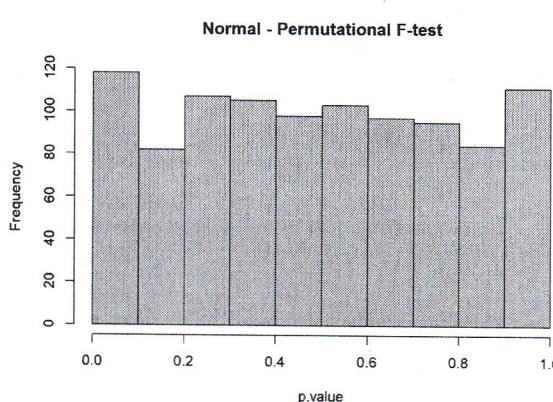
```
## [1] 0
```

```
B = 1000
seed = 1991
n = 100
alpha = 0.05
#calibration

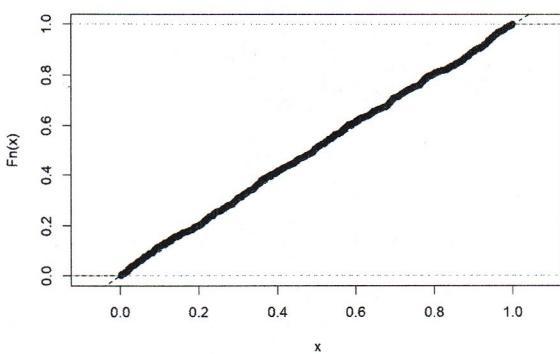
p.value = numeric(B)
pb = progress_bar$new(total=B, format = " computing [:bar] :percent eta: :eta")
pb$tick(0)
```

```
set.seed(seed)
for(j in 1:B){
  x1 = rnorm(n)
  x2 = rnorm(n)
  p.value[j] = perm_f_test(x1,x2)
  pb$tick()
}
```

```
hist(p.value, main = 'Normal - Permutational F-test')
```



```
plot(ecdf(p.value), main = 'Normal')
abline(0,1, lty=2, col='red')
```

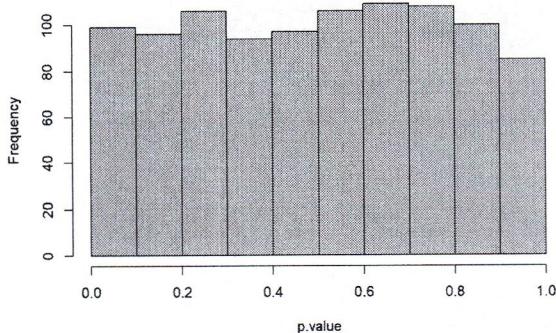
Normal

```
estimated.alpha = sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))
```

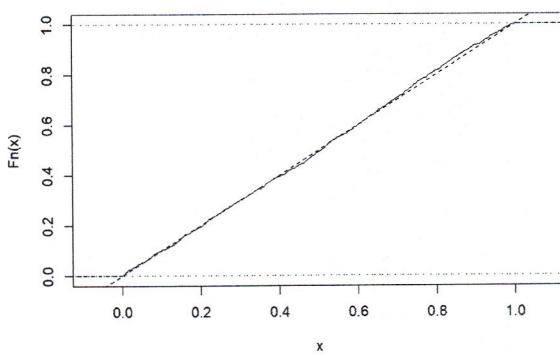
```
## [1] 0.0443961 0.0590000 0.0736039
```

```
p.value = numeric(B)
pb      = progress_bar$new(total=B)
pb$tick(0)
```

```
set.seed(seed)
for(j in 1:B){
  x1      = rnorm(n)
  x2      = rnorm(n)
  p.value[j] = var.test(x1,x2,alternative = 'greater')$p.value
  pb$tick()
}
hist(p.value, main = 'Normal - Permutational F-test')
```

Normal - Permutational F-test

```
plot(ecdf(p.value), main = 'Normal')
abline(0,1, lty=2, col='red')
```

Normal

```

estimated.alpha = sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))

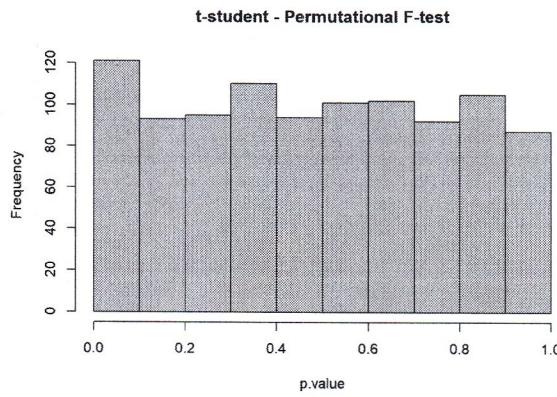
## [1] 0.03823888 0.05200000 0.06576112

p.value = numeric(B)
pb      = progress_bar$new(total=B, format = "  computing [:bar] :percent eta: :eta")
pb$tick(0)

set.seed(seed)
for(j in 1:B){
  x1      = rt(n,4)
  x2      = rt(n,4)
  p.value[j] = perm_f_test(x1,x2)
  pb$tick()
}

hist(p.value, main = 't-student - Permutational F-test')

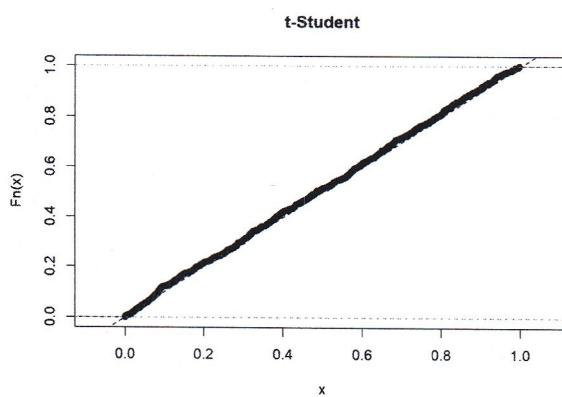
```



```

plot(ecdf(p.value), main = 't-Student')
abline(0,1, lty=2, col='red')

```



```

estimated.alpha = sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))

## [1] 0.03823888 0.05200000 0.06576112

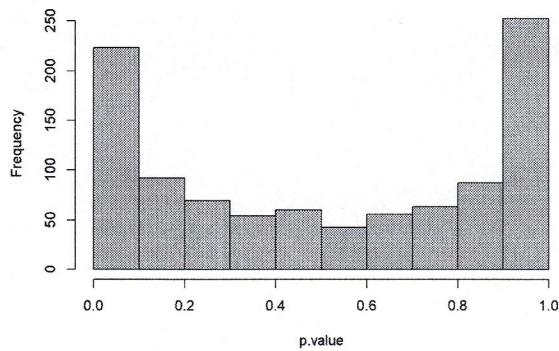
p.value = numeric(B)
pb      = progress_bar$new(total=B)
pb$tick(0)

set.seed(seed)
for(j in 1:B){
  x1      = rt(n,4)
  x2      = rt(n,4)
  p.value[j] = var.test(x1,x2,alternative = 'greater')$p.value
  pb$tick()
}

hist(p.value, main = 't-Student - Classic F-test')

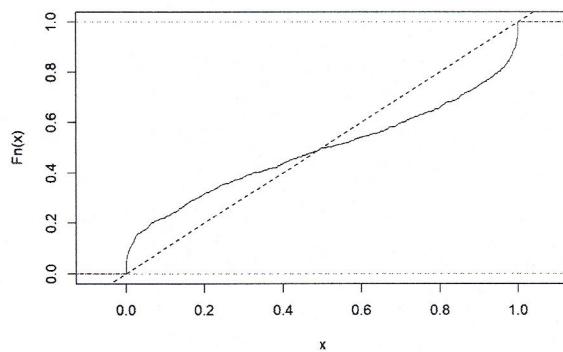
```

t-Student - Classic F-test



```
plot(ecdf(p.value), main = 't-Student - F-Test')
abline(0,1, lty=2, col='red')
```

t-Student - F-Test



```
estimated.alpha = sum(p.value < alpha)/B
c(estimated.alpha - sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975),
  estimated.alpha,
  estimated.alpha + sqrt(estimated.alpha*(1-estimated.alpha)/B)*qnorm(0.975))
```

```
## [1] 0.1514498 0.1750000 0.1985502
```

(too high!)

(a test that we would like to
perform at level 5%. is actually
performed at level ~20%.)

→ non-parametric settings!

```

### -----
### PERMUTATION TESTS: TWO INDEPENDENT POPULATIONS
### -----
## Simulated data:
# We sample n1 data from a first population X1 and n2 data from a second population X2 (n1 and n2 small)
# Test:
# H0: X1 =^d X2
# H1: X1 !=^d X2

# -----
# Case 1. H0 FALSE (we generate 2 samples from 2 different distributions → H0 false)
# To understand the behaviour of the test, we sample from two populations
# with different means

# Parameters:
n1 <- n2 <- 10
n <- n1 + n2

# Simulation:
set.seed(240279)
x1 <- runif(n1, 0, 4)
x2 <- runif(n2, 0, 4) + 3

x_pooled <- c(x1, x2)

par(mfrow=c(1,2))
boxplot(x1, x2, main='Original data')

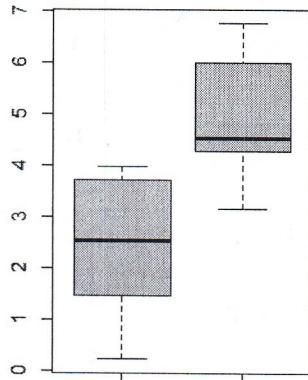
# How data change if we apply one random permutation?
permutation <- sample(1:n)

x_perm <- x_pooled[permutation]
x1_perm <- x_perm[1:n1]
x2_perm <- x_perm[(n1+1):n]

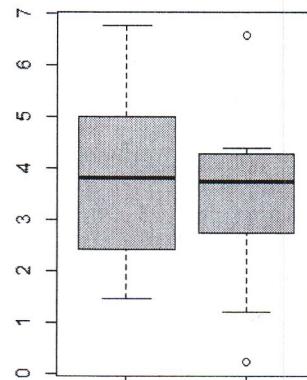
boxplot(x1_perm, x2_perm, main='Permuted data')

```

Original data



Permuted data



We have to understand if this difference between original and permuted happened by chance or it's usually like this:
if this change is remarkable (like in this case) for many permuted samples → we'll have a small p-value

As test statistic we use: $T = |\bar{x}_1 - \bar{x}_2|$

```
abs(mean(x1)-mean(x2))
```

```
## [1] 2.465564
```

(original)

```
abs(mean(x1_perm)-mean(x2_perm))
```

(permuted)

```
## [1] 0.4176959
```

The mean difference is lower.
Was that a case?

```
# TEST
# Test statistic: absolute difference between the two means
T0 <- abs(mean(x1) - mean(x2))
T0
```

```
## [1] 2.465564
```

```
# Cardinality of the permutational space:
factorial(n)
```

```
## [1] 2.432902e+18
```

```
# Number of distinct values of T*:
factorial(n)/(2*factorial(n1)*factorial(n2))
```

```
## [1] 92378
```

we can have 92K possible values
of the test statistic

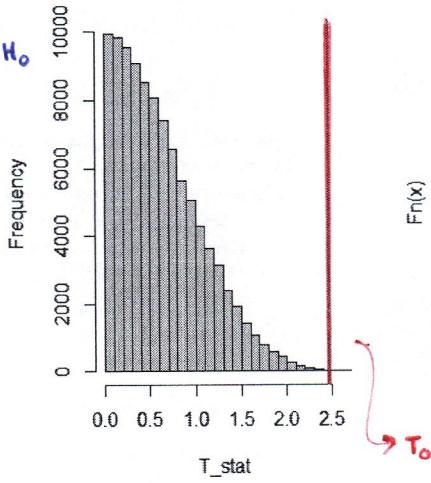
```
# Minimum achievable p-value:  
1/(factorial(n)/(2*factorial(n1)*factorial(n2)))
```

```
## [1] 1.082509e-05
```

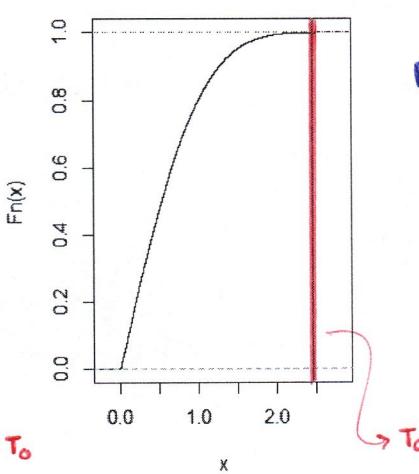
Conditional
Monte Carlo
estimate
(which randomly
permutes)

```
# CMC to estimate the p-value  
B <- 100000 # Number of permutations  
T_stat <- numeric(B) # Vector where we will store the values of T*  
  
# To estimate the p-value we use a Loop  
# Inside the Loop, we do the following:  
# 1. choose a random permutation of data (with the command sample)  
# 2. calculate and save the test statistic obtained with the permuted data  
for(perm in 1:B){  
  # permutation:  
  permutation <- sample(1:n)  
  x_perm <- x_pooled[permutation]  
  x1_perm <- x_perm[1:n1]  
  x2_perm <- x_perm[(n1+1):n]  
  # test statistic:  
  T_stat[perm] <- abs(mean(x1_perm) - mean(x2_perm))  
}  
  
# Permutational distribution of T  
hist(T_stat,xlim=range(c(T_stat,T0)),breaks=30)  
abline(v=T0,col=3,lwd=2)  
  
plot(ecdf(T_stat))  
abline(v=T0,col=3,lwd=2)
```

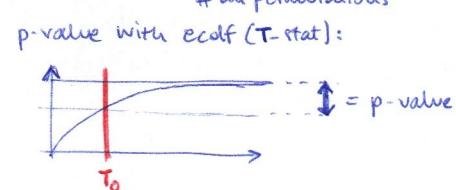
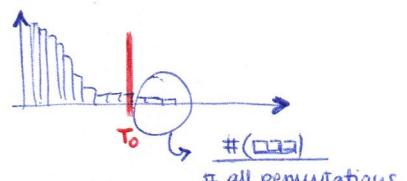
Histogram of T_stat



ecdf(T_stat)



p-value = proportion of permutations that had a T-statistic value $\geq T_0$



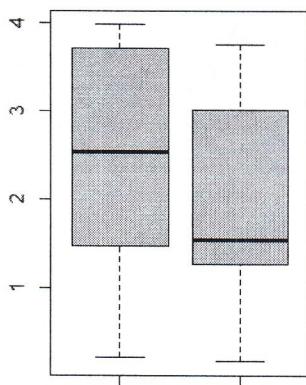
```
# p-value  
p_val <- sum(T_stat >= T0)/B  
p_val
```

```
## [1] 3e-04  $\Rightarrow$  we reject  $H_0$ 
```

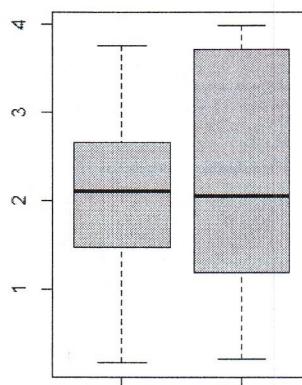
```
# -----  
# Case 2.  $H_0$  TRUE  
# now we simulate data from two populations with the same distribution
```

```
# Simulation:  
set.seed(240279)  
x1 <- runif(n1, 0, 4)  
x2 <- runif(n2, 0, 4) + 0  
x_pooled <- c(x1, x2)  
  
par(mfrow=c(1,2))  
boxplot(x1, x2, main='Original data')  
  
# How data change if we apply one random permutation?  
permutation <- sample(1:n)  
  
x_perm <- x_pooled[permutation]  
x1_perm <- x_perm[1:n1]  
x2_perm <- x_perm[(n1+1):n]  
  
boxplot(x1_perm, x2_perm, main='Permuted data')
```

Original data



Permuted data



Here it doesn't seem like the permutation influenced a lot.
let's check on other permutations!

```

abs(mean(x1)-mean(x2))

## [1] 0.5344363          (original)

abs(mean(x1_perm)-mean(x2_perm))

## [1] 0.1823041          (permuted)

# TEST
# Test statistic: absolute difference between the two means
T0 <- abs(mean(x1) - mean(x2))
T0

## [1] 0.5344363

# Cardinality of the permutational space:
factorial(n)

## [1] 2.432902e+18

# Number of distinct values of T*:
factorial(n)/(2*factorial(n1)*factorial(n2))

## [1] 92378

# Minimum achievable p-value:
1/(factorial(n)/(2*factorial(n1)*factorial(n2)))

## [1] 1.082509e-05

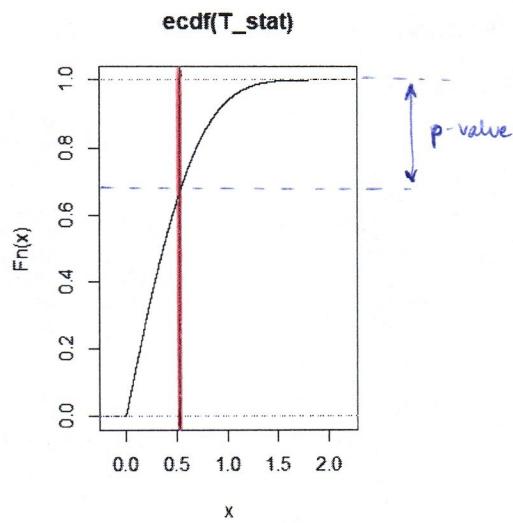
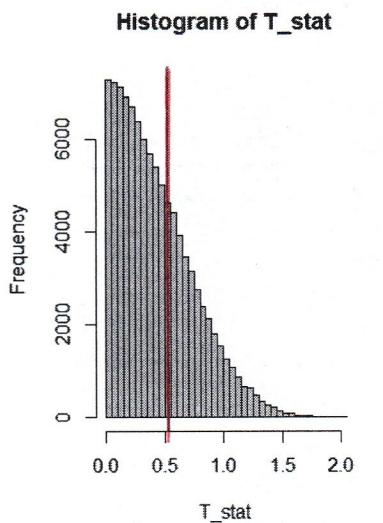
# CMC to estimate the p-value
B <- 100000 # Number of permutations
T_stat <- numeric(B) # Vector where we will store the values of T*

for(perm in 1:B){
  # permutation:
  permutation <- sample(1:n)
  x_perm <- x_pooled[permutation]
  x1_perm <- x_perm[1:n1]
  x2_perm <- x_perm[(n1+1):n]
  # test statistic:
  T_stat[perm] <- abs(mean(x1_perm) - mean(x2_perm))
}

# Permutational distribution of T
hist(T_stat,xlim=range(c(T_stat,T0)),breaks=30)
abline(v=T0,col=3,lwd=2)

plot(ecdf(T_stat))
abline(v=T0,col=3,lwd=2)

```



```
# p-value
p_val <- sum(T_stat >= T0)/B
p_val
```

```
## [1] 0.33116 → we do not reject H0
```

```
### -----
### -----
### PERMUTATION TESTS (MULTIVARIATE)
### -----
### Example 1: Two (independent) multivariate population test
# Hourly accesses to AreaC: working days vs week-end days

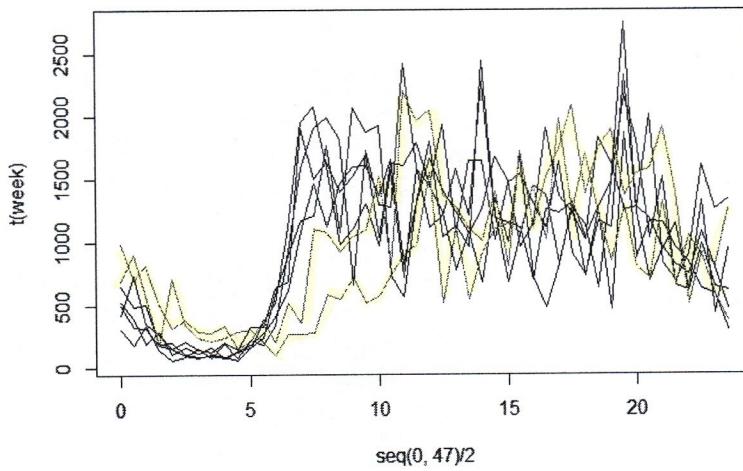
d1 <- read.csv('accessi-orari-areac-2016-09-12-00_00_00.csv', header=T)
d2 <- read.csv('accessi-orari-areac-2016-09-13-00_00_00.csv', header=T)
d3 <- read.csv('accessi-orari-areac-2016-09-14-00_00_00.csv', header=T)
d4 <- read.csv('accessi-orari-areac-2016-09-15-00_00_00.csv', header=T)
d5 <- read.csv('accessi-orari-areac-2016-09-16-00_00_00.csv', header=T)
d6 <- read.csv('accessi-orari-areac-2016-09-17-00_00_00.csv', header=T)
d7 <- read.csv('accessi-orari-areac-2016-09-18-00_00_00.csv', header=T)

week <- rbind(d1[,2], d2[,2], d3[,2], d4[,2], d5[,2], d6[,2], d7[,2])
matplot(seq(0,47)/2, t(week), type='l', col=c(1,1,1,1,2,2), lty=1)
```

Note: permutations and reflections are made on the entire unit. We permute two rows; we reflect one row; we don't work on singular column.

Two datasets: $t_1: 5 \times 48$ (working days)
 $t_2: 2 \times 48$ (weekends)

Each column contains the number of cars entering the area C of Milan in a given half an hour in a given day. Each day is made by 48 half hours. Each row corresponds to one day.



We define the two groups:

```
t1 <- week[1:5,]
t2 <- week[6:7,]

# Computing a proper test statistic
# (i.e., squared distance between the two sample mean vectors)
t1.mean <- colMeans(t1)
t2.mean <- colMeans(t2)

n1 <- dim(t1)[1]
n2 <- dim(t2)[1]
n <- n1 + n2

T20 <- as.numeric((t1.mean-t2.mean) %*% (t1.mean-t2.mean))
T20
```

```
## [1] 8976210
```

```
# Selecting a proper permutation strategy (i.e., data point permutations)
```

```
# number of possible data point permutations  
factorial(7)
```

```
## [1] 5040
```

```
# number of different values of the test statistic  
choose(7,5)
```

: How many groups of 5 and 2 can we do? We can align the 7 days in 21 possible ways \Rightarrow the permutational distribution will be a discrete uniform with just 21 values

```
## [1] 21
```

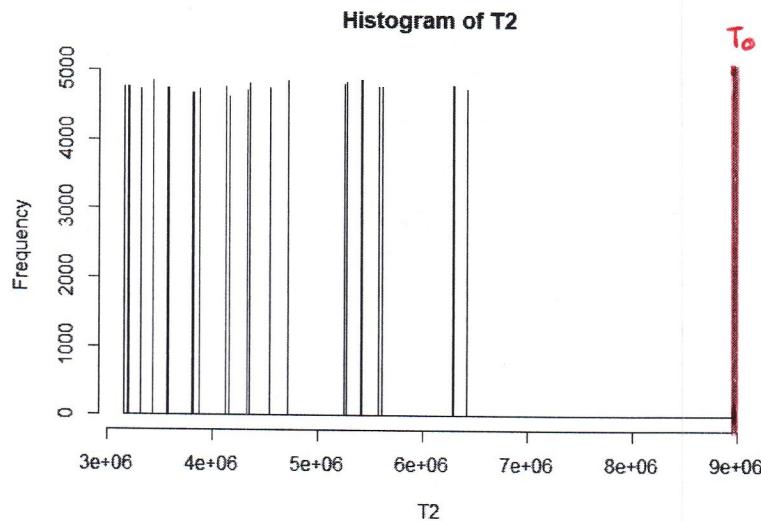
```
# Estimating the permutational distribution under H0
```

```
B <- 100000  
T2 <- numeric(B)
```

```
for(perm in 1:B){  
  # Random permutation of indexes  
  # When we apply permutations in a multivariate case, we keep the units together  
  # i.e., we only permute the rows of the data matrix  
  t_pooled <- rbind(t1,t2)  
  permutation <- sample(n)  
  t_perm <- t_pooled[permutation,]  
  t1_perm <- t_perm[1:n1,]  
  t2_perm <- t_perm[(n1+1):n,]  
  
  # Evaluation of the test statistic on permuted data  
  t1.mean_perm <- colMeans(t1_perm)  
  t2.mean_perm <- colMeans(t2_perm)  
  T2[perm] <- (t1.mean_perm-t2.mean_perm) %*% (t1.mean_perm-t2.mean_perm)  
}
```

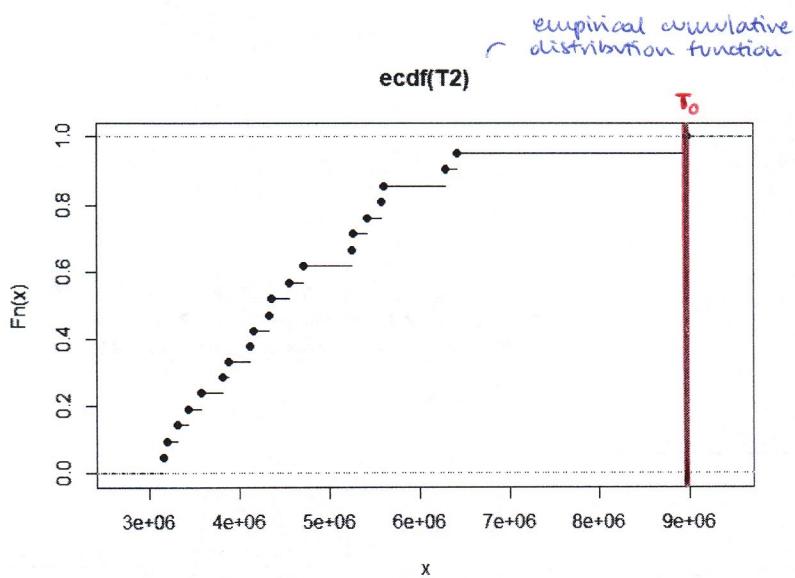
```
# plotting the permutational distribution under H0  
hist(T2,xlim=range(c(T2,T20)),breaks=1000)  
abline(v=T20,col=3,lwd=4)
```

Remember that we're permuting the days (the vectors), we're not destroying vectors! (permutations over the rows, not the columns)



← 21 test statistics
(of 21 possible test statistics
our one is the most extreme)

```
plot(ecdf(T2))  
abline(v=T20,col=3,lwd=4)
```



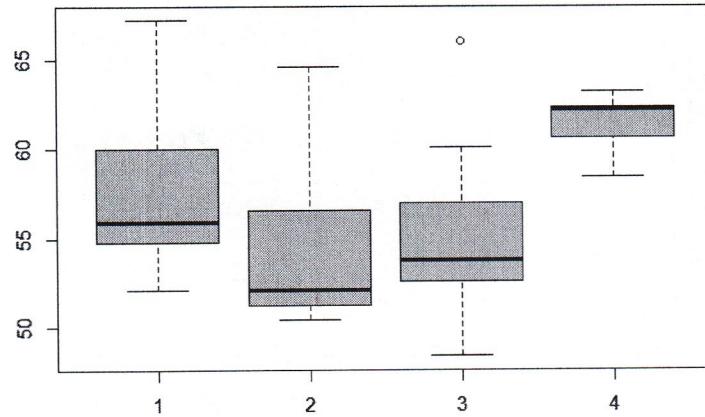
```
# p-value
p_val <- sum(T2>=T20)/B
p_val
```

```
## [1] 0.04868 ← smallest p-value achievable (we reject H0)
```

```
# 
# Example 2: Center of symmetry of one multivariate population
# Relative humidity in Milan during the summer months

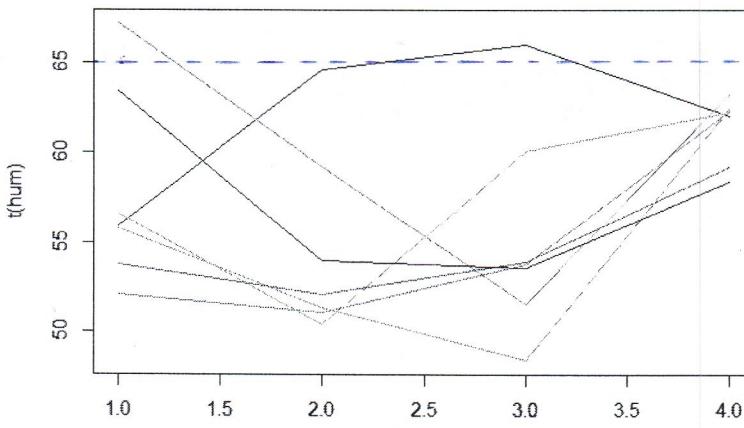
hum <- read.csv2('307_Umidita_relativa_2008_2014.csv', header=T)
hum <- hum[,3]
hum <- matrix(hum, ncol=12, byrow=T)[,6:9]

boxplot(hum)
```



```
matplot(t(hum), type='l', lty=1)
```

7 years (each line is a year) and 4 features
 (average relative humidity in the 1st, 2nd, 3rd and 4th
 trimester of the year)



$$H_0: \text{center} = [65, 65, 65, 65]^T$$

Basically we want to check:
 if we make a reflection of the lines
 w.r.t. the horizontal line (center
 $= [65, 65, 65, 65]^T$), will we obtain
 a dataset with exactly the same
 likelihood?

looking at the plot we're expecting
 strong evidence to reject.
 This time, for each curve we can decide
 to reflect it or not
 $\Rightarrow Z^N = Z^7 = 128$ possible transforms.
 (notice that this time
 are not permutations)

```
# center of symmetry under H0
mu0 <- c(65, 65, 65, 65)

# Computing a proper test statistic
# (i.e., squared distance between the sample mean vector and the hypothesized center of symmetry)
x.mean <- colMeans(hum)
n <- dim(hum)[1]
p <- dim(hum)[2]

T20 <- as.numeric((x.mean-mu0) %*% (x.mean-mu0) )

# Selecting a proper Likelihood-invariant strategy (i.e., data point reflections)
# We are assuming that under H0 the data distribution is symmetric

# number of possible data point reflections
2^7
```

[1] 128

number of different values of the test statistic
 $2^{7/2}$

} because if we reflect them all we would obtain a value
 test statistic equal to the original one

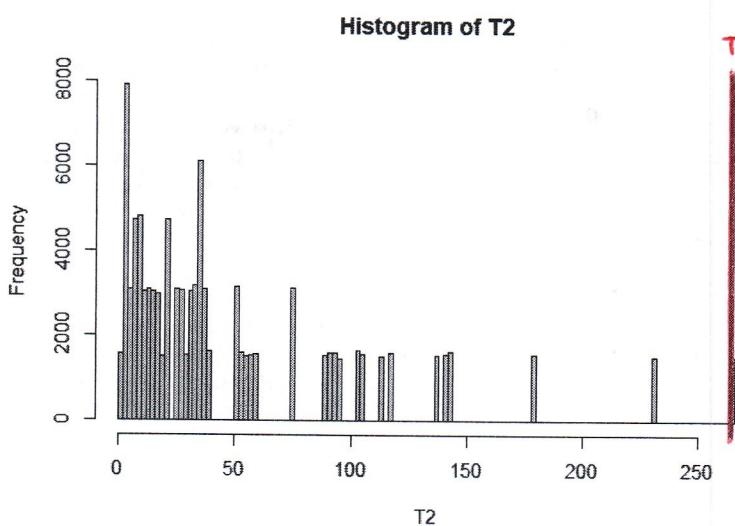
[1] 64

```
# Estimating the permutational distribution under H0
B <- 100000
T2 <- numeric(B)

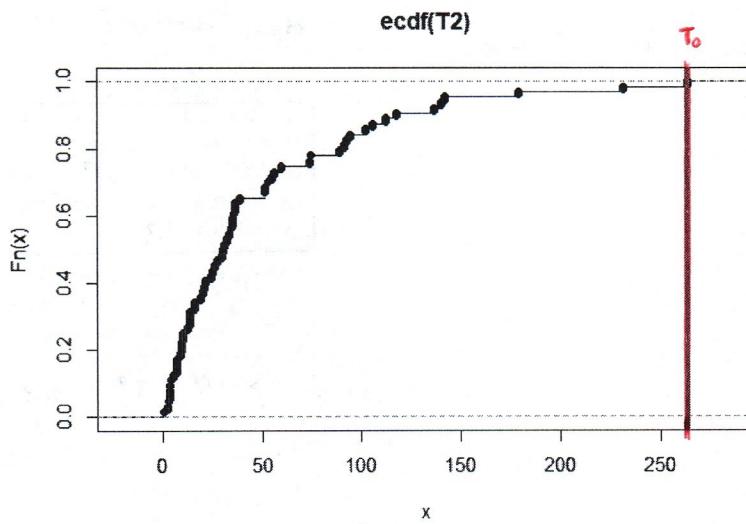
for(perm in 1:B){
  # In this case we use changes of signs in place of permutations

  # Permuted dataset
  signs.perm <- rbinom(n, 1, 0.5)*2 - 1
  hum_perm <- mu0 + (hum - mu0) * matrix(signs.perm, nrow=n, ncol=p, byrow=FALSE)
  x.mean_perm <- colMeans(hum_perm)
  T2[perm] <- (x.mean_perm-mu0) %*% (x.mean_perm-mu0)
}

# plotting the permutational distribution under H0
hist(T2, xlim=range(c(T2, T20)), breaks=100)
abline(v=T20, col=3, lwd=4)
```



```
plot(ecdf(T2))
abline(v=T20, col=3, lwd=4)
```



```
# p-value
p_val <- sum(T2 >= T20) / B
p_val
```

```
## [1] 0.01535 → we reject H0
```

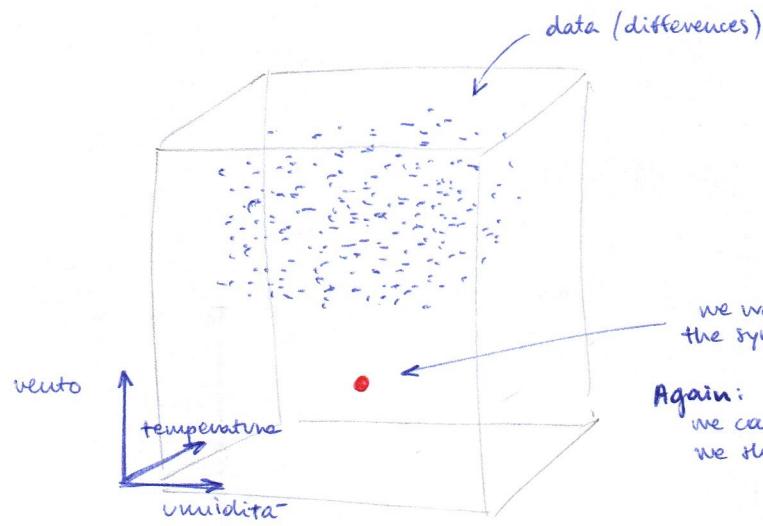
```
# Example 3: two paired multivariate population test
# (i.e., same days for Milan and Barcelona)
# The data set contains observations of temperature, humidity and wind in Milan and Barcelona
# on 50 different days
```

```
t1 <- read.table('barcellona.txt', header=T)
t2 <- read.table('milano.txt', header=T)
```

```
library(rgl)
open3d()
```

```
plot3d(t1-t2, size=3, col='orange', aspect = F)
points3d(0,0,0, size=6)
a = scene3d()
rgl.close()
x11()
rglwidget(a)
```

kind of same as before.
when we have to work with
paired populations, we actually
work with one population:
the population of differences



we want to check if this is
the symmetry of the distribution

Again: if it's the center of symmetry
we can reflect the data and
we shouldn't notice the difference

```

p <- dim(t1)[2]
n1 <- dim(t1)[1]
n2 <- dim(t2)[1]
n <- n1+n2

# Evaluate the test statistic
t1.mean <- colMeans(t1)
t2.mean <- colMeans(t2)
t1.cov <- cov(t1)
t2.cov <- cov(t2)
Sp <- ((n1-1)*t1.cov + (n2-1)*t2.cov)/(n1+n2-2)
Sinv <- solve(Sp)

delta.0 <- c(0,0,0)

diff <- t1-t2
diff.mean <- colMeans(diff)
diff.cov <- cov(diff)
diff.inv cov <- solve(diff.cov)

#T20 <- as.numeric(n1 * (diff.mean-delta.0) %*% (diff.mean-delta.0))
#T20 <- as.numeric(n1 * (diff.mean-delta.0) %*% solve(diag(diag(diff.cov))) %*% (diff.mean-delta.0))
T20 <- as.numeric(n1 * (diff.mean-delta.0) %*% diff.inv cov %*% (diff.mean-delta.0))

# With coupled data, we have to reduce the number of permutations
# the only permutations that preserve the Likelihood under H0 are
# permutations within each couple (i.e., exchange Milan and Barcelona within the same day)
# We are assuming the differences are simmetrically distributed under H0

# number of possible data point reflections
2^50

```

In this case we use the Hotelling T^2 statistic
(the sample size allows us to: $p=3, n=50$)

other possibilities

squared euclidean distance

here we're neglecting the covariance structure, we're rescaling each component by its standard deviation

here we're rescaling using also the variance/covariance structure (Mahalanobis distance² of the sample means to the red point)

[1] 1.1259e+15 \Rightarrow need of Monte Carlo

```

# Estimating the permutational distribution under H0
B <- 10000
T2 <- numeric(B)

for(perm in 1:B)
{
  # Random permutation
  # obs: exchanging data within couples means changing the sign of the difference
  signs.perm <- rbinom(n1, 1, 0.5)^2 - 1

  diff_perm <- diff * matrix(signs.perm, nrow=n1, ncol=p, byrow=FALSE)
  diff.mean_perm <- colMeans(diff_perm)
  diff.cov_perm <- cov(diff_perm)
  diff.inv cov_perm <- solve(diff.cov_perm)

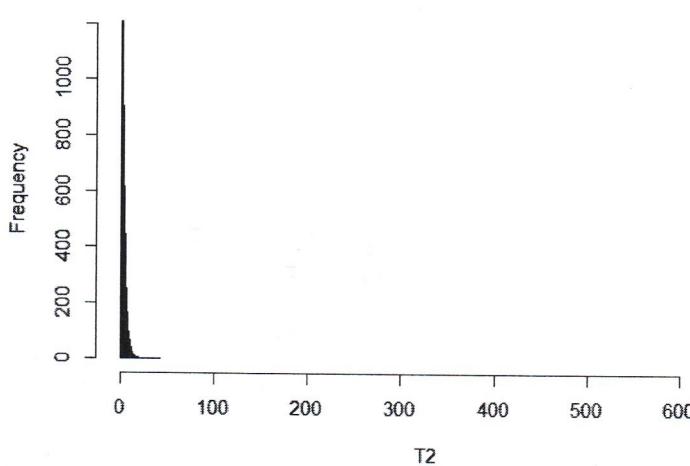
  #T2[perm] <- as.numeric(n1 * (diff.mean_perm-delta.0) %*% (diff.mean_perm-delta.0))
  #T2[perm] <- as.numeric(n1 * (diff.mean_perm-delta.0) %*% solve(diag(diag(diff.cov_perm))) %*% (diff.mean_perm-delta.0))
  T2[perm] <- as.numeric(n1 * (diff.mean_perm-delta.0) %*% diff.inv cov_perm %*% (diff.mean_perm-delta.0))
}

# plotting the permutational distribution under H0
hist(T2,xlim=range(c(T2,T20)),breaks=100)
abline(v=T20,col=3,lwd=4)

```

instead of randomly mixing data
we're randomly picking up $-1/+1$
we just multiply by $-1/+1$

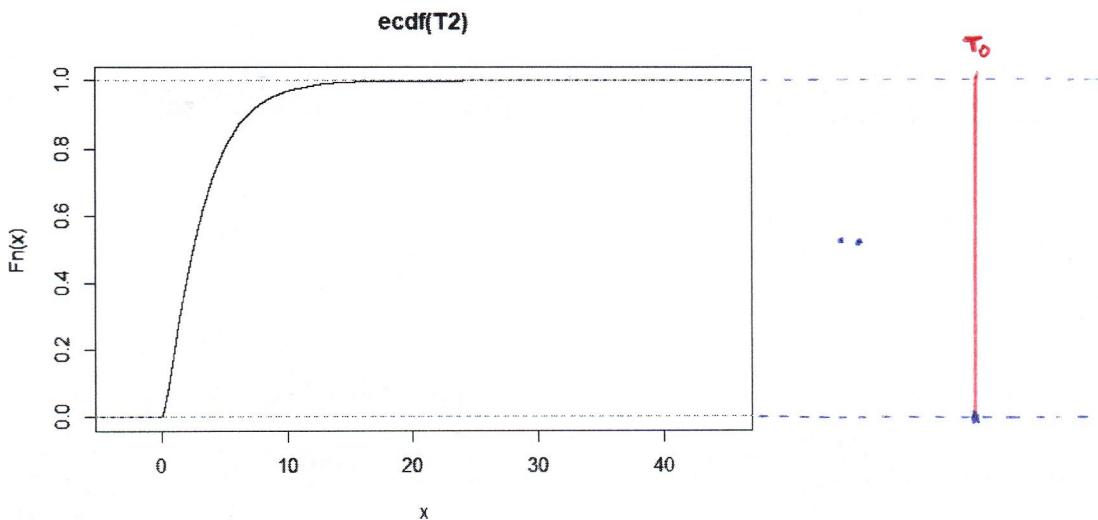
Histogram of T2



```

plot(ecdf(T2))
abline(v=T20,col=3,lwd=4)

```



```
# p-value
p_val <- sum(T2>=T20)/B
p_val
```

```
## [1] 0
```

```
2^50
```

```
## [1] 1.1259e+15
```

```
###  
###  
### PERMUTAZION TESTS: ANOVA  
###  
###
```

```
# One-way ANOVA  
# (p=1, g=6)  
head(chickwts)
```

The ANOVA is based on the F-test. F-test tends to be less robust w.r.t. the violation of normality assumptions (if normality is violated then we have a big problem \Rightarrow non-parametric settings)

generalization of the two population test

```
## weight feed
## 1 179 horsebean
## 2 160 horsebean
## 3 136 horsebean
## 4 227 horsebean
## 5 217 horsebean
## 6 168 horsebean
```

```
attach(chickwts)
summary(chickwts)
```

```
## weight feed
## Min. :108.0 casein :12
## 1st Qu.:204.5 horsebean:10
## Median :258.0 linseed :12
## Mean :261.3 meatmeal :11
## 3rd Qu.:323.5 soybean :14
## Max. :423.0 sunflower:12
```

Is there effect based on the feed?

```
g <- nlevels(feed)
n <- dim(chickwts)[1]

layout(cbind(1,2))
plot(feed, weight, xlab='treat', col=rainbow(g), main='Original Data')
```

$H_0: \tau_1 = \dots = \tau_6 = 0$: the 6 groups share the same distribution
 $H_1: \exists \tau_j \neq 0 \quad j=1, \dots, 6$

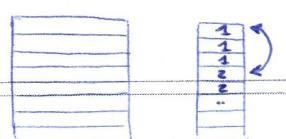
the chickens belong to the same population

H1: (H_0)^c
the chickens belong to several different populations
Parametric test:

fit <- aov(weight ~ feed) *just because we want to use the F statistic of the parametric test*

```
##          Df Sum Sq Mean Sq F value    Pr(>F)
## feed      5 231129  46226   15.37 5.94e-10 ***
## Residuals 65 195556   3009
## ...
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Under H_0 all the data come from the same distribution, so if we shuffle the observations and permute the elements it shouldn't be noticed :



we randomly change some labels

```

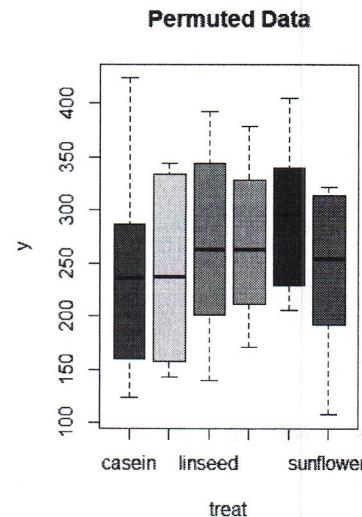
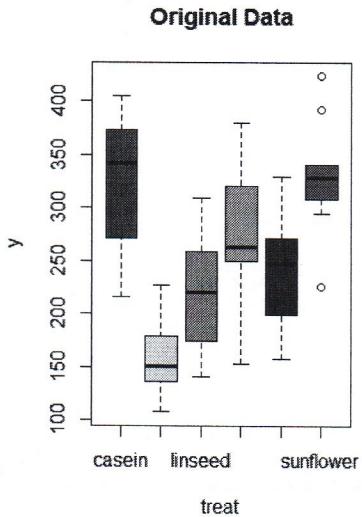
# Permutation test:
# Test statistic: F stat
T0 <- summary(fit)[[1]][1,4]
T0

## [1] 15.3648

# what happens if we permute the data?
permutazione <- sample(1:n)
weight_perm <- weight[permutazione]
fit_perm <- aov(weight_perm ~ feed)
summary(fit_perm)

##           Df Sum Sq Mean Sq F value Pr(>F)
## feed      5 34466   6893   1.142 0.347
## Residuals 65 392219   6034
plot(feed, weight_perm, xlab='treat', col=rainbow(g), main='Permuted Data')

```



Visually the two are very different,
let's make more permutations

```

# CMC to estimate the p-value
B <- 1000 # Number of permutations
T_stat <- numeric(B)
n <- dim(chickwts)[1]

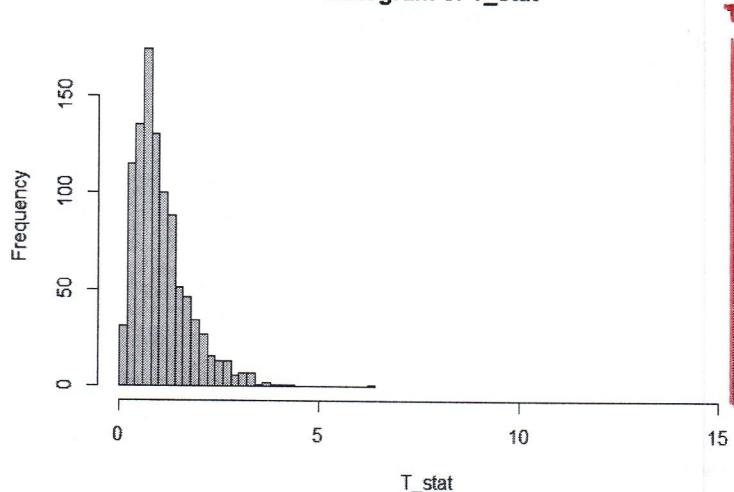
for(perm in 1:B){
  # Permutation:
  permutation <- sample(1:n)
  weight_perm <- weight[permutation]
  fit_perm <- aov(weight_perm ~ feed)

  # Test statistic:
  T_stat[perm] <- summary(fit_perm)[[1]][1,4]
}

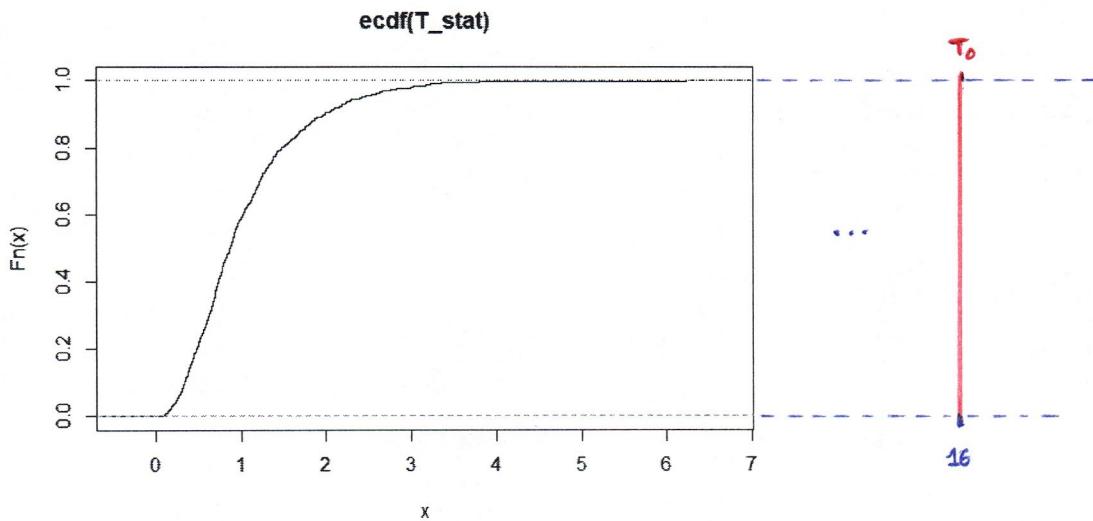
layout(1)
hist(T_stat, xlim=range(c(T_stat,T0)), breaks=30)
abline(v=T0, col=3, lwd=2)

```

Histogram of T_stat



```
plot(ecdf(T_stat))
abline(v=T20,col=3,lwd=4)
```



```
# p-value
p_val <- sum(T_stat>=T0)/B
p_val
```

[1] 0 \Rightarrow we reject H_0

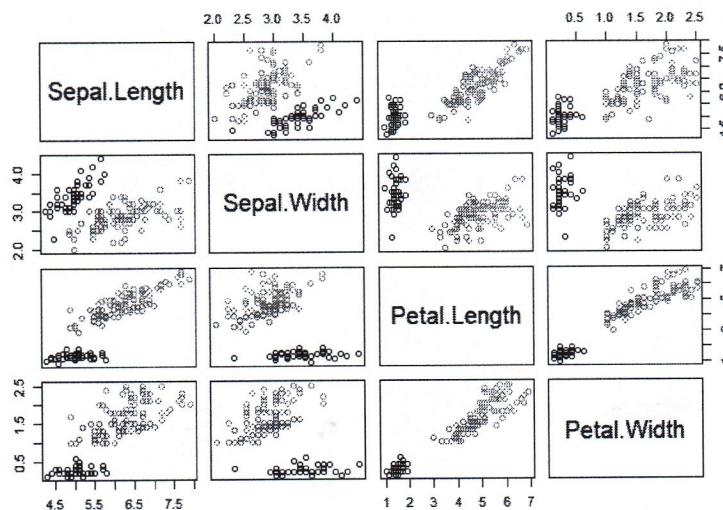
we reject the null hypothesis

```
detach(chickwts)
```

#

```
# MANOVA (MULTIVARIATE ONE-WAY ANOVA)
data(iris)
attach(iris)
species.name <- factor(Species, labels=c('setosa','versicolor','virginica'))
iris4 <- iris[,1:4]
plot(iris4,col=species.name)
```

(same code as one way ANOVA
but with a different statistic,
moreover, same way of reasoning)



```
i1 <- which(species.name=='setosa')
i2 <- which(species.name=='versicolor')
i3 <- which(species.name=='virginica')
n1 <- length(i1)
n2 <- length(i2)
n3 <- length(i3)
n <- n1+n2+n3

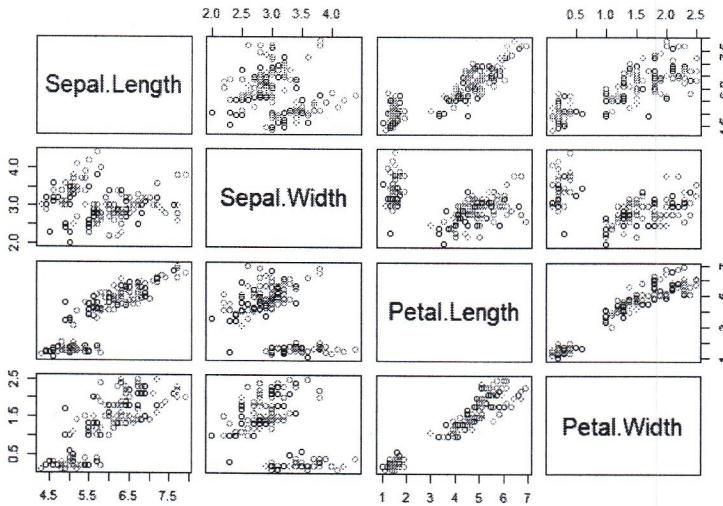
g <- length(levels(species.name))
p <- 4
detach(iris)

# MANOVA
# parametric test:
fit <- manova(as.matrix(iris4) ~ species.name)
summary.manova(fit,test="Wilks")
```

We use this as test statistic

```
##          Df    Wilks approx F num Df den Df Pr(>F)
## species.name  2  0.023439  199.15     8    288 < 2.2e-16 ***
## Residuals   147
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# How to perform a permutation test in this case?
# Multivariate framework -> We permute the Labels associated to each unit
permutation <- sample(1:n)
species.name.perm <- species.name[permutation]
plot(iris4,col=species.name.perm)
```



```
fit.perm <- manova(as.matrix(iris4) ~ species.name.perm)
summary.manova(fit.perm,test="Wilks")
```

```
##          Df    Wilks approx F num Df den Df Pr(>F)
## species.name.perm  2  0.94786  0.97683     8    288 0.4542
## Residuals       147
```

```
# TEST
# Test statistics: Wilks Lambda
T0 <- -summary.manova(fit.perm,test="Wilks")$stats[1,2]
T0
```

```
## [1] -0.02343863
```

```
# Note that wilk's Lambda is significant for small values!
# It is sufficient to change its sign to use it in a permutation test
```

high Wilks → acceptance of H_0

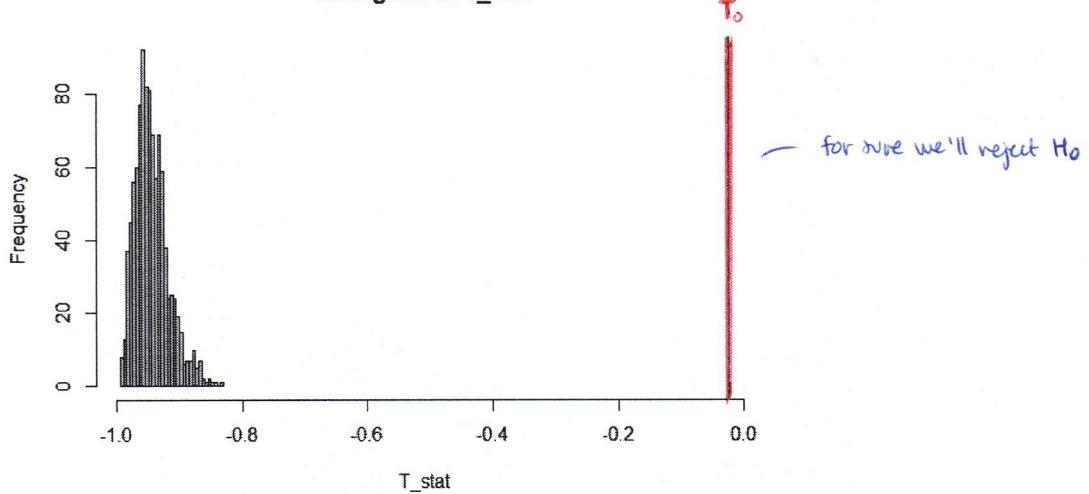
```
# Permutations
B <- 1000
T_stat <- numeric(B)

for(perm in 1:B){
  # choose random permutation
  permutation <- sample(1:n)
  species.name.perm <- species.name[permutation]
  fit.perm <- manova(as.matrix(iris4) ~ species.name.perm)
  T_stat[perm] <- -summary.manova(fit.perm,test="Wilks")$stats[1,2]
}

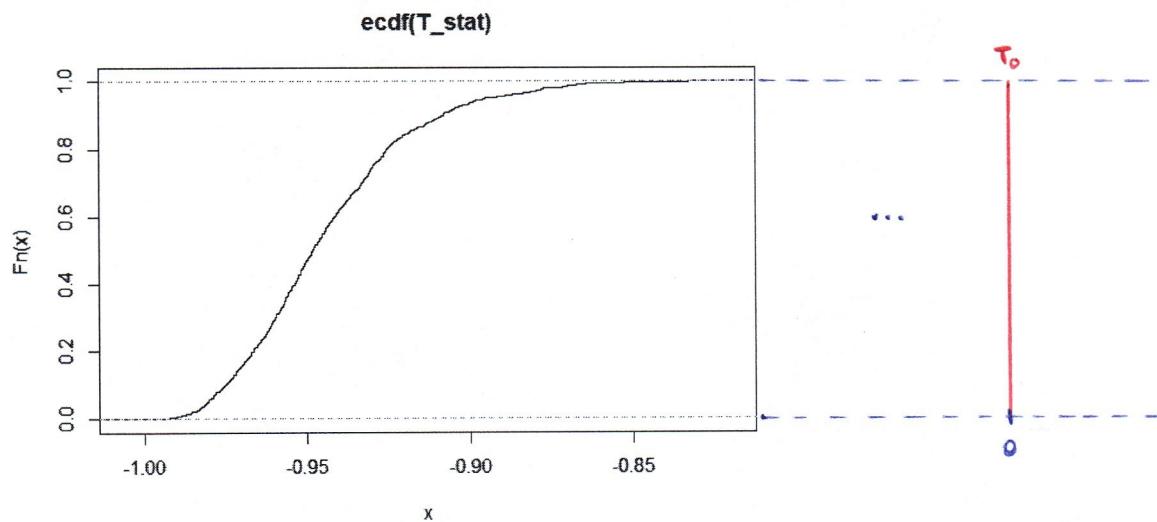
layout(1)
hist(T_stat,xlim=range(c(T_stat,T0)),breaks=30)
abline(v=T0,col=3,lwd=2)
```

→ we put a "—" , thanks to that
we reject for large Wilks values
(as in standard procedures (with
other test statistics))

Histogram of T_stat



```
plot(ecdf(T_stat))
abline(v=T0, col=3, lwd=4)
```



```
# p-value
p_val <- sum(T_stat>=T0)/B
p_val
```

```
## [1] 0
```

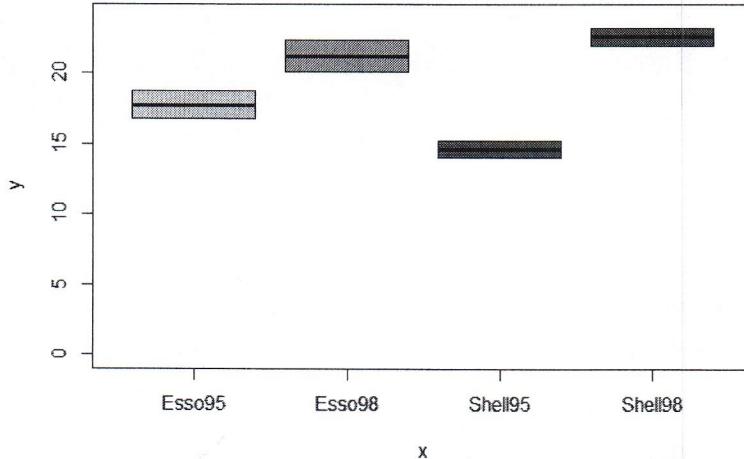
```
# 
# Two-way ANOVA (Pb. 4, 14/09/06)
# Variable: Distance [km/L]
# factor1: Fuel station (0=Esso, 1=Shell)
# factor2: Type of fuel (0=95, 1=98)
```

unidimensional response variable : fuel efficiency (km)
(based/influenced by 2 factor)

```
km      <- c(18.7, 16.8, 20.1, 22.4, 14.0, 15.2, 22.0, 23.3)
station <- factor(c('Esso','Esso','Esso','Shell','Shell','Shell','Shell'))
fuel    <- factor(c('95','95','98','98','95','95','98','98'))
station_fuel<- factor(c('Esso95','Esso95','Esso98','Esso98','Shell95','Shell95','Shell98','Shell98'))
M       <- mean(km)
Mstation <- tapply(km, station, mean)
Mfuel   <- tapply(km, fuel, mean)
Mstation_fuel <- tapply(km, station_fuel, mean)
```

```
plot(station_fuel, km, col=rainbow(5)[2:5], ylim=c(0,24))
```

← interaction between
the two factors



```
# Parametric test:
summary.aov(km ~ station + fuel + station:fuel)
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## station     1   1.53    1.53  1.018 0.37001
## fuel        1  66.70   66.70 44.357 0.00264 **
## station:fuel 1 10.35   10.35  6.884 0.05857 .
## Residuals    4   6.01    1.50
## ...
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

← we remove the interaction term

```
# Without interaction
summary.aov(km ~ station + fuel))
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## station     1   1.53    1.53  0.468 0.52440
## fuel        1  66.70   66.70 20.378 0.00632 **
## Residuals    5 16.37   3.27
## ...
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

← we remove the station

```
# Without station
summary.aov(km ~ fuel))
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## fuel        1  66.7   66.70 22.36 0.00323 **
## Residuals    6 17.9    2.98
## ...
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

try to obtain the same with a non parametric approach

• # Permutation test

```
# The model we have to test is the following:
# km = mu + alpha*station + beta*fuel + gamma*station*fuel
# We have 3 different tests:
# 1. factor station (H0: alpha=0)
# 2. factor fuel (H0: beta=0)
# 3. interaction (H0: gamma=0)
```

We use ANOVA-equivalent of Friedman-Levin method (explained later with linear models)

- fit reduced model
- compute the residuals
- permute the residuals
- summing permuted res with the fitted values
- fit the full model

(1)

```
# We apply different permutations for developing the different tests!
# We start by testing the interaction:
# H0: gamma=0 against H_1: gamma!=0

# test statistic:
summary.aov(aov(km ~ station + fuel + station:fuel))
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## station     1   1.53    1.53  1.018 0.37001
## fuel        1  66.70   66.70 44.357 0.00264 **
## station:fuel 1 10.35   10.35  6.884 0.05857 .
## Residuals    4   6.01    1.50
## ...
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
T0_station_fuel <- summary.aov(aov(km ~ station + fuel + station:fuel))[[1]][3,4]
T0_station_fuel
```

the residuals are permuted under H₀ (under H₀ we can permute them and have no effect (they're likelihood invariant under H₀))

```
## [1] 6.883624
```

```

#_permutation
# the idea is to permute the residuals under H0:
# km = mu + alpha*station + beta*fuel
# additive model
aov.H0station_fuel <- aov(km ~ station + fuel)
aov.H0station_fuel

```

reduced model

```

## Call:
##   aov(formula = km ~ station + fuel)
##
## Terms:
##   station      fuel Residuals
## Sum of Squares 1.53125 66.70125 16.36625
## Deg. of Freedom 1          1          5
##
## Residual standard error: 1.809213
## Estimated effects may be unbalanced

```

```

residuals.H0station_fuel <- aov.H0station_fuel$residuals
n <- 8
permutation <- sample(1:n)
residuals.H0station_fuel <- residuals.H0station_fuel[permutation]
# permuted y values:
km.perm.H0station_fuel <- aov.H0station_fuel$fitted + residuals.H0station_fuel
summary.aov(aov(km.perm.H0station_fuel ~ station + fuel + station:fuel))

```

residuals

permutation

$\hat{y} + \text{permuted residuals}$
new model (complete)

```

##           Df Sum Sq Mean Sq F value    Pr(>F)
## station     1  3.38   3.38  1.817 0.24891
## fuel        1 58.59  58.59 31.503 0.00495 **
## station:fuel 1  8.30   8.30  4.464 0.10216
## Residuals    4  7.44   1.86
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

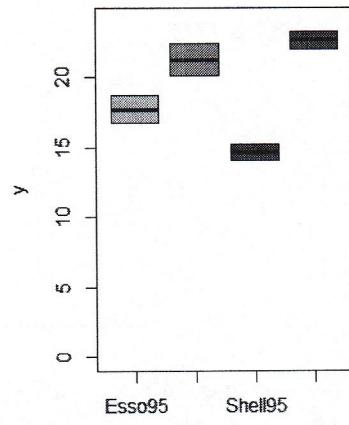
```

```

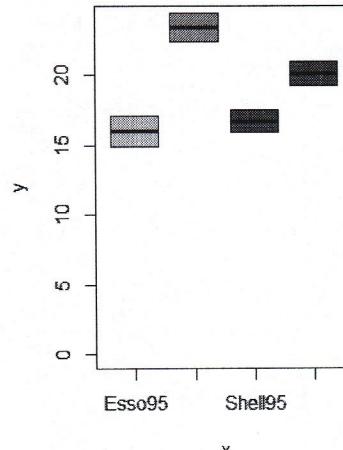
# How data has changed?
layout(rbind(1:2))
plot(station_fuel, km, col=rainbow(5)[2:5], ylim=c(0,24), main='Original data')
plot(station_fuel, km.perm.H0station_fuel, col=rainbow(5)[2:5], ylim=c(0,24), main='Permuted data')

```

Original data



Permuted data



```

# TEST of interaction
B <- 1000
T_station_fuel <- numeric(B)
for(perm in 1:B){
  permutation <- sample(n)
  residuals.H0station_fuel <- residuals.H0station_fuel[permutation]
  km.perm.H0station_fuel <- aov.H0station_fuel$fitted + residuals.H0station_fuel
  T_station_fuel[perm] <- summary.aov(aov(km.perm.H0station_fuel ~ station + fuel + station:fuel))[[1]][3,4]
}

# p-value
sum(T_station_fuel >= T0_station_fuel)/B

```

[1] 0.08 \Rightarrow we accept H_0 : we get rid of the interaction

```

# The interaction is not significant.
# We can remove it and perform a test for the two main effects

```

```

# TEST OF FACTOR STATION (H0: alpha=0)
T0_station <- summary.aov(aov(km ~ station + fuel))[[1]][1,4]
# residuals under H0:
# km = mu + beta*fuel
aov.H0station <- aov(km ~ fuel)
residuals.H0station <- aov.H0station$residuals
# permuted y values:
km.perm.H0station <- aov.H0station$fitted + residuals.H0station[permutation]
summary.aov(aov(km.perm.H0station ~ station + fuel))

```

(2)

```

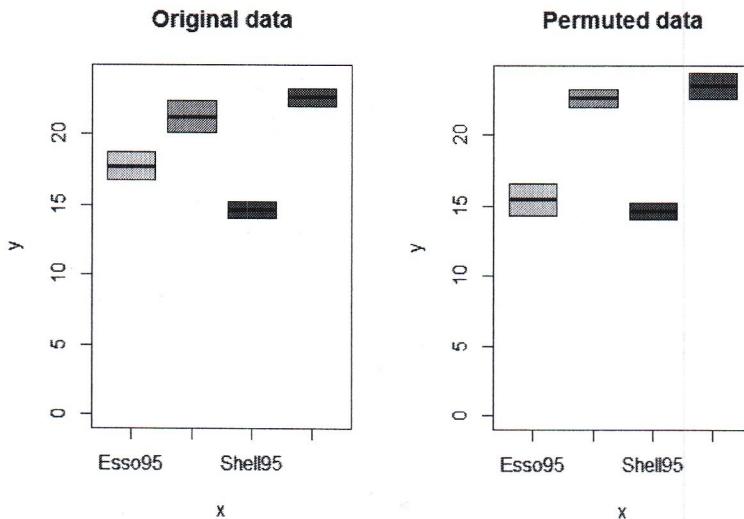
##          Df Sum Sq Mean Sq F value    Pr(>F)
## station     1   0.00   0.00   0.00 1.000000
## fuel        1 129.61 129.61  85.87 0.000246 ***
## Residuals    5   7.55   1.51
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

# How data has changed?
layout(rbind(1:2))
plot(station_fuel, km, col=rainbow(5)[2:5], ylim=c(0,24), main='Original data')
plot(station_fuel, km.perm.H0station, col=rainbow(5)[2:5], ylim=c(0,24), main='Permuted data')

```



they seem very similar.
If this happens for a lot of permutations \Rightarrow the factor station is not significant

(3)

```

# TEST OF FACTOR FUEL (H0: beta=0)
T0_fuel <- summary.aov(aov(km ~ station + fuel))[[1]][2,4]
# residuals under H0:
# km = mu + alpha*station
aov.H0fuel <- aov(km ~ station)
residuals.H0fuel <- aov.H0fuel$residuals
# permuted y values:
km.perm.H0fuel <- aov.H0fuel$fitted + residuals.H0fuel[permutation]
summary.aov(aov(km.perm.H0fuel ~ station + fuel))

```

```

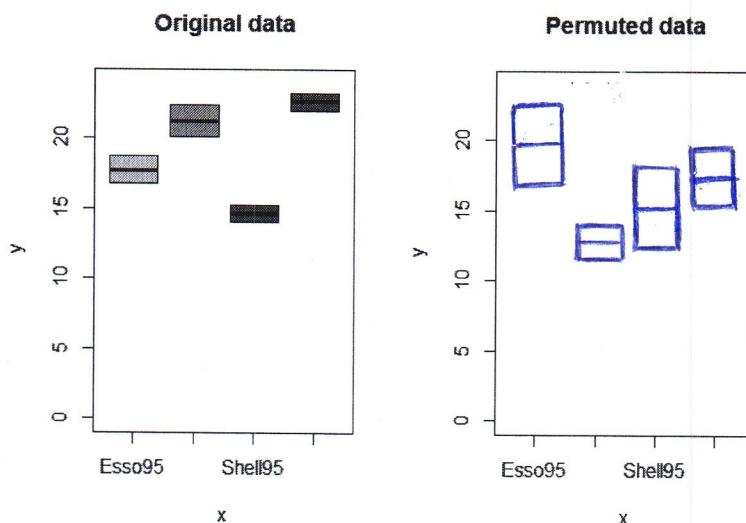
##          Df Sum Sq Mean Sq F value    Pr(>F)
## station     1  88.44  88.44  73.520 0.000356 ***
## fuel        1   10.35   10.35   8.605 0.032506 *
## Residuals    5   6.01   1.20
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

# How data has changed?
layout(rbind(1:2))
plot(station_fuel, km, col=rainbow(5)[2:5], ylim=c(0,24), main='Original data')
plot(station_fuel, km.perm.H0fuel, col=rainbow(5)[2:5], ylim=c(0,24), main='Permuted data')

```



There are a lot of differences:
if this is the case for many permutations \Rightarrow the factor fuel is significant

(2) + (3)

```
# TEST OF FACTOR STATION AND TEST OF FACTOR FUEL
# p-values
B <- 1000
T_station <- T_station <- numeric(B)
for(perm in 1:B){
  permutation <- sample(n)

  km.perm.H0station <- aov.H0station$fitted + residuals.H0station[permutation]
  T_station[perm] <- summary.aov(aov(km.perm.H0station ~ station + fuel))[[1]][1,4]

  km.perm.H0fuel <- aov.H0fuel$fitted + residuals.H0fuel[permutation]
  T_fuel[perm] <- summary.aov(aov(km.perm.H0fuel ~ station + fuel))[[1]][2,4]
}

sum(T_station >= T0_station)/B

## [1] 0.565

sum(T_fuel >= T0_fuel)/B

## [1] 0.028

# Comparison with the parametric test
summary.aov(aov(km ~ station + fuel))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)						
## station	1	1.53	1.53	0.468	0.52440						
## fuel	1	66.70	66.70	20.378	0.00632 **						
## Residuals	5	16.37	3.27								
## ---											
## Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'. '	0.1	' '	1

(4)

```
# We can remove also the factor station
# TEST ON THE FACTOR FUEL
T0_fuel <- summary.aov(aov(km ~ fuel))[[1]][1,4] → this time the model (reduced)
# residuals under H0
# km = mu
residuals.H0fuel <- km - M

# Note that in this case, permuting the residuals under H0
# and permuting the data is exactly the same:
permutation <- sample(n)
km.perm.H0fuel <- M + residuals.H0fuel[permutation]
km.perm <- km[permutation]

km.perm.H0fuel
```

```
## [1] 22.0 16.8 14.0 20.1 18.7 15.2 23.3 22.4
```

```
km.perm
```

```
## [1] 22.0 16.8 14.0 20.1 18.7 15.2 23.3 22.4
```

```
# CMC to estimate the p-value
B <- 1000
T_fuel <- numeric(B)
for(perm in 1:B){
  permutation <- sample(n)
  km.perm <- km[permutation]
  T_fuel[perm] <- summary.lm(aov(km.perm ~ fuel ))$f[1]

}
sum(T_fuel >= T0_fuel)/B
```

```
## [1] 0.008
```

```
### -  
## -  
### PERMUTATION TEST: Linear models
```

```
## -  
## -
```

```
# Linear models
```

```
# Example on simulated data
```

```
set.seed(24021979)
n <- 50
# covariate values
x1 <- runif(n,0,10)
x2 <- (1:n)/5
x3 <- rnorm(n,5,5)

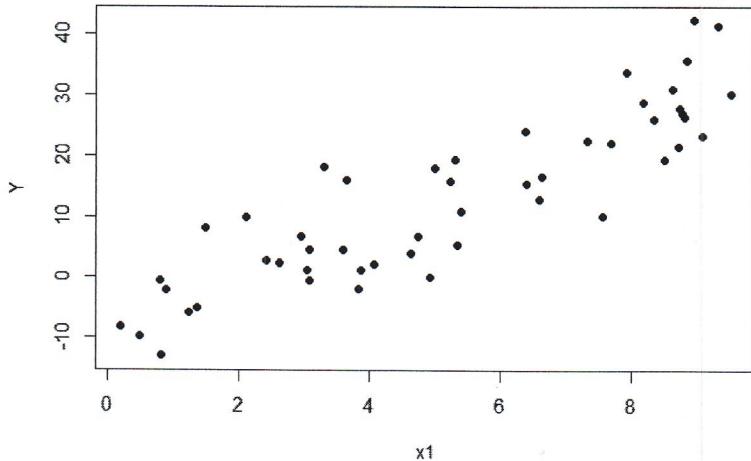
# generating model
b0 <- 2
b1 <- 4
b2 <- -2
b3 <- 0
Y <- b0 + b1*x1 + b2*x2 + b3*x3 + runif(n,-5,5)
plot(x1,Y,pch=16)
```

- Global test (all regressors): permutation of the response
- Partial test (one regressor): ASYMPTOTICALLY EXACT
we work on the residuals (permutations)

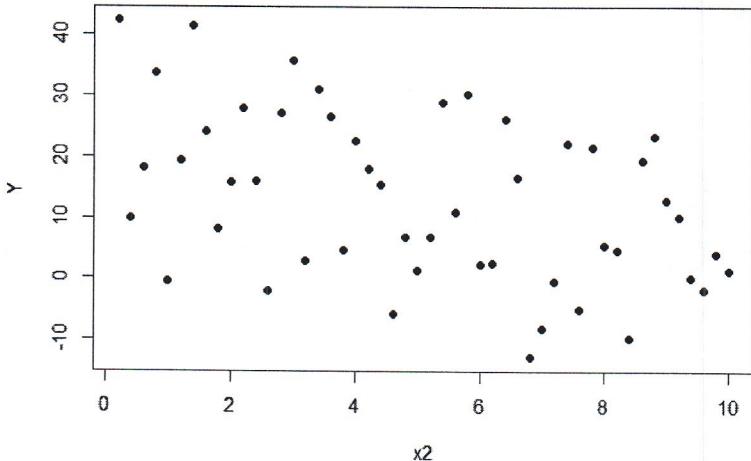
$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \epsilon$$

we're trying to get rid
of the distribution of
the error

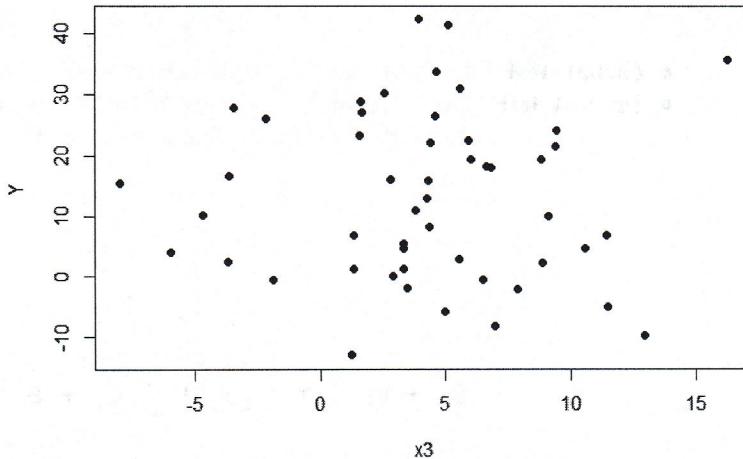
we're assuming all the
errors to have the same
distribution, just not
gaussian !



```
plot(x2,Y,pch=16)
```



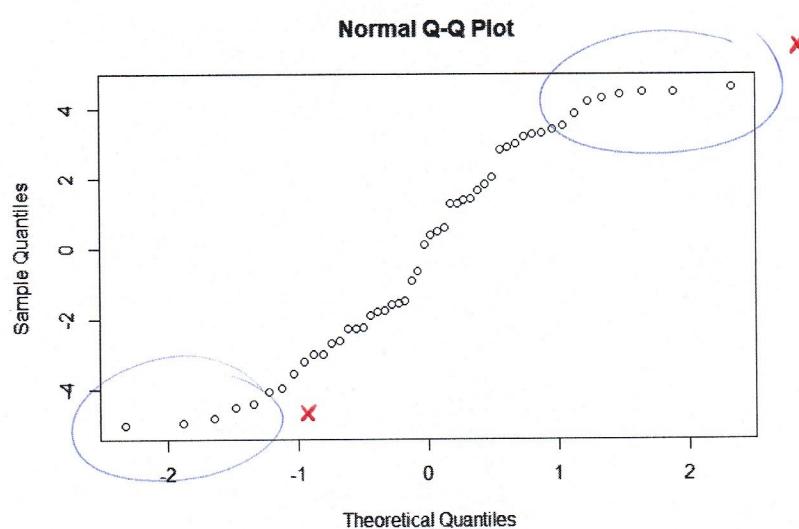
```
plot(x3,Y,pch=16)
```



```
# parametric inference (as comparison)
result <- lm(Y ~ x1 + x2 + x3)
summary(result)
```

```
## 
## Call:
## lm(formula = Y ~ x1 + x2 + x3)
## 
## Residuals:
##   Min     1Q Median     3Q    Max 
## -5.0412 -2.5301  0.2557  2.9685  4.6008 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 0.18206   1.44086   0.126   0.900    
## x1          4.17597   0.16073  25.981 < 2e-16 ***
## x2         -1.86077   0.15955 -11.663  2.45e-15 ***
## x3          0.13202   0.09578   1.378   0.175    
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 3.185 on 46 degrees of freedom
## Multiple R-squared:  0.9493, Adjusted R-squared:  0.946 
## F-statistic: 287.1 on 3 and 46 DF, p-value: < 2.2e-16 ✓
```

```
qqnorm(result$residuals)
```



```
shapiro.test(result$residuals)
```

```
## 
## Shapiro-Wilk normality test
## 
## data: result$residuals
## W = 0.93082, p-value = 0.005935
```

(Global)

```
# permutation inference
# we want to perform different tests
```

```
# Overall model
# H0: beta1 = beta2 = beta3 = 0
# test statistic
T0_glob <- summary(result)$f[1]
T0_glob
```

```
## value
## 287.1341
```



```
# permutations
permutazione <- sample(n)
Y.perm.glob <- Y[permutazione]
```

in this case permuting the responses or the residuals is the same

```
res.H0glob <- Y - mean(Y)
Y.perm.glob
```

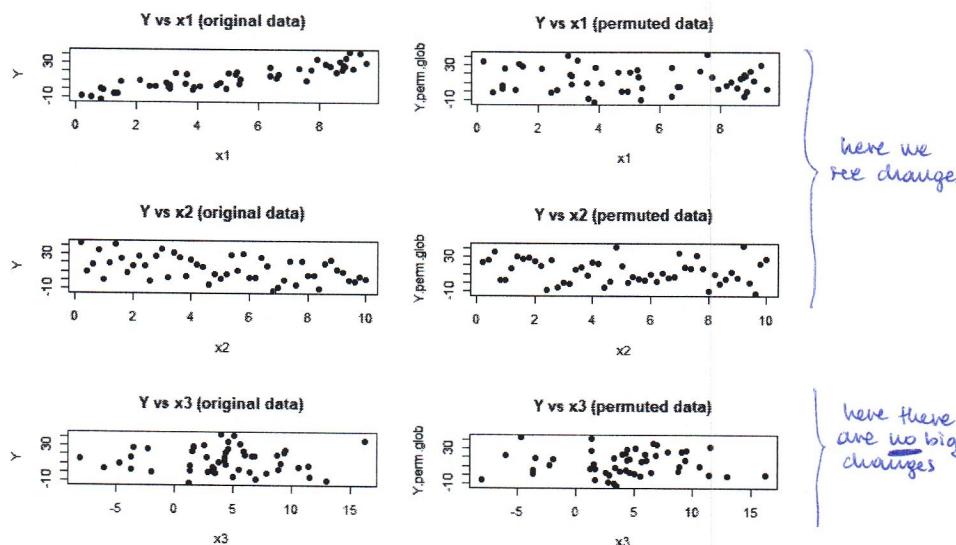
```
## [1] 22.48241717 26.38239058 35.74159214 2.21250032 2.40708113
## [6] 15.92132568 30.09713093 27.05470139 28.79799953 24.14881023
## [11] 19.37519779 -8.25398893 26.04981205 -5.07195264 -0.43453339
## [16] -1.87633222 15.37941576 18.22361867 8.08065421 23.37314319
## [21] 22.10003708 -5.75465816 1.33931967 41.39818009 19.34967537
## [26] -0.44439887 6.67252665 3.99175228 2.76803385 9.95027224
## [31] 1.32310819 10.94871942 4.58105052 6.74628679 33.66110255
## [36] 18.01712082 16.52693700 31.04898964 15.87302185 -9.81517813
## [41] 10.04966981 -2.17551423 4.53115582 12.79666281 5.43513788
## [46] 42.28576482 0.04698356 -12.97294939 21.56536855 27.77744131
```

```
mean(Y) + res.H0glob[permutazione]
```

```
## [1] 22.48241717 26.38239058 35.74159214 2.21250032 2.40708113
## [6] 15.92132568 30.09713093 27.05470139 28.79799953 24.14881023
## [11] 19.37519779 -8.25398893 26.04981205 -5.07195264 -0.43453339
## [16] -1.87633222 15.37941576 18.22361867 8.08065421 23.37314319
## [21] 22.10003708 -5.75465816 1.33931967 41.39818009 19.34967537
## [26] -0.44439887 6.67252665 3.99175228 2.76803385 9.95027224
## [31] 1.32310819 10.94871942 4.58105052 6.74628679 33.66110255
## [36] 18.01712082 16.52693700 31.04898964 15.87302185 -9.81517813
## [41] 10.04966981 -2.17551423 4.53115582 12.79666281 5.43513788
## [46] 42.28576482 0.04698356 -12.97294939 21.56536855 27.77744131
```

```
layout(matrix(1:6,nrow=3,byrow=FALSE))
plot(x1,Y,main='Y vs x1 (original data)',pch=16)
plot(x2,Y,main='Y vs x2 (original data)',pch=16)
plot(x3,Y,main='Y vs x3 (original data)',pch=16)
plot(x1,Y.perm.glob,main='Y vs x1 (permuted data)',pch=16)
plot(x2,Y.perm.glob,main='Y vs x2 (permuted data)',pch=16)
plot(x3,Y.perm.glob,main='Y vs x3 (permuted data)',pch=16)
```

+ pairs(cbind(Y,x1,x2,x3), main="Original Data")
+ pairs(cbind(Y.perm.glob, x1, x2, x3),
main="Permutated data")
(and we see that there is a difference with the perms.)



(Partial)

```
# Test on variable x1
# H0: beta1 = 0
# test statistic
summary(result)$coefficients
```

$$H_0: \beta_1 = 0$$

→ we permute the residuals
(we would like to permute the errors, but we don't have them)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.1820590	1.44085541	0.1263548	9.000018e-01
x1	4.1759699	0.16073115	25.9810872	3.913695e-29
x2	-1.8607723	0.15955127	-11.6625349	2.451497e-15
x3	0.1320238	0.09577616	1.3784623	1.747282e-01

```

T0_x1 <- abs(summary(result)$coefficients[2,3])
T0_x1

## [1] 25.98109

# permutations
# residuals of the reduced model

# reduced model:
#  $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3$ 
regr.H01 <- lm(Y ~ x1 + x3)
residui.H01 <- regr.H01$residuals
residui.H01.perm <- residui.H01[perm]

# permuted y:
Y.perm.H01 <- regr.H01$fitted + residui.H01.perm

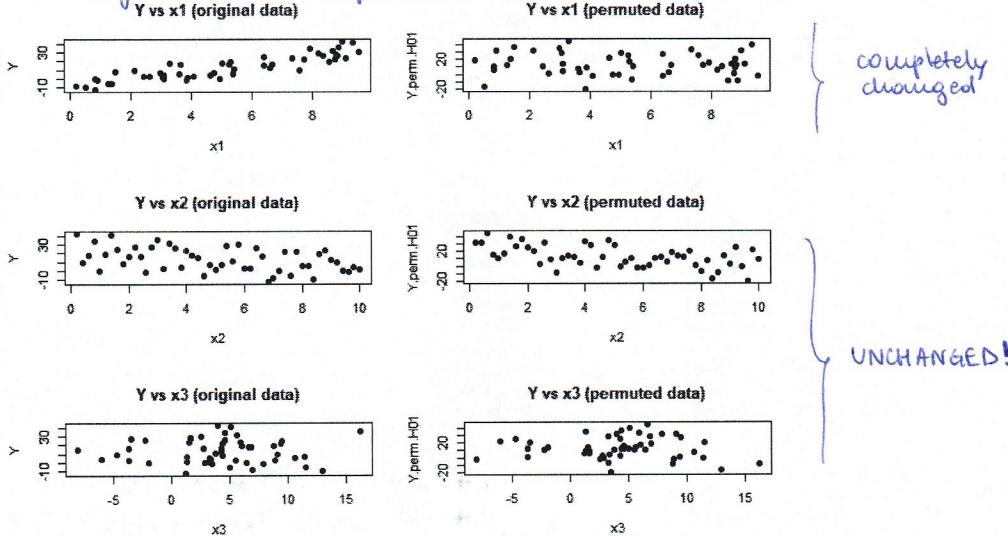
```

```

layout(matrix(1:6,nrow=3,byrow=FALSE))
plot(x1,Y,main='Y vs x1 (original data)',pch=16)
plot(x2,Y,main='Y vs x2 (original data)',pch=16)
plot(x3,Y,main='Y vs x3 (original data)',pch=16)
plot(x1,Y.perm.H01,main='Y vs x1 (permuted data)',pch=16)
plot(x2,Y.perm.H01,main='Y vs x2 (permuted data)',pch=16)
plot(x3,Y.perm.H01,main='Y vs x3 (permuted data)',pch=16)

```

Comparison: original dataset vs. permuted one:



Idea: (Friedman-Leverick idea)

- fit the reduced model (without x_1)
- compute the fitted values of the model
- compute the residuals
- permute the residuals
- sum the permuted residuals to the fitted values (*) and obtain another dataset
- introduce x_1 : fit the whole model (complete model)

Test statistic: $\hat{\beta}_1 / \hat{\beta}_2^2 / \text{std}(\hat{\beta}_1) / \dots$

Idea 2: permute the residuals of the complete model

- fit the model with x_1
- fitted values
- residuals
- permutation of the residuals
- fitted values + permuted residuals

(Partial)

```
# Test on variable x2
# H0:  $\beta_2 = 0$ 
```

$$H_0: \beta_2 = 0$$

```
# test statistic
summary(result)$coefficients
```

```

##             Estimate Std. Error    t value   Pr(>|t|)
## (Intercept) 0.1820590 1.44085541  0.1263548 9.000018e-01
## x1          4.1759699 0.16073115 25.9810872 3.913695e-29
## x2         -1.8607723 0.15955127 -11.6625349 2.451497e-15
## x3          0.1320238 0.09577616  1.3784623 1.747282e-01

```

```
T0_x2 <- abs(summary(result)$coefficients[3,3])
T0_x2
```

```
## [1] 11.66253
```

```
# permutations
# residuals of the reduced model
```

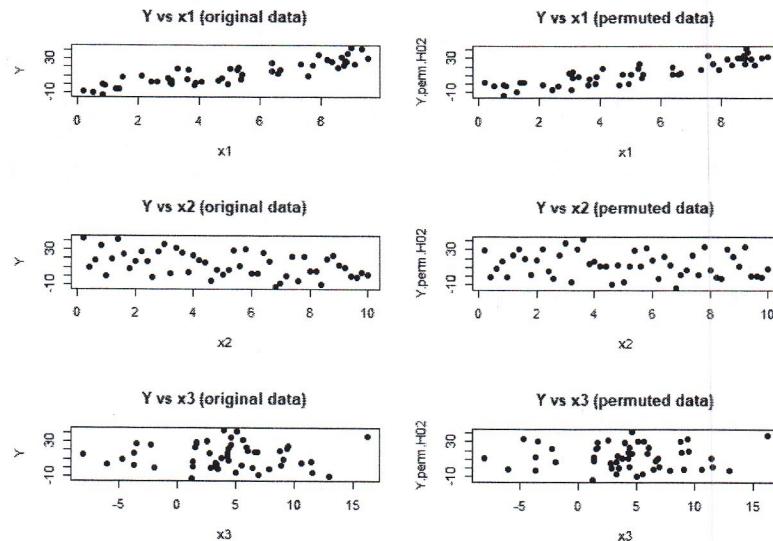
```

# reduced model:
#  $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3$ 
regr.H02 <- lm(Y ~ x1 + x3)
residui.H02 <- regr.H02$residuals
residui.H02.perm <- residui.H02[perm]

# permuted y:
Y.perm.H02 <- regr.H02$fitted + residui.H02.perm

layout(matrix(1:6,nrow=3,byrow=FALSE))
plot(x1,Y,main='Y vs x1 (original data)',pch=16)
plot(x2,Y,main='Y vs x2 (original data)',pch=16)
plot(x3,Y,main='Y vs x3 (original data)',pch=16)
plot(x1,Y.perm.H02,main='Y vs x1 (permuted data)',pch=16)
plot(x2,Y.perm.H02,main='Y vs x2 (permuted data)',pch=16)
plot(x3,Y.perm.H02,main='Y vs x3 (permuted data)',pch=16)

```



UNCHANGED

changed!

UNCHANGED

(Partial)

```
# Test on variable x3
# H0: beta3 = 0

# test statistic
summary(result)$coefficients

##             Estimate Std. Error    t value   Pr(>|t|)
## (Intercept) 0.1820590 1.44085541  0.1263548 9.000018e-01
## x1          4.1759699 0.16073115 25.9810872 3.913695e-29
## x2         -1.8607723 0.15955127 -11.6625349 2.451497e-15
## x3          0.1320238 0.09577616  1.3784623 1.747282e-01

T0_x3 <- abs(summary(result)$coefficients[4,3])
T0_x3

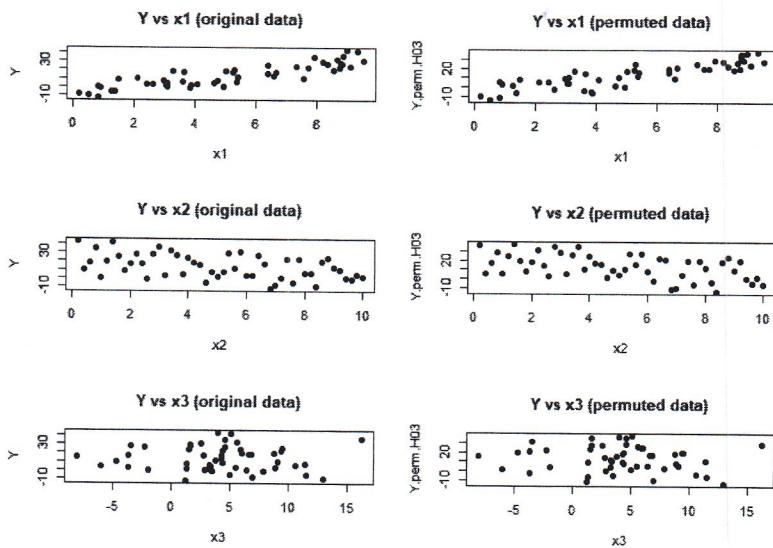
## [1] 1.378462

# permutations
# residuals of the reduced model

# reduced model:
# Y = beta0 + beta1*x1 + beta2*x2
regr.H03 <- lm(Y ~ x1 + x2)
residui.H03 <- regr.H03$residuals
residui.H03.perm <- residui.H03[permutazione]

# permuted y:
Y.perm.H03 <- regr.H03$fitted + residui.H03.perm

layout(matrix(1:6,nrow=3,byrow=FALSE))
plot(x1,Y,main='Y vs x1 (original data)',pch=16)
plot(x2,Y,main='Y vs x2 (original data)',pch=16)
plot(x3,Y,main='Y vs x3 (original data)',pch=16)
plot(x1,Y.perm.H03,main='Y vs x1 (permuted data)',pch=16)
plot(x2,Y.perm.H03,main='Y vs x2 (permuted data)',pch=16)
plot(x3,Y.perm.H03,main='Y vs x3 (permuted data)',pch=16)
```



UNCHANGED

changed a little

```

# p-values of the tests
B <- 1000
T_H0glob <- T_H01 <- T_H02 <- T_H03 <- numeric(B)

for(perm in 1:B){
  permutazione <- sample(n)

  Y.perm.glob <- Y[permutazione]
  T_H0glob[perm] <- summary(lm(Y.perm.glob ~ x1 + x2 + x3))$f[1]

  residui.H01.perm <- residui.H01[permutazione]
  Y.perm.H01 <- regr.H01$fitted + residui.H01.perm
  T_H01[perm] <- abs(summary(lm(Y.perm.H01 ~ x1 + x2 + x3))$coefficients[2,3])

  residui.H02.perm <- residui.H02[permutazione]
  Y.perm.H02 <- regr.H02$fitted + residui.H02.perm
  T_H02[perm] <- abs(summary(lm(Y.perm.H02 ~ x1 + x2 + x3))$coefficients[3,3])

  residui.H03.perm <- residui.H03[permutazione]
  Y.perm.H03 <- regr.H03$fitted + residui.H03.perm
  T_H03[perm] <- abs(summary(lm(Y.perm.H03 ~ x1 + x2 + x3))$coefficients[4,3])
}

sum(T_H0glob>=T0_glob)/B

```

```
## [1] 0 ✓
```

```
sum(T_H01>=T0_x1)/B
```

```
## [1] 0 ✓
```

```
sum(T_H02>=T0_x2)/B
```

```
## [1] 0 ✓
```

```
sum(T_H03>=T0_x3)/B
```

```
## [1] 0.168 ✗
```

comparison with the parametric test
`summary(result)`

```

## 
## Call:
## lm(formula = Y ~ x1 + x2 + x3)
## 
## Residuals:
##   Min     1Q Median     3Q    Max 
## -5.0412 -2.5301  0.2557  2.9685  4.6008 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 0.18206   1.44086   0.126   0.900    
## x1          4.17597   0.16073  25.981 < 2e-16 *** ✓
## x2         -1.86077   0.15955 -11.663 2.45e-15 *** ✓
## x3          0.13202   0.09578   1.378   0.175    
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 3.185 on 46 degrees of freedom
## Multiple R-squared:  0.9493, Adjusted R-squared:  0.946 
## F-statistic: 287.1 on 3 and 46 DF,  p-value: < 2.2e-16 ✓

```

```

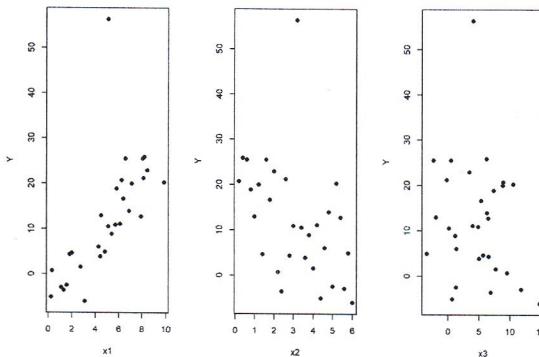
### -----
### -----
### PERMUTATION TEST: Linear models
### -----
### -
# Linear models
# Example on simulated data

set.seed(1992)
n <- 30
# covariate values
x1 <- runif(n,0,10)
x2 <- (1:n)/5
x3 <- rnorm(n,5,5)

# generating model
b0 <- 2
b1 <- 3
b2 <- -2
b3 <- 0
Y <- b0 + b1*x1 + b2*x2 + b3*x3 + stabledist::rstable(n,1.2,0)

par(mfrow=c(1,3))
plot(x1,Y,pch=16)
plot(x2,Y,pch=16)
plot(x3,Y,pch=16)

```



```

# parametric inference
result <- lm(Y ~ x1 + x2 + x3)
summary(result)

```

```

##
## Call:
## lm(formula = Y ~ x1 + x2 + x3)
##
## Residuals:
##   Min   1Q Median   3Q   Max
## -5.979 -2.335 -1.737 -0.260 44.004
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 5.2154    5.5706   0.936 0.357767
## x1          2.9268    0.6586   4.444 0.000146 ***
## x2         -2.1821    0.9874  -2.210 0.036115 *
## x3         -0.2114    0.3895  -0.543 0.591947
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 9.07 on 26 degrees of freedom
## Multiple R-squared:  0.5553, Adjusted R-squared:  0.504
## F-statistic: 10.82 on 3 and 26 DF, p-value: 8.522e-05

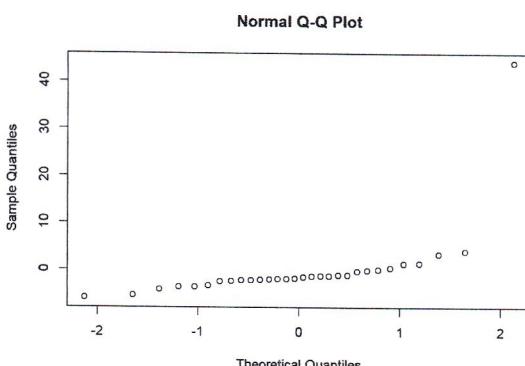
```

Let's now check if the assumptions of the model are met: normality of residuals?

```

par(mfrow=c(1,1))
qqnorm(result$residuals)

```



} it seems good. In the sense that we're rejecting where we should reject (β_3)

← This is telling us that we're doing something crazy

```
shapiro.test(result$residuals)
```

```

## Shapiro-Wilk normality test
## data: result$residuals
## W = 0.39707, p-value = 5.107e-10

```

the normality of the residuals is absolutely not respected

→ the p-values of the model do not make sense even if they lead to right conclusions in the sense that the statistical hypothesis on which the inference procedure is based are not satisfied

```

# permutation inference
# we want to perform different tests

#### -----
### Overall model
### -----
# H0: beta1 = beta2 = beta3 = 0
# test statistic
T0_glob <- summary(result)$f[1]
T0_glob

```

```

##     value
## 10.82281

```

```

# permutations
permutazione <- sample(n)
Y.perm.glob <- Y[permutazione]

# in this case permuting the responses or the residuals is the same
res.H0glob <- Y - mean(Y)
Y.perm.glob

```

```

## [1] 11.154834 13.947026 21.203688 4.953470 18.917391 20.797951 12.824425
## [8] 10.608779 -2.970116 -5.042114 8.864195 -6.038753 4.403893 6.054960
## [15] 1.580460 12.942796 10.966838 -3.562276 25.451919 4.598725 3.861026
## [22] -2.506088 20.276649 25.873357 0.717380 20.002413 25.430729 16.703224
## [29] 22.925617 56.280482

```

```

mean(Y) + res.H0glob[permutazione]

```

```

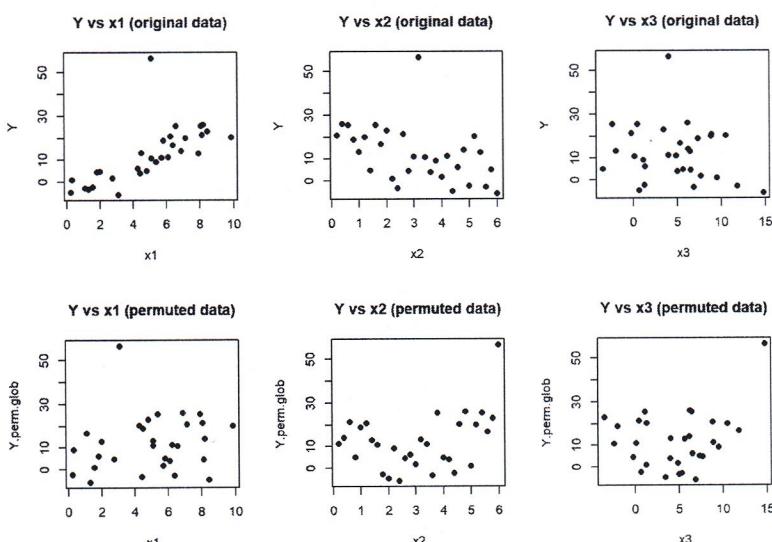
## [1] 11.154834 13.947026 21.203688 4.953470 18.917391 20.797951 12.824425
## [8] 10.608779 -2.970116 -5.042114 8.864195 -6.038753 4.403893 6.054960
## [15] 1.580460 12.942796 10.966838 -3.562276 25.451919 4.598725 3.861026
## [22] -2.506088 20.276649 25.873357 0.717380 20.002413 25.430729 16.703224
## [29] 22.925617 56.280482

```

```

par(mfrow=c(2,3))
plot(x1,Y,main='Y vs x1 (original data)',pch=16)
plot(x2,Y,main='Y vs x2 (original data)',pch=16)
plot(x3,Y,main='Y vs x3 (original data)',pch=16)
plot(x1,Y.perm.glob,main='Y vs x1 (permuted data)',pch=16)
plot(x2,Y.perm.glob,main='Y vs x2 (permuted data)',pch=16)
plot(x3,Y.perm.glob,main='Y vs x3 (permuted data)',pch=16)

```



```

#### -----
### Test on variable x1
### -----
# H0: beta1 = 0
# test statistic
summary(result)$coefficients

```

	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	5.2153780	5.5705734	0.9362372	0.3577665015
## x1	2.9267668	0.6585953	4.4439532	0.0001458356
## x2	-2.1821340	0.9874066	-2.2099650	0.0361154584
## x3	-0.2113903	0.3895040	-0.5427167	0.5919472104

```

T0_x1 <- abs(summary(result)$coefficients[2,3])
T0_x1

## [1] 4.443953

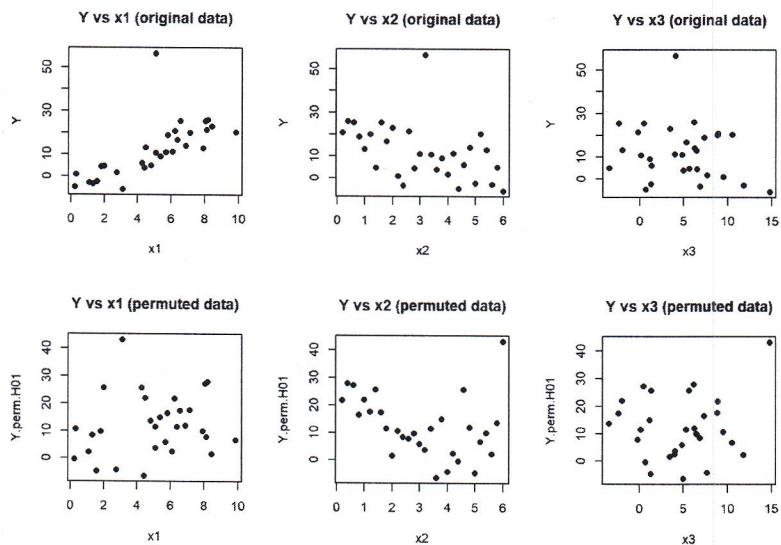
# permutations
# residuals of the reduced model

# reduced model:
#  $Y = \beta_0 + \beta_2 x_2 + \beta_3 x_3$ 
regr.H01 <- lm(Y ~ x2 + x3)
residui.H01 <- regr.H01$residuals
residui.H01.perm <- residui.H01[permutazione]

# permuted y:
Y.perm.H01 <- regr.H01$fitted + residui.H01.perm

par(mfrow=c(2,3))
plot(x1,Y,main='Y vs x1 (original data)',pch=16)
plot(x2,Y,main='Y vs x2 (original data)',pch=16)
plot(x3,Y,main='Y vs x3 (original data)',pch=16)
plot(x1,Y.perm.H01,main='Y vs x1 (permuted data)',pch=16)
plot(x2,Y.perm.H01,main='Y vs x2 (permuted data)',pch=16)
plot(x3,Y.perm.H01,main='Y vs x3 (permuted data)',pch=16)

```



from the first row to the second data changes a lot, this probably is because $\beta_1 \neq 0$

```

### -----
### Test on variable x2
### -----
# H0:  $\beta_2 = 0$ 
# test statistic
summary(result)$coefficients

```

```

##             Estimate Std. Error   t value   Pr(>|t|)    
## (Intercept) 5.2153780  5.5705734  0.9362372 0.3577665015
## x1          2.9267668  0.6585953  4.4439532 0.0001458356
## x2         -2.1821340  0.9874066 -2.2099650 0.0361154584
## x3         -0.2113903  0.3895040 -0.5427167 0.5919472104

```

```

T0_x2 <- abs(summary(result)$coefficients[3,3])
T0_x2

```

```

## [1] 2.209965

```

```

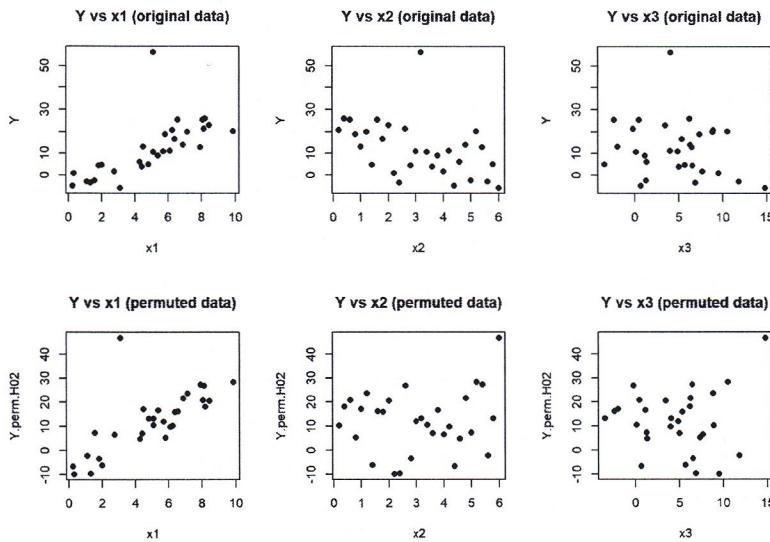
# permutations
# residuals of the reduced model

# reduced model:
#  $Y = \beta_0 + \beta_1 x_1 + \beta_3 x_3$ 
regr.H02 <- lm(Y ~ x1 + x3)
residui.H02 <- regr.H02$residuals
residui.H02.perm <- residui.H02[permutazione]

# permuted y:
Y.perm.H02 <- regr.H02$fitted + residui.H02.perm

par(mfrow=c(2,3))
plot(x1,Y,main='Y vs x1 (original data)',pch=16)
plot(x2,Y,main='Y vs x2 (original data)',pch=16)
plot(x3,Y,main='Y vs x3 (original data)',pch=16)
plot(x1,Y.perm.H02,main='Y vs x1 (permuted data)',pch=16)
plot(x2,Y.perm.H02,main='Y vs x2 (permuted data)',pch=16)
plot(x3,Y.perm.H02,main='Y vs x3 (permuted data)',pch=16)

```



```
### -----
### Test on variable x3
### -----
# H0: beta3 = 0
# test statistic
summary(result)$coefficients
```

```
##             Estimate Std. Error   t value   Pr(>|t|) 
## (Intercept) 5.2153780  5.5705734  0.9362372 0.3577665015
## x1          2.9267668  0.6585953  4.4439532 0.0001458356
## x2         -2.1821340  0.9874066 -2.2099650 0.0361154584
## x3         -0.2113903  0.3895040 -0.5427167 0.5919472104
```

```
T0_x3 <- abs(summary(result)$coefficients[4,3])
T0_x3
```

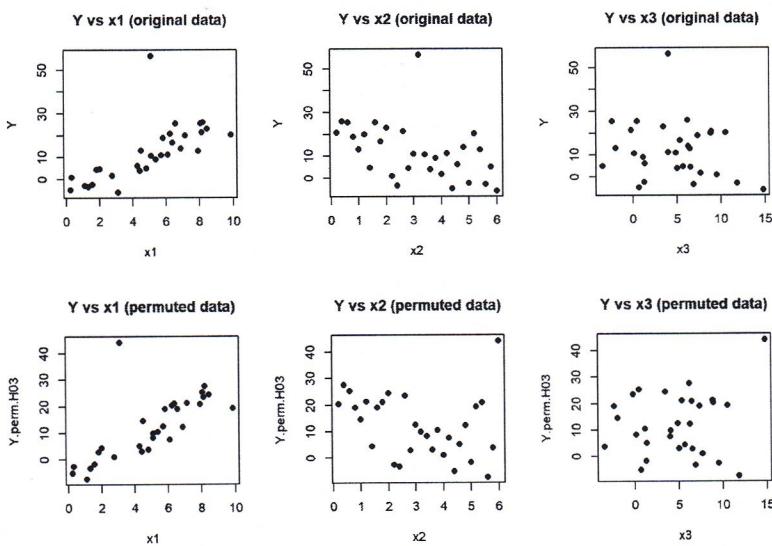
```
## [1] 0.5427167
```

```
# permutations
# residuals of the reduced model
```

```
# reduced model:
# Y = beta0 + beta1*x1 + beta2*x2
regr.H03 <- lm(Y ~ x1 + x2)
residui.H03 <- regr.H03$residuals
residui.H03.perm <- residui.H03[permutazione]

# permuted y:
Y.perm.H03 <- regr.H03$fitted + residui.H03.perm

par(mfrow=c(2,3))
plot(x1,Y,main='Y vs x1 (original data)',pch=16)
plot(x2,Y,main='Y vs x2 (original data)',pch=16)
plot(x3,Y,main='Y vs x3 (original data)',pch=16)
plot(x1,Y.perm.H03,main='Y vs x1 (permuted data)',pch=16)
plot(x2,Y.perm.H03,main='Y vs x2 (permuted data)',pch=16)
plot(x3,Y.perm.H03,main='Y vs x3 (permuted data)',pch=16)
```



from the first row to the second data does not really change, this probably will lead us to say $\beta_3 = 0$

```

#### -----
### p-values of the tests
#####
B <- 1000
T_H0glob <- T_H01 <- T_H02 <- T_H03 <- numeric(B)

for(perm in 1:B){
  permutazione <- sample(n)

  Y.perm.glob <- Y[permutazione]
  T_H0glob[perm] <- summary(lm(Y.perm.glob ~ x1 + x2 + x3))$f[1]

  residui.H01.perm <- residui.H01[permutazione]
  Y.perm.H01 <- regr.H01$fitted + residui.H01.perm
  T_H01[perm] <- abs(summary(lm(Y.perm.H01 ~ x1 + x2 + x3))$coefficients[2,3])

  residui.H02.perm <- residui.H02[permutazione]
  Y.perm.H02 <- regr.H02$fitted + residui.H02.perm
  T_H02[perm] <- abs(summary(lm(Y.perm.H02 ~ x1 + x2 + x3))$coefficients[3,3])

  residui.H03.perm <- residui.H03[permutazione]
  Y.perm.H03 <- regr.H03$fitted + residui.H03.perm
  T_H03[perm] <- abs(summary(lm(Y.perm.H03 ~ x1 + x2 + x3))$coefficients[4,3])
}

p-values of the tests:
sum(T_H0glob>=T0_glob)/B

## [1] 0

sum(T_H01>=T0_x1)/B

## [1] 0

sum(T_H02>=T0_x2)/B

## [1] 0.025

sum(T_H03>=T0_x3)/B

## [1] 0.61      → we accept  $\beta_3 = 0$ 

# comparison with the parametric test
summary(result)

```

```

##
## Call:
## lm(formula = Y ~ x1 + x2 + x3)
##
## Residuals:
##   Min     1Q Median     3Q    Max 
## -5.979 -2.335 -1.737 -0.260 44.004 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 5.2154    5.5706   0.936 0.357767    
## x1          2.9268    0.6586   4.444 0.000146 ***  
## x2         -2.1821    0.9874  -2.210 0.036115 *   
## x3          -0.2114    0.3895  -0.543 0.591947    
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
##
## Residual standard error: 9.07 on 26 degrees of freedom
## Multiple R-squared:  0.5553, Adjusted R-squared:  0.504 
## F-statistic: 10.82 on 3 and 26 DF,  p-value: 8.522e-05

```

General approach:

- calculate the statistic on the complete model
- calculate the residuals of the reduced model
- permute the residuals
 - ↑
↓
→ $\hat{y}_{\text{reduced}} + \text{permuted residuals}$
= new data
- test statistic on the complete model with ~~new data~~ new data
- p-value:
 $\frac{\text{sum}(\text{stat_perm} \geq \text{stat_original})}{\text{\# iterations (loop)}}$

on loop

```

### -----
### -----
### PERMUTATION TEST: Linear models - Homework
### -----
### -----
# Homework: find a test for  $H_0: \beta = \mu$  (where  $\mu$  is generic and not = 0)

# Example on simulated data
set.seed(1992)
n <- 30
x1 <- runif(n,0,10)
x2 <- (1:n)/5
x3 <- rnorm(n,5,5)

# Generating model
b0 <- 2
b1 <- 3
b2 <- -2
b3 <- 0
Y <- b0 + b1*x1 + b2*x2 + b3*x3 + stabledist::rstable(n,1.2,0)

### Test on variable x1
###  $H_0: \beta_1 = a$ 
###  $H_0: Y = b_0 + a*x1 + \beta_2*x2 + \beta_3*x3$ 
###  $H_0: Y - a*x1 = b_0 + \beta_2*x2 + \beta_3*x3$ 

### -----
a = 2.8 # the true value is a = 3
### -----
### Parametric inference
result <- lm(Y - a*x1 ~ x1 + x2 + x3)
summary(result)

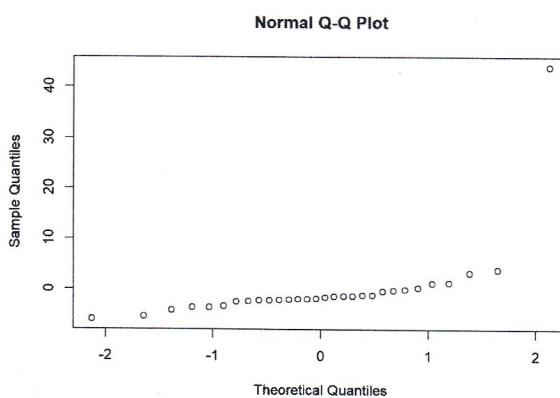
```

```

## 
## Call:
## lm(formula = Y - a * x1 ~ x1 + x2 + x3)
##
## Residuals:
##   Min     1Q Median     3Q    Max 
## -5.979 -2.335 -1.737 -0.260 44.004 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept)  5.2154    5.5706   0.936   0.3578    
## x1          0.1268    0.6586   0.192   0.8489    
## x2         -2.1821    0.9874  -2.210   0.0361 *  
## x3         -0.2114    0.3895  -0.543   0.5919    
## --- 
## Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 9.07 on 26 degrees of freedom
## Multiple R-squared:  0.1892, Adjusted R-squared:  0.09559 
## F-statistic: 2.022 on 3 and 26 DF, p-value: 0.1355

```

```
qqnorm(result$residuals)
```



```

shapiro.test(result$residuals)

## 
## Shapiro-Wilk normality test
## 
## data: result$residuals
## W = 0.39707, p-value = 5.107e-10

### Permutation inference
# test statistic
summary(result)$coefficients

```

```

##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 5.2153780 5.5705734 0.9362372 0.35776650
## x1          0.1267668 0.6585953 0.1924805 0.84886136
## x2         -2.1821340 0.9874066 -2.2099650 0.03611546
## x3         -0.2113903 0.3895040 -0.5427167 0.59194721

T0_x1 <- abs(summary(result)$coefficients[2,3])
T0_x1

## [1] 0.1924805

# permutations
# residuals of the reduced model

# reduced (H0) model:
# Y = beta0 + 3*x1 + beta2*x2 + beta3*x3
# Y - 3*x1 = beta0 + beta2*x2 + beta3*x3
regr.H01 <- lm(Y-a*x1 ~ x2 + x3)
residui.H01 <- regr.H01$residuals
permutazione <- sample(n)
residui.H01.perm <- residui.H01[permutazione]

# permuted Y:
Y.perm.H01 <- regr.H01$fitted + residui.H01.perm + a*x1

# p-values of the tests
B <- 1000
T_H01 <- numeric(B)

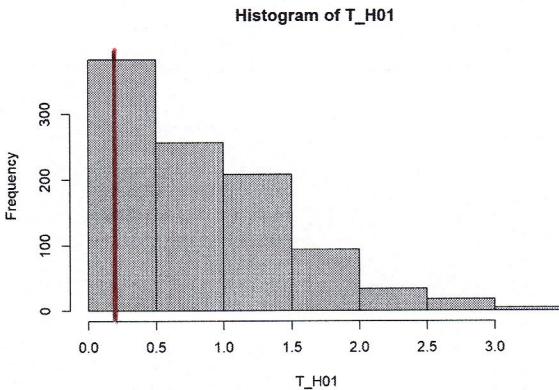
for(perm in 1:B){
  permutazione <- sample(1:n)
  residui.H01.perm <- residui.H01[permutazione]
  Y.perm.H01 <- regr.H01$fitted + residui.H01.perm + a*x1
  T_H01[perm] <- abs(summary(lm(Y.perm.H01 - a*x1 ~ x1 + x2 + x3))$coefficients[2,3])
}

hist(T_H01, xlim=range(c(T_H01, T0_x1)))
sum(T_H01>=T0_x1)/B

## [1] 0.819

abline(v=T0_x1, col=3, lwd=4)

```



```

# comparison with the parametric test
summary(result)

## 
## Call:
## lm(formula = Y - a * x1 ~ x1 + x2 + x3)
## 
## Residuals:
##   Min     1Q Median     3Q    Max 
## -5.979 -2.335 -1.737 -0.260 44.004 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 5.2154     5.5706   0.936   0.3578    
## x1          0.1268     0.6586   0.192   0.8489    
## x2         -2.1821     0.9874  -2.210   0.0361 *  
## x3         -0.2114     0.3895  -0.543   0.5919    
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
## 
## Residual standard error: 9.07 on 26 degrees of freedom
## Multiple R-squared:  0.1892, Adjusted R-squared:  0.09559 
## F-statistic: 2.022 on 3 and 26 DF,  p-value: 0.1355

```

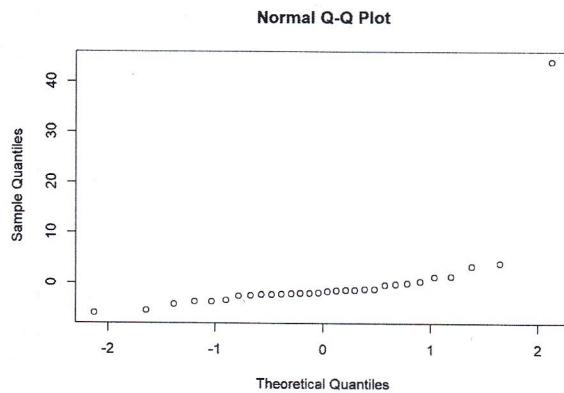
```

### -----
a = 1.8 # the true value is a = 3
### -----
### Parametric inference
result <- lm(Y ~ a*x1 ~ x1 + x2 + x3)
summary(result)

## 
## Call:
## lm(formula = Y ~ a * x1 ~ x1 + x2 + x3)
##
## Residuals:
##   Min     1Q Median     3Q    Max 
## -5.979 -2.335 -1.737 -0.260 44.004 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 5.2154    5.5706   0.936  0.3578    
## x1          1.1268    0.6586   1.711  0.0990 .  
## x2         -2.1821    0.9874  -2.210  0.0361 *  
## x3         -0.2114    0.3895  -0.543  0.5919    
## ---      
## Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 9.07 on 26 degrees of freedom
## Multiple R-squared:  0.2983, Adjusted R-squared:  0.2174 
## F-statistic: 3.685 on 3 and 26 DF,  p-value: 0.02463

```

```
qqnorm(result$residuals)
```



```
shapiro.test(result$residuals)
```

```

## 
## Shapiro-Wilk normality test
## 
## data: result$residuals
## W = 0.39707, p-value = 5.107e-10

```

```

### Permutation inference
# test statistic
summary(result)$coefficients

```

```

##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 5.2153780 5.5705734  0.9362372 0.35776650
## x1          1.1267668 0.6585953  1.7108636 0.09901390
## x2         -2.1821340 0.9874066 -2.2099650 0.03611546
## x3         -0.2113903 0.3895040 -0.5427167 0.59194721

```

```
T0_x1 <- abs(summary(result)$coefficients[2,3])
T0_x1
```

```
# [1] 1.710864
```

```

# permutations
# residuals of the reduced model

# reduced (H0) model:
# Y      = beta0 + beta1*x1 + beta2*x2 + beta3*x3
# Y - 3*x1 = beta0           + beta2*x2 + beta3*x3
regr.H01      <- lm(Y-a*x1 ~ x2 + x3)
residui.H01    <- regr.H01$residuals
permutazione   <- sample(n)
residui.H01.perm <- residui.H01[permutazione]

# permuted Y:
Y.perm.H01 <- regr.H01$fitted + residui.H01.perm + a*x1

# p-values of the tests
B      <- 1000
T_H01 <- numeric(B)

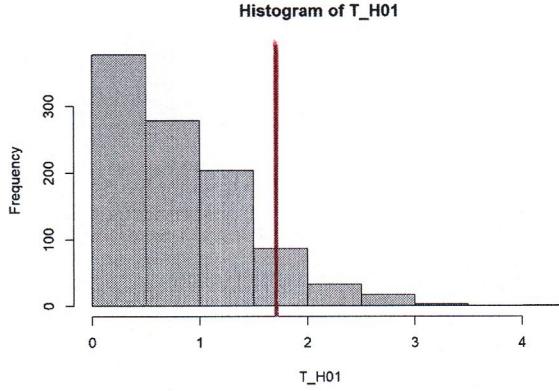
for(perm in 1:B){
  permutazione   <- sample(1:n)
  residui.H01.perm <- residui.H01[permutazione]
  Y.perm.H01      <- regr.H01$fitted + residui.H01.perm + a*x1
  T_H01[perm]     <- abs(summary(lm(Y.perm.H01 - a*x1 ~ x1 + x2 + x3))$coefficients[2,3])
}

hist(T_H01, xlim=range(c(T_H01, T0_x1)))
sum(T_H01>=T0_x1)/B

## [1] 0.09

abline(v=T0_x1, col=3, lwd=4)

```



```

# comparison with the parametric test
summary(result)

## 
## Call:
## lm(formula = Y - a * x1 ~ x1 + x2 + x3)
##
## Residuals:
##   Min     1Q Median     3Q    Max 
## -5.979 -2.335 -1.737 -0.260 44.004 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 5.2154    5.5706   0.936   0.3578    
## x1          1.1268    0.6586   1.711   0.0990 .  
## x2         -2.1821    0.9874  -2.210   0.0361 *  
## x3         -0.2114    0.3895  -0.543   0.5919    
## ---        
## Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
## 
## Residual standard error: 9.07 on 26 degrees of freedom
## Multiple R-squared:  0.2983, Adjusted R-squared:  0.2174 
## F-statistic: 3.685 on 3 and 26 DF,  p-value: 0.02463

```

this method is preferred

```

### -----
### ----- Correction of the assignment (alternative)
### -----
### -----
set.seed(1992)
n <- 30
# Covariate values
x1 <- runif(n,0,10)
x2 <- (1:n)/5
x3 <- rnorm(n,5,5)

# Generating model
b0 <- 2
b1 <- 3
b2 <- -2
b3 <- 0
Y <- b0 + b1*x1 + b2*x2 + b3*x3 + stabledist::rstable(n,1.2,0)
result <- lm(Y ~ x1 + x2 + x3)

### -----
### Test on variable x2
### H0: beta2 = -2
### -----
# Test statistic
mu0 <- -2
summary(result)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	5.2153780	5.5705734	0.9362372	0.3577665015
## x1	2.9267668	0.6585953	4.4439532	0.0001458356
## x2	-2.1821340	0.9874066	-2.2099650	0.0361154584
## x3	-0.2113903	0.3895040	-0.5427167	0.5919472104

- T0_X2 <- abs(summary(result)\$coefficients[3,3]-mu0) ← test statistic

```

## [1] 0.209965
```

```

# Permutations
# Residuals of the reduced model

# Reduced model:
# Y = beta0 + (mu0*x2) + beta1*x1 + beta3*x3
regr.H02 <- lm(Y ~ I(1+x2*mu0) + x1 + x3 - 1) # we replace the intercept with a fixed value
summary(regr.H02)
```

	Estimate	Std. Error	t value	Pr(> t)
## I(1 + x2 * mu0)	0.85980	0.39302	2.188	0.0375 *
## x1	3.32389	0.41780	7.956	1.5e-08 ***
## x3	-0.09765	0.35888	-0.272	0.7876
## ---				
## Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'
##	0.1	' '	1	

```

## Residual standard error: 9.005 on 27 degrees of freedom
## Multiple R-squared: 0.761, Adjusted R-squared: 0.7344
## F-statistic: 28.65 on 3 and 27 DF, p-value: 1.527e-08
```

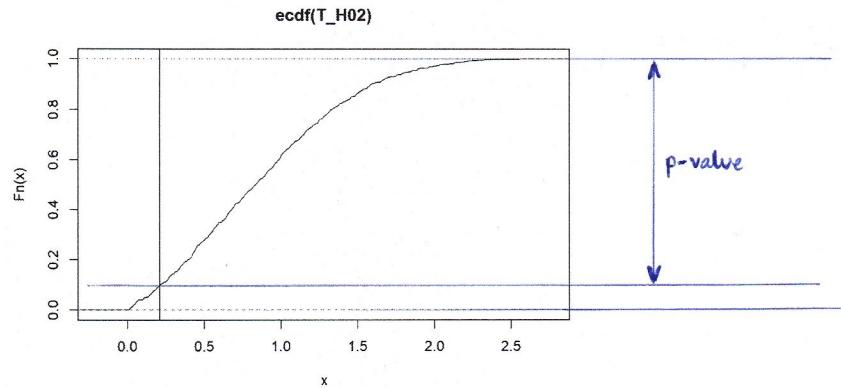
```

residui.H02 <- regr.H02$residuals

B <- 1000
T_H02 <- numeric(B)

for(perm in 1:B){
  permutazione <- sample(n)
  residui.H02.perm <- residui.H02[permutazione]
  Y.perm.H02 <- regr.H02$fitted + residui.H02.perm
  T_H02[perm] <- abs(summary(lm(Y.perm.H02 ~ x1 + x2 + x3))$coefficients[3,3]-mu0)
}

plot(ecdf(T_H02))
abline(v=T0_X2)
```



```
pval=(sum(T_H02>T0_x2))/B  
pval
```

```
## [1] 0.901
```

```

### -----
### PERMUTATIONAL (UNIVARIATE) CONFIDENCE INTERVAL
### -----
# Let's compute the confidence interval for the mean of a univariate population.
# The idea is to assign to all points in R a p value, and then threshold over that.

```

```

uni_t_perm = function(data,mu0,B=10000){
  data_trans = data-mu0
  T0         = abs(mean(data_trans))
  T_perm     = numeric(B)
  n          = length(data)
  for(perm in 1:B){
    refl      = rbinom(n, 1, 0.5)*2 - 1
    T_perm[perm] = abs(mean(data_trans*refl))
  }
  return(sum(T_perm>=T0)/B)
}

library(pbapply)
library(parallel)

grid = seq(-3,3,by=0.001)

cl = makeCluster(2)
clusterExport(cl,varlist=list("data","uni_t_perm"))

data = stabledist::rstable(30,1.5,0)

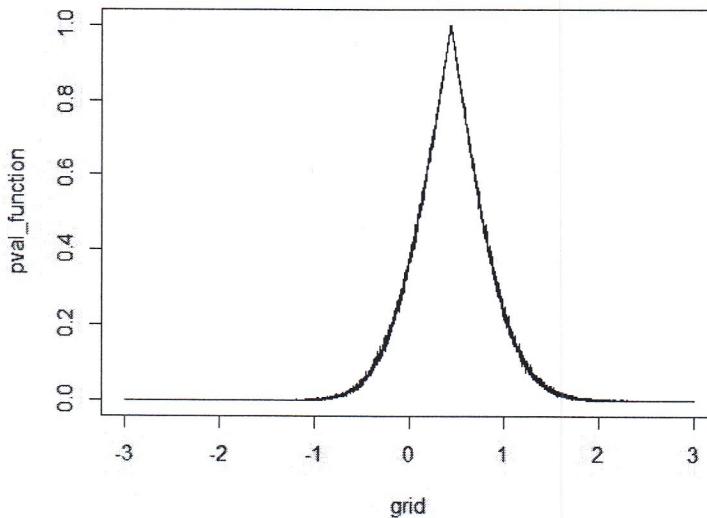
perm_wrapper = function(grid_point){uni_t_perm(data,grid_point,B=2000)}
pval_function = pbapply(grid,perm_wrapper,cl=cl)
pval_function
plot(grid,pval_function,type='l')

```

$$\begin{aligned} \text{confidence region} &= \{x : p\text{-value}(x) \geq \text{threshold}\} \\ &\subseteq \{x \text{ where we're not rejecting } H_0\} \end{aligned}$$

→ we look for a p-value function:
 $x \mapsto p\text{-value}(x)$

→ we test a grid of points and we see which tests are significant and which are not



```

range(grid[pval_function>0.05])
[1] -0.526   1.425

```